

Relational Algebra

Presented by Stéphane Bressan

Introduction to Database Systems

Relational Query Languages

- Two mathematical Query Languages form the basis for practical languages like SQL:
 - Relational Algebra: Operational**, useful for representing execution plans
 - Relational Calculus: Declarative: Describe what you want**, rather than how to compute it.
- Query languages are **NOT programming languages**:
 - Not expected to be Turing complete**

Introduction to Database Systems

Operations (Operators)

- Operations on a single relation
 - selection σ , projection π
- Usual set operations (relations are sets):
 - union \cup , intersection \cap , and difference $-$ (non-symmetric)
- Operations combining two or more relations
 - Cartesian product \times , join \bowtie and natural join \bowtie_n
- A renaming operation ρ
- A division operation $/$

Introduction to Database Systems

Example: employee

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

Introduction to Database Systems

Example: plane

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

Introduction to Database Systems

Example: canFly

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC9

Introduction to Database Systems

Example: assigned

eNumber	date	fNumber
1001	Nov 1	100
1001	Oct 31	100
1002	Nov 1	100
1002	Oct 31	100
1003	Oct 31	100
1003	Oct 31	337
1004	Oct 31	337
1005	Oct 31	337
1006	Nov 1	991
1006	Oct 31	337

Introduction to Database Systems

Projection

Keeps vertical slices of a relation according to a list L of attributes (i.e. **a list of columns**) of the **relation R**.

$$\pi_L(R) = \{t \mid \exists t_1$$

$$(t_1 \in R \wedge_{A \in L} t.A = t_1.A)\}$$

Calculus equivalent
Cutting a vertical slice of table
SELECT ...

Projection (Example)

$\pi_{\text{eNumber}, \text{fNumber}}$ (assigned)

eNumber	date	fNumber
1001	Nov 1	100
1001	Oct 31	100
1002	Nov 1	100
1002	Oct 31	100
1003	Oct 31	100
1003	Oct 31	337
1004	Oct 31	337
1005	Oct 31	337
1006	Nov 1	991
1006	Oct 31	337

Projection (Result)

$\pi_{\text{eNumber}, \text{fNumber}}$ (assigned)

eNumber	fNumber
1001	100
1002	100
1003	100
1003	337
1004	337
1005	337
1006	991
1006	337

Projection (SQL)

SELECT DISTINCT eNumber, fNumber
FROM assigned

same as

$\pi_{\text{eNumber}, \text{fNumber}}$ (assigned)

eliminate duplicates

Introduction to Database Systems

Selection

Selects the t-uples of a relation verifying a condition c:

$$\sigma_c(R) = \{t \mid t \in R \wedge c\}$$

c is any Boolean expression (\wedge , \vee , \neg) involving t ($<$, $=$, $>$, \neq , \leq , \geq)

calculus equivalent

Selection (Example)

$\sigma_{\text{salary} < 100000}(\text{employee})$

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

Selection (Result)

$\sigma_{\text{salary} < 100000}(\text{employee})$

name	salary	eNumber
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003

Selection (SQL)

SELECT *
FROM employee
WHERE salary < 100000

Same as

$\sigma_{\text{salary} < 100000}(\text{employee})$

Introduction to Database Systems

Selection (Example)

$\sigma_{\text{salary} > 100000 \wedge \neg(\text{name} = \text{'Gates'})}(\text{employee})$

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

Introduction to Database Systems

Selection (Result)

$\sigma_{\text{salary} > 100000 \wedge \neg(\text{name} = \text{'Gates'})}(\text{employee})$

name	salary	eNumber
Clark	150000	1006
Warnock	500000	1007

Introduction to Database Systems

Selection (SQL)

SELECT *
FROM employee
WHERE salary > 100000
AND name <> 'Gates'

Same as

$\sigma_{\text{salary} > 100000 \wedge \neg(\text{name} = \text{'Gates'})}(\text{employee})$

Introduction to Database Systems

Remark: Composability

The result of a query is a relation

$\sigma_{\text{salary} < 50000}(\text{employee})$

$\pi_{\text{name, salary}}(\sigma_{\text{salary} < 50000}(\text{employee}))$

Introduction to Database Systems

Remark: Commutativity

$\pi_{\text{name, salary}}(\sigma_{\text{salary} < 50000}(\text{employee}))$

$\sigma_{\text{salary} < 50000}(\pi_{\text{name, salary}}(\text{employee}))$

Can we always do this?

MUST BE VERY CAREFUL ABOUT ORDER!!

simple mod:

$\sigma_{\text{salary} < 100000}(\pi_{\text{name}}(\text{employee}))$

DOES NOT WORK!!

Introduction to Database Systems

Remark: SQL

$\pi_{\text{name, salary}}(\sigma_{\text{salary} < 50000}(\text{employee}))$

SELECT DISTINCT name, salary
FROM employee
WHERE salary < 50000

Projection

Selection

Introduction to Database Systems

Union, Intersection, Set-difference

- $R_1 \cup R_2 = \{t \mid t \in R_1 \vee t \in R_2\}$
- $R_1 \cap R_2 = \{t \mid t \in R_1 \wedge t \in R_2\}$
- $R_1 - R_2 = \{t \mid t \in R_1 \text{ and } \neg(t \in R_2)\}$

The relations R_1 and R_2 must be

union compatible:

Must have same col. & same type

- Same number of attributes
- Corresponding attributes have the same type (but not necessarily the same name)

*** Playing with sets: NO DUPLICATES!!**

Union (Example)

$\text{plane}_1 \cup \text{plane}_2$

plane_1

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Boeing	B747
Boeing	B757

plane_2

maker	mNumber
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
MD	DC10
MD	DC9

Introduction to Database Systems

Union (Result)

$\text{plane}_1 \cup \text{plane}_2$

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

Introduction to Database Systems

Union (SQL)

```
SELECT *  
FROM plane1  
UNION  
SELECT *  
FROM plane2
```

What about duplicates?

eliminated. both SQL & Algebra

Introduction to Database Systems

Intersection (Example)

$\text{plane}_1 \cap \text{plane}_2$

Plane₁

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Boeing	B747
Boeing	B757

Plane₂

maker	mNumber
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
MD	DC10
MD	DC9

Introduction to Database Systems

Intersection (Result)

$\text{plane}_1 \cap \text{plane}_2$

maker	mNumber
Airbus	A330
Boeing	B747

Introduction to Database Systems

Intersection (SQL)

```
SELECT *  
FROM plane1  
INTERSECT  
SELECT *  
FROM plane2
```

What about duplicates?

Also eliminated

Introduction to Database Systems

Difference (Example)

$\text{plane}_1 - \text{plane}_2$

Plane₁

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Boeing	B747
Boeing	B757

Plane₂

maker	mNumber
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
MD	DC10
MD	DC9

Introduction to Database Systems

Difference (Result)

$\text{plane}_1 - \text{plane}_2$

maker	mNumber
Airbus	A310
Airbus	A320
Boeing	B757

Introduction to Database Systems

ORACLE: MINUS

Difference (SQL)

```
SELECT *
FROM plane1
MINUS
SELECT *
FROM plane2
```

In other systems
EXCEPT

What about duplicates?

Also eliminated.

Introduction to Database Systems

Cartesian Product

Combines the t-uples of two relations in all possible ways

$$R_1 \times R_2 = \{t \mid \exists t_1 \exists t_2 \\ (t_1 \in R_1 \wedge t_2 \in R_2 \\ \wedge_{A \in R_1} t.A = t_1.A \\ \wedge_{A \in R_2} t.A = t_2.A)\}$$

Calculus equivalent

Introduction to Database Systems

Cartesian Product (Example)

canFly × plane

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

↓
10 tuples

↓
9 tuples

Introduction to Database Systems

Cartesian Product (Result)

canFly × plane

90 t-uples

↓
9 × 10 tuples

eNumber	mNumber	maker	mNumber
1001	B727	Airbus	A310
1001	B727	Airbus	A320
1001	B727	Airbus	A330
1001	B727	Airbus	A340
1001	B727	Boeing	B727
1001	B727	Boeing	B747
1001	B727	Boeing	B757
1001	B727	MD	DC10
1001	B727	MD	DC9
1001	B747	Airbus	A310
1001	B747	Airbus	A320
1001	B747	Airbus	A330
1001	B747	Airbus	A340
1001	B747	Boeing	B727
1001	B747	Boeing	B747
1001	B747	Boeing	B757
1001	B747	MD	DC10
1001	B747	MD	DC9
1001	B727	Airbus	A310
1001	B727	Airbus	A320
...

Introduction to Database Systems

Cartesian Product (SQL)

```
SELECT *
FROM canFly, plane
```

Introduction to Database Systems

Join (θ-Join)

Combines the t-uples of two relations that verify a condition → c

$$R_1 \bowtie_c R_2 = \{t \mid \exists t_1 \exists t_2 \\ (t_1 \in R_1 \wedge t_2 \in R_2 \wedge c \\ \wedge_{A \in R_1} t.A = t_1.A \\ \wedge_{A \in R_2} t.A = t_2.A)\}$$

$= \sigma_c(R_1 \times R_2)$ → what it can also mean

Introduction to Database Systems

Join (θ -Join) (Example)

canFly \bowtie canFly.mNnumber=plane.mNumber plane

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

Introduction to Database Systems

Join (Result)

canFly \bowtie canFly.mNnumber=plane.mNumber plane

eNumber	mNumber	maker	mNumber
1001	B727	Boeing	B727
1001	B747	Boeing	B747
1001	DC10	MD	DC10
1002	A320	Airbus	A320
1002	A340	Airbus	A340
1002	B757	Boeing	B757
1002	DC9	MD	DC9
1003	A310	Airbus	A310
1003	DC9	MD	DC9
1003	DC10	MD	DC10

Introduction to Database Systems

Join (SQL)

```
SELECT *
FROM canFly c, plane p
WHERE c.mNumber = p.mNumber
```

Introduction to Database Systems

Don't use equi-join, just know it exists!!

Equi-join

- Combines two relations on a condition composed only of equalities of attributes of the first and second relation
- Projects only one of the redundant attributes (since they are equal)

$$R1 \bowtie_{E(A1.1=A2.1 \wedge \dots \wedge A1.n=A2.n)} R2$$

Introduction to Database Systems

Equi-join (Example)

canFly $\bowtie_{E(\text{canFly.mNnumber}=\text{plane.mNumber})}$ plane

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

Introduction to Database Systems

Equi-join (Result)

canFly $\bowtie_{E(\text{canFly.mNnumber}=\text{plane.mNumber})}$ plane

eNumber	mNumber	maker
1001	B727	Boeing
1001	B747	Boeing
1001	DC10	MD
1002	A320	Airbus
1002	A340	Airbus
1002	B757	Boeing
1002	DC9	MD
1003	A310	Airbus
1003	DC9	MD
1003	DC10	MD

Introduction to Database Systems

Don't use it also... ٧

Natural Join

- Combines two relations on a condition composed only of equalities of attributes with the same name in the first and second relation
- Projects only one of the redundant attributes (since they are equal)

$$R_1 \bowtie_N R_2$$

Introduction to Database Systems

Natural Join (Example)

canFly \bowtie_N plane

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

Introduction to Database Systems

Natural Join (Result)

canFly \bowtie_N plane

eNumber	mNumber	maker
1001	B727	Boeing
1001	B747	Boeing
1001	DC10	MD
1002	A320	Airbus
1002	A340	Airbus
1002	B757	Boeing
1002	DC9	MD
1003	A310	Airbus
1003	DC9	MD

Introduction to Database Systems

Renaming

Renaming a relation or its attributes:

$$\rho(R'(N_1 \rightarrow N'_1, \dots, N_n \rightarrow N'_n), R)$$

The new relation R' has the same instance as R , but its schema has attribute N'_i instead of attribute N_i

Useful for : Combining same tables/
shorten names for
convenience

Introduction to Database Systems

Renaming (Example)

$\rho(\text{staff}(\text{salary} \rightarrow \text{wages}), \text{employee})$

employee

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

Introduction to Database Systems

Renaming (Result)

$\rho(\text{staff}(\text{salary} \rightarrow \text{wages}), \text{employee})$

name	wages	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

Introduction to Database Systems


```
SELECT name, salary AS wages, eNumber
FROM employee
```

Introductory to Database Systems

Find for each employee, her name and the model numbers of the planes she can fly

name	salary	eNumber	eNumber	mNumber
Clark	150000	1006	1001	B727
Gates	5000000	1005	1001	B747
Jones	50000	1001	1001	DC10
Peter	45000	1002	1002	A320
Phillips	25000	1004	1002	A340
Rowe	35000	1003	1002	B757
Warnock	500000	1007	1002	DC9
			1003	A310
			1003	DC9

Introduction to Database Systems

```

graph TD
    A["πname,mNumber"] --- B["σemployee.eNumber=canFly.eNumber"]
    B --- C["×"]
    C --- D["employee"]
    C --- E["canFly"]
  
```

Introductory to Database Systems

```
SELECT name, mNumber ← Projection
FROM employee, canFly ← Cartesian Product ✈️
WHERE
    employee.eNumber=canFly.eNumber
    ← Selection
```

Introduction to Database Systems

- $\pi_{\text{eNumber}} (\sigma_{\text{canFly.mNumber=plane.mNumber} \wedge \text{maker='Airbus'}} (\text{canFly} \times \text{plane}))$
Which one is more efficient?
- $\pi_{\text{eNumber}} (\sigma_{\text{canFly.mNumber=plane.mNumber}} (\text{canFly} \times \sigma_{\text{maker='Airbus'}} (\text{plane})))$
- $\pi_{\text{eNumber}} ((\text{canFly} \bowtie_{\text{canFly.mNumber=plane.mNumber}} \sigma_{\text{maker='Airbus'}} (\text{plane})))$
- $\pi_{\text{eNumber}} (\text{canFly} \bowtie_{\text{canFly.mNumber=plane.mNumber} \wedge \text{maker='Airbus'}} \text{plane})$

Different ways To write the same thing

Find the employee numbers of employees who can fly Airbus planes


eNumber	mNumber	maker	mNumber
1001	B727	Airbus	A310
1001	B747	Airbus	A320
1001	DC10	Airbus	A330
1002	A320	Airbus	A340
1002	A340	Boeing	B727
1002	B757	Boeing	B747
1002	DC9	Boeing	B757
1003	A310	MD	DC10
1003	DC9	MD	DC9
1003	DC10		

Introduction to Database Systems

```
SELECT eNumber
FROM employee e, confly c
WHERE e.mNumber = c.mNumber AND
      c.make = 'Airbus';
```

Example

Find the employee numbers of employees who can fly Airbus planes



Introduction to Database Systems

Example

Find the employee numbers of employees who can fly Boeing **or** MD planes


eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

Introduction to Database Systems

Example


Find the employee numbers of employees who can fly Boeing **or** MD planes



Introduction to Database Systems

Example

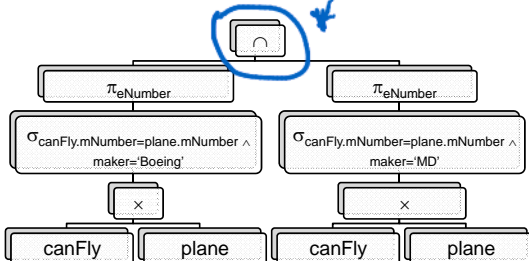
Find the employee numbers of employees who can fly Boeing **and** MD planes



Introduction to Database Systems

Example

Find the employee numbers of employees who can fly Boeing **and** MD planes



Introduction to Database Systems

Example

Find the names of employees who can fly two planes or more

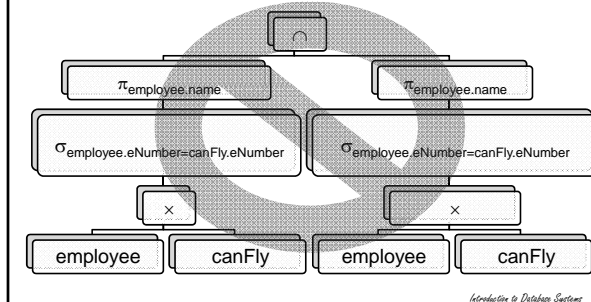
name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

Introduction to Database Systems

Example

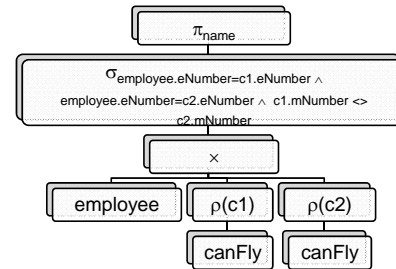
Find the names of employees who can fly two planes or more



NOT THE ANSWER!!

Example

Find the names of employees who can fly two planes or more



Example (SQL)

Find the names of employees who can fly two planes or more

```
SELECT employee.name
FROM canFly c1, canFly c2, employee
WHERE c1.eNumber=employee.eNumber
AND c2.eNumber=employee.eNumber
AND c1.mNumber <> c2.mNumber
```

cannot use aggregates

Example (SQL with aggregates)

Find the names of employees who can fly two planes or more (no aggregates in Algebra)

```
SELECT employee.name
FROM canFly, employee
WHERE employee.eNumber=canFly.eNumber
GROUP BY employee.eNumber, employee.name
HAVING COUNT(*) >= 2
```

Division

- Compute all possible combinations of the first column of R_1 and R_2 .

$$(\pi_A(R_1) \times R_2)$$

- Then remove those rows that exist in R_1

$$(\pi_A(R_1) \times R_2) - R_1$$

- Keep only the first column of the result. These are the disqualified values

$$\pi_A((\pi_A(R_1) \times R_2) - R_1)$$

- A/B is the first column of R_1 except the disqualified values

$$R_1 / R_2 = \pi_A(R_1) - \pi_A((\pi_A(R_1) \times R_2) - R_1)$$

Example

Find the employment numbers of employees who can fly all MD planes

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

Instead of NOT EXIST NOT EXIST
can use NOT EXIST EXCEPT

Division (Example)

Find the Employment numbers of the pilots who can fly **all** MD planes

$\text{canFly} / \pi_{\text{mNumber}}(\sigma_{\text{maker}='MD'}(\text{plane}))$

$$\pi_{\text{eNumber}}(\text{canFly}) - \pi_{\text{eNumber}}((\pi_{\text{eNumber}}(\text{canFly}) \times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) - \text{canFly})$$

Introduction to Database Systems

Example: step-by-step

$$\pi_{\text{eNumber}}(\text{canFly}) \times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))$$

eNumber	mNumber
1001	DC9
1001	DC10
1002	DC9
1002	DC10
1003	DC9
1003	DC10

Introduction to Database Systems

Example: step-by-step

$$(\pi_{\text{eNumber}}(\text{canFly}) \times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) - \text{canFly}$$

eNumber	mNumber
1001	DC9
1002	DC10

Introduction to Database Systems

Example: step-by-step

$$\pi_{\text{eNumber}}((\pi_{\text{eNumber}}(\text{canFly}) \times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) - \text{canFly})$$

eNumber
1001
1002

Introduction to Database Systems

Example: step-by-step

$$\pi_{\text{eNumber}}(\text{canFly}) - \pi_{\text{eNumber}}((\pi_{\text{eNumber}}(\text{canFly}) \times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) - \text{canFly})$$

eNumber
1003

Introduction to Database Systems

Division (SQL)

```
SELECT DISTINCT c1.eNumber
FROM canFly c1
WHERE NOT EXISTS
    (SELECT c.eNumber, p.mNumber
     FROM canFly c, plane p
     WHERE p.maker='MD'
           AND c1.eNumber=c.eNumber
     EXCEPT
     SELECT enumber, mNumber
     FROM canFly)
```

Introduction to Database Systems

Division (SQL)

```
SELECT DISTINCT c1.eNumber
FROM canFly c1
WHERE NOT EXISTS(
  SELECT *
  FROM plane p
  WHERE p.maker='MD' AND NOT EXISTS(
    SELECT *
    FROM canFly c2
    WHERE c1.eNumber=c2.eNumber
      AND p.mNumber=c2.mNumber))
```

Introduction to Database Systems

Credits

The content of this lecture is based
on chapter 4 of the book
"Introduction to database
Systems"

By
S. Bressan and B. Catania,
McGraw Hill publisher

Animated characters are animated
using VocaliseTTS under
license from Digital Curiosity

Clipart and media are licensed from
Microsoft Office Online Clipart
and Media

Copyright © 2013 by Stéphane Bressan



Introduction to Database Systems