

# *CS2102 Database Systems*

## *The Relational Model*



School of  
Computing

Leading The World With Asia's Best

# *Relational Model*

- ❖ Introduced by Edgar Codd from IBM Research Lab in 1970
- ❖ Data is modeled in terms of relations, tables with rows and columns
- ❖ Degree/Arity = No. of columns
- ❖ Cardinality = No. of rows

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

# *Relational Model Basics*

- ❖ A relational database is a set of relations
- ❖ Each relation has a definition called a relation schema
  - Schema specifies name of relation, attributes and data constraints.
  - Data constraints include domain constraints  
Students (sid: **string**, sname: **string**, login: **string**, age: **integer**, cap: **real**)
- ❖ Each row in a relation is called a tuple/record (all rows are distinct)

# *Relational Model Basics*

- ❖ A domain is a set of atomic values
  - Special value null is a member of each domain
  - A null value means value is not applicable or unknown
- ❖ A relation is a set of tuples
  - Consider a relation schema  $R(A_1, A_2, \dots, A_n)$  with  $n$  attributes
  - Let  $D_i$  be the domain of attribute  $A_i$  (a set of possible values for  $A_i$ )
  - Each instance of  $R$  is a subset of  $\{(a_1, a_2, \dots, a_n) \mid a_i \in D_i\}$
- ❖ A relational database schema is a set of relation schemas

# *Integrity Constraints (ICs)*

- ❖ Integrity constraint is a condition that must be true for any instance of the database;
  - ICs are specified when schema is defined.
  - ICs are checked when relations are updated.
- ❖ A legal instance of a relation is one that satisfies all specified ICs.
- ❖ DBMS enforces ICs - allows only legal instances to be stored.

# *Integrity Constraints (ICs)*

- ❖ ICs are based on the semantics of the real-world enterprise that is being described in the database relations.
- ❖ We can check a database instance to see if an IC is violated, but we can **NEVER** infer that an IC is true by looking at an instance.
  - An IC is a statement about *all possible* instances!
- ❖ Common types of ICs
  - Domain constraints - restricts attribute values of relations
  - Key constraints, foreign key constraints

# ICs: Key Constraints

- ❖ A superkey is a subset of attributes in a relation that uniquely identifies its tuples
- ❖ No two distinct tuples in a relation have same values in all attributes of superkey
- ❖ What are the possible superkeys for the following relation?

<i>sid</i>	<i>cid</i>	<i>grade</i>
1	CS204	C
1	CS101	A
2	CS204	C
3	CS101	B
3	MA112	B

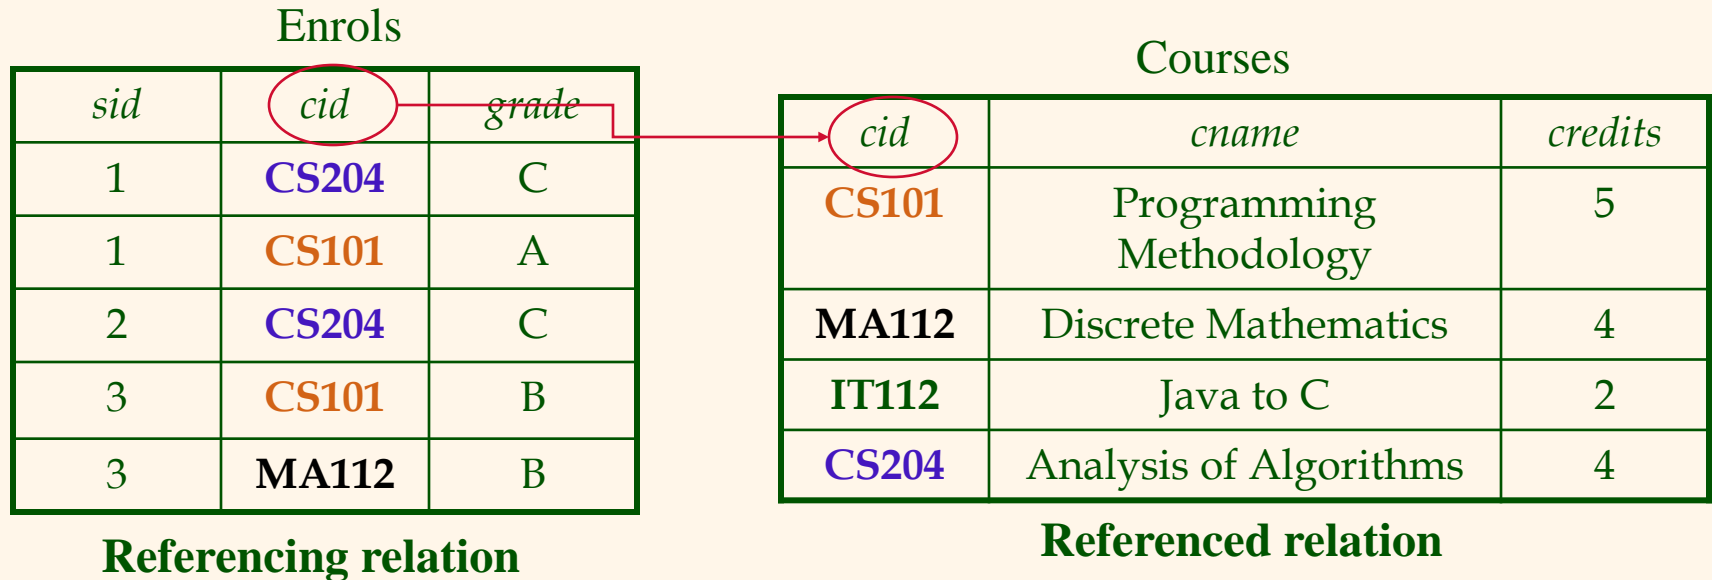
# *ICs: Key Constraints*

- ❖ A key is a superkey that satisfies the additional property: No proper subset of the key is a superkey
- ❖ That is, a key is a minimal subset of attributes in a relation that uniquely identifies its tuples
- ❖ Key attribute values cannot be null
- ❖ A relation could have multiple keys called candidate keys
- ❖ One of the candidate keys is chosen as the primary key



# ICs: Foreign Keys Constraints

- ❖ A subset of attributes in a relation is a foreign key if it refers to the primary key of a second relation



- ❖ *cid* is a foreign key in Enrols and a primary key in Courses

# ICs: Foreign Keys Constraints

- ❖ Foreign key constraint (referential integrity constraint): Each foreign key value in referencing relation must either
  - Appear as primary key value in the referenced relation, or
  - Be a null value
- ❖ Can occur within same relation
- ❖ Constraints on Course\_Prereq
  - cid is the primary key
  - Prereq is the foreign key

Course\_Prereq

<i>cid</i>	<i>prereq</i>
CS101	null
CS204	CS101
CS332	CS101
MA101	null
MA211	MA101

# Primary and Foreign Keys Constraints

- ❖ Which of the following updates violates the primary/foreign key constraints?
- Insert tuple (CS101, “Java Programming”, 4) into Courses
  - Delete tuple (MA112, “Discrete Mathematics”, 4) from Courses
  - Delete tuple (IT112, “Java to C”, 2) from Courses
  - Modify tuple (3, MA112, B) to (3, MA113, B) in Enrols
  - Insert tuple (1, null, C) into Enrols

Enrols			Courses		
<i>sid</i>	<i>cid</i>	<i>grade</i>	<i>cid</i>	<i>cname</i>	<i>credits</i>
1	CS204	C	CS101	Programming Methodology	5
1	CS101	A	MA112	Discrete Mathematics	4
2	CS204	C	IT112	Java to C	2
3	CS101	B	CS204	Analysis of Algorithms	4
3	MA112	B			

# *How to Handle Foreign Key Violation?*

- ❖ Which can be done if a tuple  $t$  in Courses is to be deleted?
- ❖ Options:
  - Disallow deletion if some row in Enrols refers to  $t$
  - Delete all rows in Enrols that refer to  $t$
  - For each row in Enrols that refers to  $t$ , replace its *cid* value with an existing default value
  - For each row in Enrols that refers to  $t$ , replace its *cid* value with the null value, provided *cid* is not a primary key attribute in Enrols
- ❖ Above options also apply if the primary key of a tuple in Courses is to be modified

# *Relational Query Languages*

- ❖ A major strength of the relational model: it supports simple, powerful *querying* of data.
- ❖ Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.

# *Structured Query Language (SQL)*

- ❖ Most widely used commercial relational database language
- ❖ Developed by IBM San Jose Research Laboratory in the 1970s
- ❖ Standards:
  - SQL-86
  - SQL-89 (minor revision)
  - SQL-92 (major revision)
  - SQL-99 (major extensions, current standard)

# *SQL: Table Definition*

- ❖ **CREATE TABLE** statement

- Specify attributes and constraints of table

- ❖ Basic attribute types

- CHAR(n), VARCHAR(n), BOOLEAN, INTEGER, REAL, DATE, TIME

- ❖ Basic constraints

- Primary, candidate, foreign key constraints

# SQL: Table Definition

```
CREATE TABLE Courses (  
    cid          CHAR(5)          DEFAULT `C0000`,  
    cname       VARCHAR(50),  
    credits     INTEGER,  
    prereq      CHAR(5),  
    PRIMARY KEY (cid),  
    UNIQUE (cname),  
    FOREIGN KEY (prereq) REFERENCES Courses(cid)  
)
```

- Type of each field is specified, and enforced by the DBMS when tuples are added or modified.
- Candidate keys are *cid* and *cname* (specified using **UNIQUE**)
- Each course *cid* has at most one prerequisite course



# *SQL: Table Definition*

```
CREATE TABLE Students (  
    sid          CHAR(5),  
    sname        VARCHAR(20),  
    age          INTEGER,  
    PRIMARY KEY (sid)  
)
```

```
CREATE TABLE Enrols (  
    sid CHAR(5), cid CHAR(5), grade CHAR,  
    PRIMARY KEY (sid, cid),  
    FOREIGN KEY (sid) REFERENCES Students(sid),  
    FOREIGN KEY (cid) REFERENCES Courses(cid)  
)
```

# *Enforcing Referential Integrity*

- ❖ *sid* in Enrols is a foreign key that references Students.
- ❖ What should be done if an Enrols tuple with a non-existent student id is inserted? (*Reject it!*)
- ❖ What should be done if a Students tuple is deleted?
  - Also delete all Enrols tuples that refer to it
  - Disallow deletion of a Students tuple that is referred to
  - Set *sid* in Enrols tuples that refer to it to a *default sid*
  - Set *sid* in Enrols tuples that refer to it to *null*
- ❖ Similar if primary key of Students tuple is updated.

# *SQL: Options for Foreign Key Violations*

```
CREATE TABLE Enrols (  
    sid          CHAR(5),  
    cid          CHAR(5),  
    grade        REAL,  
    PRIMARY KEY (sid, cid),  
    FOREIGN KEY (sid) REFERENCES Students(sid),  
    FOREIGN KEY (cid) REFERENCES Courses(cid)  
        ON DELETE/UPDATE option  
)
```

## Options:

- **NO ACTION** reject any violations (default)
- **CASCADE** propagate delete/update to referencing tuples
- **SET DEFAULT** update foreign keys of referencing tuples to some default values
- **SET NULL** update foreign keys of referencing tuples to null values

# *SQL: Options for Foreign Key Violations*

```
CREATE TABLE Students (  
    sid          CHAR(5),  
    sname        VARCHAR(20),  
    age          INTEGER,  
    PRIMARY KEY (sid),  
)  
  
CREATE TABLE Courses (  
    cid          CHAR(5)  DEFAULT `C0000`,  
    cname        VARCHAR(50),  
    credits      INTEGER,  
    PRIMARY KEY (cid),  
)  
  
CREATE TABLE Enrols (  
    sid CHAR(5), cid CHAR(5), grade REAL,  
    PRIMARY KEY (sid, cid),  
    FOREIGN KEY (sid) REFERENCES Students(sid)  
        ON DELETE CASCADE  
        ON UPDATE NO ACTION,  
    FOREIGN KEY (cid) REFERENCES Courses(cid)  
        ON UPDATE CASCADE  
        ON DELETE SET DEFAULT  
)
```

# *SQL: Query Single Relation*

- ❖ Find all 18 year old students

```
SELECT *  
FROM Students S  
WHERE S.age = 18
```

Students

sid	name	login	age	cap
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❖ Find student names and logins

```
SELECT S.name, S.login  
FROM Students S
```

# SQL: Query Multiple Relations

Students

sid	name	login	age	cap
53831	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Enrols

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53688	History105	B

```
SELECT S.name, E.cid  
FROM   Students S, Enrols E  
WHERE  S.sid = E.sid AND E.grade = "A"
```

S.name	E.cid
Smith	Topology112

# *SQL: Destroy Table*

**DROP TABLE** Students

- ❖ Destroys the relation Students.
- ❖ The schema information and the tuples are deleted.

# *SQL: Alter Table*

```
ALTER TABLE Students  
  ADD COLUMN firstYear: integer
```

- ❖ Schema of Students is altered by adding a new attribute
- ❖ Every tuple in the current instance of the relation is extended with a *null* value in the new field.



# *SQL: Adding and Deleting Tuples*

- ❖ Insert a single tuple

```
INSERT INTO Students (sid, name, login, age, gpa)  
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

- ❖ Delete all tuples satisfying some condition  
(e.g., name = Smith)

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```


# *Summary*

- ❖ Relational model is a tabular representation of data.
- ❖ Integrity constraints is based on application semantics.
  - Domain constraints
  - Primary/candidate key constraints
  - Foreign key constraints
- ❖ Options for dealing with foreign key constraint violations
  - Disallow violations
  - Update/delete referencing tuples

# *From ER to RELATIONAL*

# *Map Entity Types*


- ❖ Map an entity type to a relation
- ❖ Entity type attributes become attributes of the relation
- ❖ Key of entity type becomes key of relation



Students (matric, name, age)

```
CREATE TABLE Students (  
    matric VARCHAR(5),  
    name   VARCHAR(20),  
    age    INTEGER,  
    PRIMARY KEY (matric)  
)
```

# *Map Entity Types*




```
CREATE TABLE Person (  
    firstName VARCHAR(32),  
    lastName VARCHAR(32),  
    address VARCHAR(128),  
    PRIMARY KEY (firstName, lastName)  
)
```

# *Map Relationship Sets*

- ❖ Map a relationship set to a relation
- ❖ Attributes of relationship set become attributes of relation
- ❖ Identifier of each participating entity type form a superkey for the relation
  - They are also foreign keys


# *Many-to-Many Relationship*



**Enrol (matric, code, grade)**

```
CREATE TABLE Enrol (  
    matric CHAR(5), code CHAR(5), grade CHAR,  
    PRIMARY KEY (matric, code),  
    FOREIGN KEY (matric) REFERENCES Student(matric),  
    FOREIGN KEY (code) REFERENCES Course(code)  
)
```


# *Many-to-One Relationship*



- ❖ Each student is supervised by at most one professor
- ❖ Map relationship type Supervise to a relation
- ❖ Key is sid
- ❖ Separate tables for Professor and Student.
- ❖ Since each student has a unique professor, we could also combine Student and Supervise.



# Many-to-One Relationship



Professor (pid, pname, expertise)

Student (sid, name, age)


Supervise (sid, pid, since)

OR

Professor (pid, pname, expertise)

SupervisedStudent (sid, sname, age, pid, since)

# Many-to-One Relationship




OR

```
CREATE TABLE Supervise (  
  sid CHAR(5),  
  pid CHAR(8),  
  since DATE,  
  PRIMARY KEY (sid),  
  FOREIGN KEY (sid)  
    REFERENCES Student(sid),  
  FOREIGN KEY (pid)  
    REFERENCES Professor(pid)  
)
```

```
CREATE TABLE SupervisedStudent (  
  sid CHAR(5), sname CHAR(20),  
  age INTEGER,  
  pid CHAR(8),  
  since DATE,  
  PRIMARY KEY (sid),  
  FOREIGN KEY (pid)  
    REFERENCES Professor(pid)  
)
```

# Map Participation Constraints

- ❖ Does every department have a manager?
- ❖ If so, the participation of Dept in Manage is *mandatory or total* (vs. *optional or partial*)
- ❖ Every *dname* value in Dept table must appear in a row of the Manage table (with a non-null *eid* value)




# *Map Participation Constraints*

- ❖ Capture participation constraints involving one entity type in a binary relationship


```
CREATE TABLE DeptManager (  
    dname          CHAR(20),  
    budget         REAL,  
    eid            CHAR(5) NOT NULL,  
    PRIMARY KEY (dname),  
    FOREIGN KEY (eid) REFERENCES Employee,  
    ON DELETE NO ACTION  
)
```

# Exercise

What are the relations obtained from this ER diagram?



# Map Recursive Relationship



Mentor (senior\_sid, junior\_sid)

```
CREATE TABLE Mentor (  
    senior_sid CHAR(5), junior_sid CHAR(5),  
    PRIMARY KEY (senior_sid, junior_sid),  
    FOREIGN KEY (senior_sid) REFERENCES Student(sid),  
    FOREIGN KEY (junior_sid) REFERENCES Student(sid)  
)
```

# *Map Weak Entity Types*

- ❖ Owner entity type and weak entity type must participate in a one-to-many relationship set (1 owner, many weak entities).
- ❖ Weak entity type must have total participation in this identifying relationship set.
- ❖ Weak entity type and identifying relationship set are mapped to a single relation
  - When the owner entity is deleted, all owned weak entities must also be deleted.


# Map Weak Entity Types

Book (isbn, author, title)

Chapter (isbn, cno, ctitle)

```
CREATE TABLE Book (  
  isbn CHAR(10), title CHAR(50),  
  author CHAR(50),  
  PRIMARY KEY (isbn) )
```

```
CREATE TABLE Chapter (  
  isbn CHAR(10), cno INTEGER, ctitle CHAR(50),  
  PRIMARY KEY (isbn, cno),  
  FOREIGN KEY (isbn) REFERENCES Book(isbn),  
  ON DELETE CASCADE) )
```






# Map ISA Hierarchies

- ❖ One relation per subclass or superclass

Employee (eid, ename)

Contract\_Emp (eid, contractid)

Temp (eid, hrs\_worked,  
hourly\_rate)



# *Map ISA Hierarchies*

- ❖ One relation per subclass or superclass

```
CREATE TABLE Employee (  
    eid      CHAR(5) PRIMARY KEY,  
    ename    CHAR(30),  
)
```

```
CREATE TABLE Contract_Emp (  
    eid      CHAR(5) PRIMARY KEY  
            REFERENCES Employee,  
    contractid CHAR(10)  
)
```

```
CREATE TABLE Temp (  
    eid      CHAR(5) PRIMARY KEY  
            REFERENCES Employee,  
    hrs_worked    INTEGER,  
    hourly_rate   REAL  
)
```

# *Map ISA Hierarchies*

- ❖ One relation per subclass if each employee must be in one of these two subclasses.

Contract\_Emp (eid, ename, contractid)

Temp (eid, ename, hrs\_worked, hourly\_rate)

```
CREATE TABLE ContractEmp (  
    eid    CHAR(5) PRIMARY KEY,  
    ename  CHAR(30),  
    contractid CHAR(10)  
)
```

```
CREATE TABLE Temp (  
    eid    CHAR(5) PRIMARY KEY,  
    ename  CHAR(30),  
    hrs_worked    INTEGER,  
    hourly_rate   REAL  
)
```

# *Summary*

- ❖ Rules for mapping ER to relational model
  - Entity types and relationship sets
  - Relationship cardinalities
  - Participation constraints
  - Recursive relationship
  - Weak entity types
  - ISA hierarchies