

**CS2102**  
**Structured Query Language (SQL)**  
**Part 1**

# Structured Query Language

- Designed by D. Chamberlin & R. Boyce (IBM Research) in 1974
  - Original name: SEQUEL (Structured English QUERy Language)
- SQL is not a general-purpose language but a **domain-specific language (DSL)**
  - Designed for computations on relations
- Unlike relational algebra which is a procedural language, SQL is a **declarative language**
  - Focuses on *what* to compute, not on *how* to compute

# Using SQL

- **Directly write SQL statements**
  - Command line interface
    - PostgreSQL's **psql**  
<https://www.postgresql.org/docs/current/static/app-psql.html>
  - Graphical user interface
    - PostgreSQL's **pgAdmin**  
<https://www.pgadmin.org/>
- **Include SQL in application programs**
  - Statement-Level Interface (SLI)
    - Embedded SQL
    - Dynamic SQL
  - Call-Level Interface (CLI)
    - JDBC (Java DataBase Connectivity)
    - ODBC (Open DataBase Connectivity)

# SQL

- Current ANSI/ISO standard for SQL is called SQL:2016
  - Different DBMSs may have minor variations in SQL syntax
- SQL consists of two main parts: DDL & DML
- **Data Definition Language (DDL):**  
create/delete/modify schemas
- **Data Manipulation Language (DML):** ask queries,  
insert/delete/modify data

# Create/Drop Table

-- Comments in SQL are preceded  
-- by two hyphens

```
create table Students (  
    studentId      integer,  
    name           varchar(100),  
    birthDate      date  
);
```

```
/* SQL also supports  
C-style comments */  
drop table Students;
```

# Data Types

- Examples of built-in data types:
  - boolean
  - integer, numeric, real
  - char(50), varchar(50), text
  - date, time, timestamp
- SQL also supports user-defined data types
- The domain of each data type includes the special value **null**

# Null Values

- SQL uses a three-valued logic system: *true*, *false*, & *unknown*

x	y	x AND y	x OR y	NOT x
FALSE	FALSE	FALSE	FALSE	TRUE
FALSE	UNKNOWN	FALSE	UNKNOWN	
FALSE	TRUE	FALSE	TRUE	
UNKNOWN	FALSE	FALSE	UNKNOWN	UNKNOWN
UNKNOWN	UNKNOWN	UNKNOWN	UNKNOWN	
UNKNOWN	TRUE	UNKNOWN	TRUE	
TRUE	FALSE	FALSE	TRUE	FALSE
TRUE	UNKNOWN	UNKNOWN	TRUE	
TRUE	TRUE	TRUE	TRUE	

# Null Values (cont.)

- Result of a **comparison operation** involving *null* value is *unknown*
- Result of an **arithmetic operation** involving *null* value is *null*
- Examples: Assume that the value of *x* is *null*
  - ***x* < 100** evaluates to *unknown*
  - ***x* = null** evaluates to *unknown*
  - ***x* <> null** evaluates to *unknown*
  - ***x* + 20** evaluates to *null*



# IS NULL Comparison Predicate

- How to check if a value is equal to *null*?
- Use the **IS NULL** comparison predicate
- Examples:
  - **x IS NULL** evaluates to *true* if *x* is a null value
  - **x IS NULL** evaluates to *false* if *x* is a non-null value
  - **x IS NOT NULL** evaluates to *false* if *x* is a null value
  - **x IS NOT NULL** evaluates to *true* if *x* is a non-null value

# IS DISTINCT FROM Comparison Predicate

- How to treat *null* values as ordinary values for comparison?
- Use the **IS DISTINCT FROM** comparison predicate
- The comparison “**x IS DISTINCT FROM y**”
  - is equivalent to “ $x \neq y$ ” if both *x* & *y* are non-null values
  - evaluates to *false* if both the values are *null*
  - evaluates to *true* if only one of the values is *null*

# Constraints in Data Definitions

- **Constraint Types**
  - Not-null constraints
  - Unique constraints
  - Primary key constraints
  - Foreign key constraints
  - General constraints
- **Constraint Specifications**
  - Column constraints
  - Table constraints
  - Assertions (not covered)
- A constraint is **violated** if it evaluates to *false*

# Not-Null Constraints

```
create table Students (  
    studentId      integer,  
    name           varchar(100) not null,  
    birthDate      date  
);
```

```
create table Students (  
    studentId      integer,  
    name           varchar(100),  
    birthDate      date,  
    check (name is not null)  
);
```

# Unique Constraints

```
create table Students (  
    studentId      integer unique,  
    name           varchar(100),  
    birthDate      date  
);
```

The unique constraint is violated if there exists two records  $x, y \in Students$  where “ $x.studentId \neq y.studentId$ ” evaluates to *false*

studentId	name	birthDate
100	Alice	1999-12-25
200	Bob	2000-04-01
200	Carol	2001-11-28

studentId	name	birthDate
100	Alice	1999-12-25
null	Bob	2000-04-01
null	Carol	2001-11-28

# Unique Constraints (cont.)

```
create table Census (  
  city      varchar(50),  
  state     char(2),  
  population integer,  
  unique (city,state)  
);
```

city	state	population
New York	NY	8537673
null	CA	3976322
Chicago	IL	2704958
Houston	TX	2303482
null	CA	1406630

The unique constraint is violated if there exists two records  $x, y \in \text{Census}$  where  
“(x.city  $\neq$  y.city) or (x.state  $\neq$  y.state)”  
evaluates to *false*

# Primary Key Constraints

```
create table Students (  
    studentId      integer primary key,  
    name           varchar(100),  
    birthDate      date  
);
```

```
create table Students (  
    studentId      integer unique not null,  
    name           varchar(100),  
    birthDate      date  
);
```

# Primary Key Constraints (cont.)

```
create table Enrolls (  
    sid      integer,  
    cid      integer,  
    grade    char(2),  
    primary key (sid, cid)  
);
```



# Foreign Key Constraints

```
create table Students (  
    studentId      integer,  
    name           varchar(100),  
    birthDate      date,  
    primary key (studentId));
```

```
create table Courses (  
    courseId       integer,  
    name           varchar(80),  
    credits         integer,  
    primary key (courseId));
```

```
create table Enrolls (  
    sid            integer references Students (studentId),  
    cid            integer,  
    grade          char(2),  
    primary key (sid, cid),  
    foreign key (cid) references Courses (courseId));
```

# Foreign Key Constraints (cont.)



- Consider a **foreign key constraint** with the subset of attributes  $A_f$  in referencing table  $T_f$  referring to the subset of attributes  $A_p$  in referenced table  $T_p$
- $A_p$  must be declared as **primary key** or **unique** in table  $T_p$
- For each tuple  $t_f \in T_f$  with non-null values for all the attributes  $A_f$ , the referenced table  $T_p$  must contain a tuple  $t_p$  where the values of  $t_p$ 's attributes  $A_p$  agree with those of  $t_f$ 's attributes  $A_f$

# General Constraints: Example

```
create table Movies (  
    title          varchar(100),  
    director       varchar(80),  
    releaseYear integer,  
    rating         numeric (3,1),  
    primary key (title),  
    check (releaseYear > 2010 or rating > 8.0 )  
);
```

# Database Modifications: Insert

```
create table Students (  
    studentId      integer primary key,  
    name           varchar(100) not null,  
    birthDate      date,  
    dept           varchar(20) default 'CS'  
);
```

```
insert into Students  
values      (12345, 'Alice', '1999-12-25', 'Maths');
```

```
insert into Students (name, studentId)  
values      ('Bob', 67890);
```

studentId	name	birthDate	dept
12345	Alice	1999-12-25	Maths
67890	Bob	null	CS

# Database Modifications: Delete

```
create table Students (  
    studentId      integer primary key,  
    name           varchar(100) not null,  
    birthDate      date,  
    dept           varchar(20) default 'CS'  
);
```

```
-- Remove all students  
delete from Students;
```

```
-- Remove all students from Maths department  
delete from Students  
where      dept = 'Maths';
```

# Database Modifications: Update

```
create table Accounts (  
    accountId      integer primary key,  
    name           varchar(100) not null,  
    birthDate      date,  
    balance        numeric (10,2) default 0.00  
);
```

-- Add 2% interest to all accounts

```
update Accounts  
set      balance = balance * 1.02;
```

-- Add \$500 to account 12345

```
update   Accounts  
set      balance = balance + 500  
where    accountId = 12345;
```

# Transactions

- A **transaction** consists of one or more update/retrieval operations (i.e., SQL statements)
- Abstraction for representing a logic unit of work
- The **begin** command starts a new transaction
- Each transaction must end with either a **commit** or **rollback** command

```
begin;  
update Accounts  
set balance = balance + 1000  
where accountId = 456;  
  
update Accounts  
set balance = balance - 1000  
where accountId = 123;  
commit;
```

# Transactions (cont.)

```
begin;  
update Accounts  
set balance = balance + 1000  
where accountId = 456;  
  
update Accounts  
set balance = balance - 1000  
where accountId = 123;  
commit;
```

vs.

```
update Accounts  
set balance = balance + 1000  
where accountId = 456;  
  
update Accounts  
set balance = balance - 1000  
where accountId = 123;
```




# Transactions: ACID Properties

- **Atomicity**: Either all the effects of the transactions are reflected in the database or none are
- **Consistency**: The execution of a transaction in isolation preserves the consistency of the database
- **Isolation**: The execution of a transaction is isolated from the effects of other concurrent transaction executions
- **Durability**: The effects of a committed transaction persists in the database even in the presence of system failures

# Checking of Constraints

- By default, constraints are checked **immediately** at the end of SQL statement execution
  - A violation will cause the statement to be rolledback
- Constraint checking could also be **deferred** to the end of transaction execution
  - A violation will cause the transaction to be aborted
- The type of constraint checking could be specified as part of constraint declaration or configure using **set constraints** command

# Handling foreign key constraint violations



- Deletion/update of a referenced tuple could violate a foreign key constraint
- Specify action to deal with violation as part of a foreign key constraint declaration:

FOREIGN KEY ... REFERENCES ... ON DELETE/UPDATE action

# Handling foreign key constraint violations (cont.)

- **NO ACTION**: rejects delete/update if it violates constraint (default option)
- **RESTRICT**: similar to *NO ACTION* except that constraint checking can't be deferred
- **CASCADE**: propagates delete/update to referencing tuples
- **SET DEFAULT**: updates foreign keys of referencing tuples to some default value
- **SET NULL**: updates foreign keys of referencing tuples to *null* value

# Handling foreign key constraint violations (cont.)

```
create table Students (  
  studentId      integer,  
  name           varchar(100),  
  birthDate      date,  
  primary key (studentId));
```

```
create table Courses (  
  courseId       integer,  
  name           varchar(80),  
  credits         integer,  
  primary key (courseId));
```

```
create table Enrolls (  
  sid integer, cid integer, grade char(2),  
  primary key (sid, cid),  
  foreign key (sid) references Students  
    on delete cascade  
    on update no action,  
  foreign key (cid) references Courses  
    on update cascade  
    on delete set null);
```

# Modifying Schema with ALTER TABLE Command

- Add/remove/modify columns
  - **alter table** Students **alter column** dept **drop default**;
  - **alter table** Students **drop column** dept;
  - **alter table** Students **add column** faculty **varchar**(20);
  - etc.
- Add/remove constraints
- etc.

# Simple Queries

- Basic form of SQL query consists of three clauses:

**select**      [ **distinct** ] select-list  
**from**        from-list  
[ **where**    qualification ]

- **from-list**: Specifies list of relations
- **qualification**: Specifies conditions on relations
- **select-list**: Specifies columns to be included in output table
- Output relation could contain duplicate records if **distinct** is not used in the select clause

# Simple Queries (cont.)

**select distinct**     $a_1, a_2, \dots, a_m$   
**from**                     $r_1, r_2, \dots, r_n$   
**where**                     $c$

$$\pi_{a_1, a_2, \dots, a_m} \left( \sigma_c \left( r_1 \times r_2 \times \dots \times r_n \right) \right)$$



# Simple Queries (cont.)

Find the names of restaurants, the pizzas that they sell and their prices, where the price is under \$20

```
select rname, pizza, price
from Sells
where price < 20;
```

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	pizza	price
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Pizza King	Diavola	17

# Simple Queries (cont.)

Find the names of restaurants, the pizzas that they sell and their prices, where the price is under \$20

```
select rname, pizza, price  
from Sells  
where price < 20;
```

```
select *  
from Sells  
where price < 20;
```

# Simple Queries (cont.)

Find all tuples from `Sells` relation such that (1) the price is under \$20 and the restaurant name is not “Pizza King” or (2) the restaurant name is “Corleone Corner”.

```
select * from Sells
where ((price < 20) and (rname <> 'Pizza King'))
or      rname = 'Corleone Corner';
```

`Sells`

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16

# Removing Duplicate Records

Q1: **select** A, C **from** R;

Q2: **select** **distinct** A, C **from** R;

R			Q1		Q2	
A	B	C	A	C	A	C
10	1	2	10	2	10	2
10	7	2	10	2	20	null
20	3	null	20	null	30	2
20	9	null	20	null	30	9
30	3	2	30	2		
30	5	9	30	9		

Two tuples  $(a_1, c_1)$  and  $(a_2, c_2)$  in  $R$  are considered to be distinct if the following evaluates to *true*:

“( $a_1$  IS DISTINCT FROM  $a_2$ ) or ( $c_1$  IS DISTINCT FROM  $c_2$ )”

# Expressions in SELECT Clause

**select** item, price \* qty **as** cost  
**from** Orders;

Orders

item	price	qty
A	2.50	100
B	4.00	100
C	7.50	100

item	cost
A	250.00
B	400.00
C	750.00

# Expressions in SELECT Clause (cont.)

**select** 'Price of ' || pizza || ' is ' || **round**(price / 1.3) || ' USD' **as** menu  
**from** Sells  
**where** rname = 'Corleone Corner';

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

menu
Price of Diavola is 11 USD
Price of Hawaiian is 20 USD
Price of Margherita is 15 USD

# Pattern Matching with LIKE Operator

Find customer names ending with “e” that consists of at least four characters

**select** cname **from** Customers **where** cname **like** '\_\_\_\_%e';

Customers	
cname	area
Homer	West
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Willie	North

cname
Maggie
Willie

- The underscore symbol `_` matches any single character
- The percent symbol `%` matches any sequence of 0 or more characters

# Set Operations

- Let  $Q_1$  &  $Q_2$  denote SQL queries that output union-compatible relations
- $Q_1$  **union**  $Q_2 = Q_1 \cup Q_2$
- $Q_1$  **intersect**  $Q_2 = Q_1 \cap Q_2$
- $Q_1$  **except**  $Q_2 = Q_1 - Q_2$
- **union**, **intersect**, and **except** eliminate duplicate records
- **union all**, **intersect all**, and **except all** preserves duplicate records



# Set Operations (cont.)

Example 1: Find all customer/restaurant names

**select** cname **from** Customers

**union**

**select** rname **from** Restaurants;

Example 2: Find pizzas that contain both cheese and chilli

**select** pizza **from** Contains **where** ingredient = 'cheese'

**intersect**

**select** pizza **from** Contains **where** ingredient = 'chilli';

Example 3: Find pizzas that contain cheese but not chilli

**select** pizza **from** Contains **where** ingredient = 'cheese'

**except**

**select** pizza **from** Contains **where** ingredient = 'chilli';

# Set Operations (cont.)

Q1: select B from R **except** select B from S;

Q2: select B from R **except all** select B from S;

R		S		Q1	Q2
A	B	A	B	B	B
10	1	10	1	4	1
20	1	20	5		1
30	1	30	2		4
40	2	40	2		4
50	3	50	3		
60	4				
70	4				

# Multi-relation Queries

Find customers and restaurants that are located in the same area

```
select  cname, rname
from    Customers, Restaurants
where    Customers.area = Restaurants.area;
```

```
select  cname, rname
from    Customers as C, Restaurants as R
where    C.area = R.area;
```

# Multi-relation Queries (cont.)

Find distinct restaurant pairs (R1,R2) where  $R1 < R2$   
and they sell some common pizza

```
select  distinct S1.rname, S2.rname
from    Sells S1, Sells S2
where   S1.rname < S2.rname
and     S1.pizza = S2.pizza;
```

# Join Operators in Relational Algebra

- A **join operator** combines cross-product, selection, and possibly projection operators
- More convenient to use than plain cross-product operator
- Join operators:
  - Natural join,  $\bowtie$
  - Inner join (aka join, theta join, or condition join),  $\bowtie_c$
  - Left outer join (aka left join)  $\rightarrow_c$
  - Right outer join (aka right join)  $\leftarrow_c$
  - Full outer join (aka full join)  $\leftrightarrow_c$
  - Others: Natural left/right/outer joins

# Natural Join: $R \bowtie S$

The **natural join** of  $R$  and  $S$  is defined as

$$R \bowtie S = \pi_{\ell}(\sigma_c(R \times S))$$

where

$A$  = common attributes between  $R$  &  $S = \{a_1, a_2, \dots, a_n\}$ ,

$c$  is “ $(R.a_1 = S.a_1) \wedge (R.a_2 = S.a_2) \wedge \dots \wedge (R.a_n = S.a_n)$ ”,

$\ell$  is the list of attributes in  $A$ , followed by the list of attributes in  $R$  (excluding those in  $A$ ) and the list of attributes in  $S$  (excluding those in  $A$ )

# Natural Join: Example

For each restaurant, find its name, area, and the pizzas it sells together with their prices

Restaurants

rname	area
Corleone Corner	North
Gambino Oven	Central
Lorenzo Tavern	Central
Mamma's Place	South
Pizza King	East

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Restaurants ⋈ Sells

rname	area	pizza	price
Corleone Corner	North	Diavola	24
Corleone Corner	North	Hawaiian	25
Corleone Corner	North	Margherita	19
Gambino Oven	Central	Siciliana	16
Lorenzo Tavern	Central	Funghi	23
Mamma's Place	South	Marinara	22
Pizza King	East	Diavola	17
Pizza King	East	Hawaiian	19

$\text{Restaurants} \bowtie \text{Sells} = \pi_{\text{Restaurants.rname, area, pizza, price}}(\sigma_c(\text{Restaurants} \times \text{Sells}))$   
 where  $c$  is "Restaurant.rname = Sells.rname"

# Natural Join: Example (cont.)

For each restaurant, find its name, area, and the pizzas it sells together with their prices

```
select  R.rname, R.area, S.pizza, S.price
from    Restaurants R, Sells S
where   R.rname = S.rname;
```

```
select  *
from    Restaurants natural join Sells;
```



# Natural Join: Example (cont.)

Find distinct names of restaurants that sell some pizza for under \$20 that Homer likes

```
select  distinct S.rname
from    Sells S, Likes L
where   S.price < 20
and     L.cname = 'Homer'
and     S.pizza = L.pizza;
```

```
select  distinct S.rname
from    Sells S natural join Likes L
where   S.price < 20
and     L.cname = 'Homer';
```

# Inner Join: $R \bowtie_c S$

The **inner join** of  $R$  and  $S$  is defined as

$$R \bowtie_c S = \sigma_c(R \times S)$$

**Example:** Find customer pairs  $(C1, C2)$  such that they like some common pizza and  $C1 < C2$

Likes

cname	pizza
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Sciliana
Ralph	Diavola

$\pi_{cname, cname2}(Likes \bowtie_c \rho_{Likes2(cname2, pizza2)}(Likes))$   
where  $c = (pizza = pizza2) \wedge (cname < cname2)$

cname	cname2
Lisa	Maggie
Lisa	Moe
Maggie	Moe

# Inner Join: $R \bowtie_c S$ (cont.)

Find customer pairs ( $C1, C2$ ) such that they like some common pizza and  $C1 < C2$

```
select  distinct L1.cname, L2.cname
from    Likes L1, Likes L2
where   L1.cname < L2.cname
and     L1.pizza = L2.pizza;
```

```
select  distinct L1.cname, L2.cname
from    Likes L1 inner join Likes L2
on      (L1.pizza = L2.pizza) and (L1.cname < L2.cname);
```

# Left Outer Join: $R \rightarrow_c S$

- $attr(R)$  denote the list of attributes in the schema of  $R$ 
  - Example:  $attr(Sells) = \text{rname, pizza, price}$
- Let  $null(R)$  denote a  $n$ -component tuple of null values, where  $n = \text{arity of relation } R$ 
  - Example:  $null(Sells) = (\text{null}, \text{null}, \text{null})$

The **left outer join** of  $R$  and  $S$  is defined as

$$R \rightarrow_c S = (R \bowtie_c S) \cup ((R \triangleright_c S) \times \{null(S)\})$$

where

$$R \triangleright_c S = R - (R \bowtie_c S), \quad // \text{ left antijoin of } R \text{ \& } S$$

$$R \bowtie_c S = \pi_{attr(R)}(R \bowtie S) \quad // \text{ left semijoin of } R \text{ \& } S$$

# Left Outerjoin: Example

- Find customers and the pizzas they like; include also customers who don't like any pizza

Customers		Likes			
cname	area	cname	pizza	cname	pizza
Homer	West	Homer	Hawaiian	Homer	Hawaiian
Lisa	South	Homer	Margherita	Homer	Margherita
Maggie	East	Lisa	Funghi	Lisa	Funghi
Moe	Central	Maggie	Funghi	Maggie	Funghi
Ralph	Central	Moe	Funghi	Moe	Funghi
Willie	North	Moe	Sciliana	Moe	Sciliana
		Ralph	Diavola	Ralph	Diavola
				Willie	null

$\pi_{Customers.cname, pizza}(Customers \rightarrow_c Likes)$   
where c is “Customers.cname = Likes.cname”

# Left Outerjoin: Example (cont.)

Find customers and the pizzas they like; include also customers who don't like any pizza

```
select  C.cname, L.pizza
from    Customers C left outer join Likes L
on C.cname = L.cname;
```

```
select  C.cname, L.pizza
from    Customers C natural left outer join Likes L;
```

# Full Outer Join: $R \leftrightarrow_c S$

The **full outer join** of  $R$  and  $S$  is defined as

$$R \leftrightarrow_c S = (R \rightarrow_c S) \cup (\{null(R)\} \times (S \triangleright_c R))$$

# Full Outerjoin: Example

Find customer-restaurant pairs (C,R) where C and R are located in the same area. Include customers that are not co-located with any restaurant, and include restaurants that are not co-located with any customer

Customers

cname	area
Homer	North
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Willie	North

Restaurants

rname	area
Corleone Corner	West
Gambino Oven	East
Lorenzo Tavern	Central
Mamma's Place	South
Pizza King	East

cname	rname
Homer	null
Lisa	Mamma's Place
Maggie	Gambino Oven
Maggie	Pizza King
Moe	Lorenzo Tavern
Ralph	Lorenzo Tavern
Willie	null
null	Corleone Corner



# Full Outerjoin: Example (cont.)

Find customer-restaurant pairs (C,R) where C and R are located in the same area. Include customers that are not co-located with any restaurant, and include restaurants that are not co-located with any customer

```
select  C.cname, R.rname
from    Customers C full outer join Restaurants R
on      C.area = R.area;
```

```
select  C.cname, R.rname
from    Customers C natural full outer join Restaurants R;
```

# Summary

- SQL is the standard query language for relational DBMS
- Basic form for querying consists of SELECT, FROM, & WHERE clauses

**select distinct**     $a_1, a_2, \dots, a_m$   
**from**                     $r_1, r_2, \dots, r_n$   
**where**                     $C$

$$\pi_{a_1, a_2, \dots, a_m} \left( \sigma_C \left( r_1 \times r_2 \times \dots \times r_n \right) \right)$$