

CS2102 Lecture 11

Transactions

Transactions: Recap

- Abstraction for representing a logical unit of work
- **ACID Properties**
 - **Atomicity**: Either all the effects of the transactions are reflected in the database or none are
 - **Consistency**: The execution of a transaction in isolation preserves the consistency of the database
 - **Isolation**: The execution of a transaction is isolated from the effects of other concurrent transaction executions
 - **Durability**: The effects of a committed transaction persists in the database even in the presence of system failures

Transaction Example: Money Transfer

```
1 int Transfer (int fromAcctId, int toAcctId, int amount)
2 {
3     EXEC SQL BEGIN DECLARE SECTION;
4         int fromBalance;    int toBalance;
5     EXEC SQL END DECLARE SECTION;
6     EXEC SQL WHENEVER SQLERROR GOTO query_error;
7
8     EXEC SQL SELECT balance INTO :fromBalance FROM Accounts
9         WHERE accountId = :fromAcctId;
10    if (fromBalance < amount) {
11        EXEC SQL ROLLBACK;    return 1;
12    }
13    EXEC SQL SELECT balance INTO :toBalance FROM Accounts
14        WHERE accountId = :toAcctId;
15    EXEC SQL UPDATE Accounts SET balance = :toBalance + :amount
16        WHERE accountId = :toAcctId;
17    EXEC SQL UPDATE Accounts SET balance = :fromBalance - :amount
18        WHERE accountId = :fromAcctId;
19    EXEC SQL COMMIT;
20    return 0;
21    query_error: printf ("SQL error: %ld\n", sqlca->sqlcode); exit();
22 }
```

Transaction Example: Money Transfer

Two possible execution outcomes:

```
begin transaction;
```

```
select    balance into :fromBalance  
from      Accounts  
where     accountId = :fromAcctId;
```

```
rollback;
```

```
begin transaction;
```

```
select    balance into :fromBalance  
from      Accounts  
where     accountId = :fromAcctId;
```

```
select    balance into :toBalance  
from      Accounts  
where     accountId = :toAcctId;
```

```
update    Accounts  
set       balance = :toBalance + :amount  
where     accountId = 1;
```

```
update    Accounts  
set       balance = :fromBalance - :amount  
where     accountId = 2;
```

```
commit;
```

ANSI SQL Isolation Levels

- The **isolation level** for a transaction affects what the transaction will read
- ANSI SQL defines four isolation levels
 - Read Uncommitted (**weakest isolation level**)
 - Read Committed
 - Repeatable Read
 - Serializable (**strongest isolation level**)
- Choice of isolation level affects correctness vs performance tradeoff
- In many DBMSs, the default isolation level is **Read Committed**
- Configure using **set transaction isolation level** statement

Serial Transaction Executions

- Consider a set of transactions $S = \{T_1, \dots, T_n\}$
- An execution of S is a **serial execution** if the execution of the transactions in S are not interleaved
 - For any pair of transactions T_i and T_j in S , either T_i completes execution before T_j starts its execution, or T_j completes execution before T_i starts its execution

Serial Transaction Executions

Two possible serial executions of **Transfer(1,2,100)** & **Transfer(2,1,100)**

```
begin transaction;
select      balance into :xbal
from        Accounts
where       accountId = 1;
select      balance into :ybal
from        Accounts
where       accountId = 2;
update      Accounts
set         balance = :ybal + 100
where       accountId = 2;
update      Accounts
set         balance = :xbal - 100
where       accountId = 1;
commit;

begin transaction;
select      balance into :ybal2
from        Accounts
where       accountId = 2;
select      balance into :xbal2
from        Accounts
where       accountId = 1;
update      Accounts
set         balance = :xbal2 + 100
where       accountId = 1;
update      Accounts
set         balance = :ybal2 - 100
where       accountId = 2;
commit;
```

```
begin transaction;
select      balance into :ybal2
from        Accounts
where       accountId = 2;
select      balance into :xbal2
from        Accounts
where       accountId = 1;
update      Accounts
set         balance = :xbal2 + 100
where       accountId = 1;
update      Accounts
set         balance = :ybal2 - 100
where       accountId = 2;
commit;

begin transaction;
select      balance into :xbal
from        Accounts
where       accountId = 1;
select      balance into :ybal
from        Accounts
where       accountId = 2;
update      Accounts
set         balance = :ybal + 100
where       accountId = 2;
update      Accounts
set         balance = :xbal - 100
where       accountId = 1;
commit;
```

Interleaved Transaction Executions

An interleaved execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

```
begin transaction;  
select    balance into :xbal  
from      Accounts  
where     accountId = 1;  
select    balance into :ybal  
from      Accounts  
where     accountId = 2;  
update    Accounts  
set       balance = :ybal + 100  
where     accountId = 2;
```

```
update    Accounts  
set       balance = :xbal - 100  
where     accountId = 1;  
commit;
```

```
begin transaction;  
select    balance into :ybal2  
from      Accounts  
where     accountId = 2;  
select    balance into :xbal2  
from      Accounts  
where     accountId = 1;
```

```
update    Accounts  
set       balance = :xbal2 + 100  
where     accountId = 1;  
update    Accounts  
set       balance = :ybal2 - 100  
where     accountId = 2;  
commit;
```


Serializable Transaction Executions

- Consider a set of transactions $S = \{T_1, \dots, T_n\}$
- An execution of S is **serializable** if it is equivalent to some serial execution of S
- Let E_1 & E_2 denote two executions of S
- E_1 and E_2 are equivalent executions of S if
 1. both executions produce the same final database state, &
 2. both executions retrieve the same values: for every value read by some T_i in E_1 , the corresponding read by T_i in E_2 returns the same value
- **Serializable executions** guarantee correctness of transaction executions

Interleaved Executions: Example 1

An interleaved execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

```
begin transaction;  
select    balance into :xbal  
from      Accounts  
where     accountId = 1;  
select    balance into :ybal  
from      Accounts  
where     accountId = 2;  
update    Accounts  
set       balance = :ybal + 100  
where     accountId = 2;
```

```
update    Accounts  
set       balance = :xbal - 100  
where     accountId = 1;  
commit;
```

```
begin transaction;  
select    balance into :ybal2  
from      Accounts  
where     accountId = 2;  
select    balance into :xbal2  
from      Accounts  
where     accountId = 1;
```

```
update    Accounts  
set       balance = :xbal2 + 100  
where     accountId = 1;  
update    Accounts  
set       balance = :ybal2 - 100  
where     accountId = 2;  
commit;
```

Non-Serializable Executions: Example 1a

Transfer(1,2,100)	Transfer(2,1,100)	Comments
begin transaction;		Accounts: (1,1000), (2,2000)
select balance into :xbal from Accounts where accountId = 1;		xbal = 1000
select balance into :ybal from Accounts where accountId = 2;		ybal = 2000
update Accounts set balance = :ybal + 100 where accountId = 2;		Accounts: (1,1000), (2,2100)
	begin transaction;	
	select balance into :ybal2 from Accounts where accountId = 2;	ybal2 = 2100
	select balance into :xbal2 from Accounts where accountId = 1;	xbal2 = 1000
update Accounts set balance = :xbal - 100 where accountId = 1;		Accounts: (1,900), (2,2100)
commit;		
	update Accounts set balance = :xbal2 + 100 where accountId = 1;	Accounts: (1,1100), (2,2100)
	update Accounts set balance = :ybal2 - 100 where accountId = 2;	Accounts: (1,1100), (2,2000)
	commit;	

Non-Serializable Executions: Example 1b

Transfer(1,2,100)	Transfer(2,1,100)	Comments
begin transaction;		Accounts: (1,1000), (2,2000)
select balance into :xbal from Accounts where accountId = 1;		xbal = 1000
select balance into :ybal from Accounts where accountId = 2;		ybal = 2000
update Accounts set balance = :ybal + 100 where accountId = 2;		Accounts: (1,1000), (2,2100)
	begin transaction;	
	select balance into :ybal2 from Accounts where accountId = 2;	ybal2 = 2000
	select balance into :xbal2 from Accounts where accountId = 1;	xbal2 = 1000
update Accounts set balance = :xbal - 100 where accountId = 1;		Accounts: (1,900), (2,2100)
commit;		
	update Accounts set balance = :xbal2 + 100 where accountId = 1;	Accounts: (1,1100), (2,2100)
	update Accounts set balance = :ybal2 - 100 where accountId = 2;	Accounts: (1,1100), (2,1900)
	commit;	

Interleaved Executions: Example 2

Another interleaved execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

```
begin transaction;  
select    balance into :xbal  
from      Accounts  
where     accountId = 1;  
select    balance into :ybal  
from      Accounts  
where     accountId = 2;  
update    Accounts  
set       balance = :ybal + 100  
where     accountId = 2;
```

```
update    Accounts  
set       balance = :xbal - 100  
where     accountId = 1;  
commit;
```

```
begin transaction;  
select    balance into :ybal2  
from      Accounts  
where     accountId = 2;
```

```
select    balance into :xbal2  
from      Accounts  
where     accountId = 1;  
update    Accounts  
set       balance = :xbal2 + 100  
where     accountId = 1;  
update    Accounts  
set       balance = :ybal2 - 100  
where     accountId = 2;  
commit;
```

Serializable Executions: Example 2a

Transfer(1,2,100)	Transfer(2,1,100)	Comments
begin transaction;		Accounts: (1,1000), (2,2000)
select balance into :xbal from Accounts where accountId = 1;		xbal = 1000
select balance into :ybal from Accounts where accountId = 2;		ybal = 2000
update Accounts set balance = :ybal + 100 where accountId = 2;		Accounts: (1,1000), (2,2100)
	begin transaction;	
	select balance into :ybal2 from Accounts where accountId = 2;	ybal2 = 2100
update Accounts set balance = :xbal - 100 where accountId = 1;		Accounts: (1,900), (2,2100)
commit;		
	select balance into :xbal2 from Accounts where accountId = 1;	xbal2 = 900
	update Accounts set balance = :xbal2 + 100 where accountId = 1;	Accounts: (1,1000), (2,2100)
	update Accounts set balance = :ybal2 - 100 where accountId = 2;	Accounts: (1,1000), (2,2000)
	commit;	

Non-Serializable Executions: Example 2b

Transfer(1,2,100)	Transfer(2,1,100)	Comments
begin transaction;		Accounts: (1,1000), (2,2000)
select balance into :xbal from Accounts where accountId = 1;		xbal = 1000
select balance into :ybal from Accounts where accountId = 2;		ybal2 = 2000
update Accounts set balance = :ybal + 100 where accountId = 2;		Accounts: (1,1000), (2,2100)
	begin transaction;	
	select balance into :ybal2 from Accounts where accountId = 2;	ybal2 = 2000
update Accounts set balance = :xbal - 100 where accountId = 1;		Accounts: (1,900), (2,2100)
commit;		
	select balance into :xbal2 from Accounts where accountId = 1;	xbal2 = 900
	update Accounts set balance = :xbal2 + 100 where accountId = 1;	Accounts: (1,1000), (2,2100)
	update Accounts set balance = :ybal2 - 100 where accountId = 2;	Accounts: (1,1000), (2,1900)
	commit;	

Quiz 1

Is the the following execution of transactions T_1 & T_2 serializable?

- T_1 : Withdraw \$100 from accountId 1
- T_2 : Withdraw \$500 from accountId 1

```
begin transaction;  
select balance into :xbal from Accounts  
where accountId = 1;
```

```
update Accounts  
set balance = :xbal - 100  
where accountId = 1;  
commit;
```

```
begin transaction;  
select      balance into :xbal from Accounts  
where      accountId = 1;
```

```
update      Accounts  
set         balance = :xbal - 500  
where      accountId = 1;  
commit;
```


Quiz 2

Assume that the balances for accountId 1 and accountId 2 are \$1000 & \$2000, respectively. Consider a concurrent execution of transactions T_3 and T_4 where the values read by T_4 are \$1000 & \$2100. Is this concurrent execution serializable?

T_3 : **begin transaction;**
update Accounts
set balance = balance + 100
where accountId = 2;

update Accounts
set balance = balance - 100
where accountId = 1;
commit;

T_4 : **begin transaction;**
select balance
from Accounts
where accountId = 1;

select balance
from Accounts
where accountId = 2;
commit;

Quiz 3

Is the following execution of T_5 & T_6 serializable?

T_5	T_6	Comments
begin transaction;		Accounts: (1,1000), (2,2000)
select balance into :ybal from Accounts where accountId = 2;		ybal = 2000
	begin transaction;	
	select balance into :xbal2 from Accounts where accountId = 1;	xbal2 = 1000
	select balance into :ybal2 from Accounts where accountId = 2;	ybal2 = 2000
update Accounts set balance = :ybal + 200; where accountId = 2;		Accounts: (1,1000), (2,2200)
commit;		
	update Accounts set balance = :xbal2 + 100 where accountId = 1;	Accounts: (1,1100), (2,2200)
	commit;	

Quiz 4

Is the following execution of T_5 , T_6 & T_7 serializable?

T_5	T_6	T_7	Comments
begin transaction;			Accounts: (1,1000), (2,2000)
select balance into :ybal from Accounts where accountId = 2;			ybal = 2000
	begin transaction;		
	select balance into :xbal2 from Accounts where accountId = 1;		xbal2 = 1000
	select balance into :ybal2 from Accounts where accountId = 2;		ybal2 = 2000
update Accounts set balance = :ybal + 200; where accountId = 2;			Accounts: (1,1000), (2,2200)
commit;			
		begin transaction;	
		select balance from Accounts where accountId = 1;	1000
		select balance from Accounts where accountId = 2;	2200
		commit;	
	update Accounts set balance = :xbal2 + 100 where accountId = 1;		Accounts: (1,1100), (2,2200)
	commit;		

ANSI SQL Isolation Levels

- The isolation level for a transaction affects what the transaction will read
- **Dirty read** = read value is produced by a transaction that has not yet committed
- **Non-repeatable read** = successive reads of the same tuple yield different values
- **Phantom read** = successive reads of a set of tuples satisfying a predicate yield different values

Dirty Read: Example

Consider the following execution of transactions T_1 & T_2

- T_1 : Transfer(1,2,100)
- T_2 : Read the balances for accountIds 1 & 2

T_1	T_2	Comments
begin transaction;		Accounts: (1,1000), (2,2000)
select balance into :xbal from Accounts where accountId = 1;		xbal = 1000
select balance into :ybal from Accounts where accountId = 2;		ybal = 2000
update Accounts set balance = :ybal + 100 where accountId = 2;		Accounts: (1,1000), (2,2100)
	begin transaction;	
	select balance from Accounts where accountId = 1;	1000
	select balance from Accounts where accountId = 2;	2100 (dirty read!)
	commit;	
update Accounts set balance = :xbal - 100 where accountId = 1;		Accounts: (1,900), (2,2100)
commit;		

Non-Repeatable Read: Example

Consider the following execution of transactions T_1 & T_2

- T_1 : Transfer \$100 from accountId 1 to accountId 2
- T_2 : Two successive reads of the balance of accountId 2

T_1	T_2	Comments
	begin transaction;	Accounts: (1,1000), (2,2000)
	select balance from Accounts where accountId = 2;	2000
begin transaction;		
update Accounts set balance = balance + 100 where accountId = 2;		Accounts: (1,1000), (2,2100)
update Accounts set balance = balance - 100 where accountId = 1;		Accounts: (1,900), (2,2100)
commit;		
	select balance from Accounts where accountId = 2;	2100 (non-repeatable read!)
	commit;	

Phantom Read: Example

Consider the following execution of transactions T_1 & T_2

- T_1 : Two successive reads of tuples with balance more than \$1000.
- T_2 : Insert a new tuple with balance more than \$1000.

T_1	T_2	Comments
	begin transaction;	Accounts: (1,100), (2,2000)
	select accountId from Accounts where balance > 1000;	{2}
insert into Accounts values (3, 3000);		Accounts: (1,100), (2,2000), (3, 3000)
	select accountId from Accounts where balance > 1000;	{2, 3} (phantom read!)
	commit;	

ANSI SQL Isolation Levels

- ANSI SQL defines four isolation levels
- Configure using **set transaction isolation level** statement

Isolation Level	Dirty Read	Non-repeatable Read	Phantom Read
READ UNCOMMITTED	possible	possible	possible
READ COMMITTED	not possible	possible	possible
REPEATABLE READ	not possible	not possible	possible
SERIALIZABLE	not possible	not possible	not possible

Comments on SQL Isolation Levels

- In many DBMSs, running transactions at SQL's `SERIALIZABLE` isolation level guarantees that the execution is serializable
- However, depending on the application's transaction workload, a non-serializable execution could still be correct
- Ideally, use the weakest isolation level to obtain correct transaction executions for application

Transaction Example: Seat Booking

- Customers (custId, name, . . .)
- Seats (seatNumber, price)
- Bookings (eventDate, seatNumber, custId)

```
1 begin transaction;  
2 — Customer 123 wants to book a seat for June 20 2018  
3  
4 — Display available seats  
5 select seatNumber from Seats  
6 where seatNumber not in  
7     (select seatNumber from Bookings  
8      where eventDate = '2018-06-20');  
9  
10 — Customer 123 books seat 10  
11 insert into Bookings values ( '2018-06-20', 10, 123);  
12 commit;
```

Transaction Example: Seat Booking

- T_1 : Customer 123 books seat 10 for 2018-06-20
- T_2 : Customer 456 books seat 20 for 2018-06-20

begin transaction;

set transaction isolation level read committed;

select seatNumber **from** Seats
where seatNumber **not in** (
 select seatNumber **from** Bookings
 where eventdate = '2018-06-20');

insert into Bookings **values** ('2018-06-20',10,123);
commit;

begin transaction;

set transaction isolation level read committed;

select seatNumber **from** Seats
where seatNumber **not in** (
 select seatNumber **from** Bookings
 where eventdate = '2018-06-20');

insert into Bookings **values** ('2018-06-20',20,456);
commit;
