

CS2102

Database Systems

Slides adapted from Prof. Chan Chee Yong

LECTURE 03

SQL #1

Relational data model

Relation schema

- Each relation has a definition called a **relation schema**
 - Schema specifies **attributes** and **data constraints**
 - Data constraints include **domain constraints**
 - Each row in a relation is called a **tuple/record**

Relations

- A **relation** is defined as a set of tuples

Relational database schema

- A **relational database schema** consists of a set of schemas
- *Relational database schema*
= *relational schemas* + *data constraints*

Relational database

- A **relational database** is a collection of tables

Integrity constraints (ICs)

Definitions

- Integrity constraint
 - A condition that restricts the data that can be stored in database instance

Types

- Domain constraints
 - Restrict attribute values of relations
- Key constraints
- Foreign key constraints
- Other general constraints

Key constraints

Superkey

- A **superkey** is a subset of attributes in a relation that uniquely identifies its tuples
- No two distinct tuples of relation have the same values in all attributes of superkey

Key

- A **key** is a superkey that satisfies the additional property
 - Not null & no proper subset of a key is a superkey
 - Minimal subset of attributes that uniquely identifies its tuples
- A relation can have multiple keys
 - These are called **candidate keys**
 - One of the candidate keys is then selected as the **primary keys**

Foreign key constraints

Foreign key

- A subset of attributes in a relation is a **foreign key** if it refers to the primary key of a second relation
- **Foreign key constraints**
 - Each foreign key value in **referencing relation** must either
 - Appear as primary key value in **referenced relation**, or
 - Be a **null** value
 - Referencing & referenced relations could be the same relation
 - Also called **referential integrity constraints**

Relational algebra

Selection: σ_c

- $\sigma_c(R)$ selects tuples from relation R that satisfies selection condition c

Projection: π_ℓ

- $\pi_\ell(R)$ projects attributes given by a list ℓ of attributes from relation R

Renaming: $\rho_{S(B_1, B_2, \dots, B_n)}(R)$

- $\rho_{S(B_1, B_2, \dots, B_n)}(R)$ renames $R(A_1, A_2, \dots, A_n)$ to $S(B_1, B_2, \dots, B_n)$
 - When the attributes are not renamed $\rho_S(R)$
 - When the table is not renamed $\rho_{(B_1, B_2, \dots, B_n)}(R)$

Relational algebra

Union: $R \cup S$

- Returns a relation containing all tuples that occur in R , S , or both

Intersection: $R \cap S$

- Returns a relation containing all tuples that occur in both R and S

Set-difference: $R - S$

- Returns a relation containing all tuples that occur in R but not in S

❖ Union (\cup), intersection (\cap), and set-difference ($-$) operators require input relations to be **union compatible**

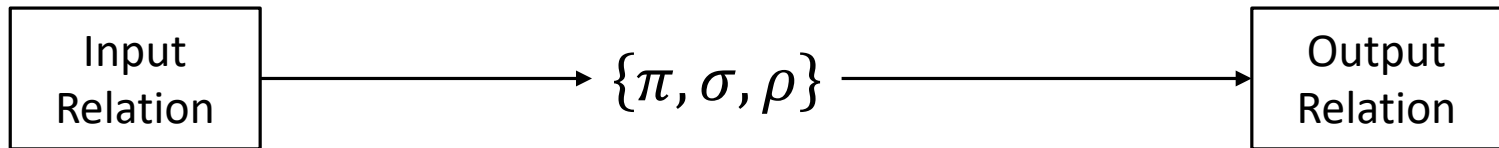
Cross-product: \times

- Consider a relation $R_1(A, B, C)$ and $R_2(X, Y)$
- $R_1 \times R_2$ returns a relation with schema (A, B, C, X, Y) defined as follows:
 - $R_1 \times R_2 = \{(a, b, c, x, y) \mid (a, b, c) \in R_1, (x, y) \in R_2\}$
- Also known as **cartesian product**

Closure properties

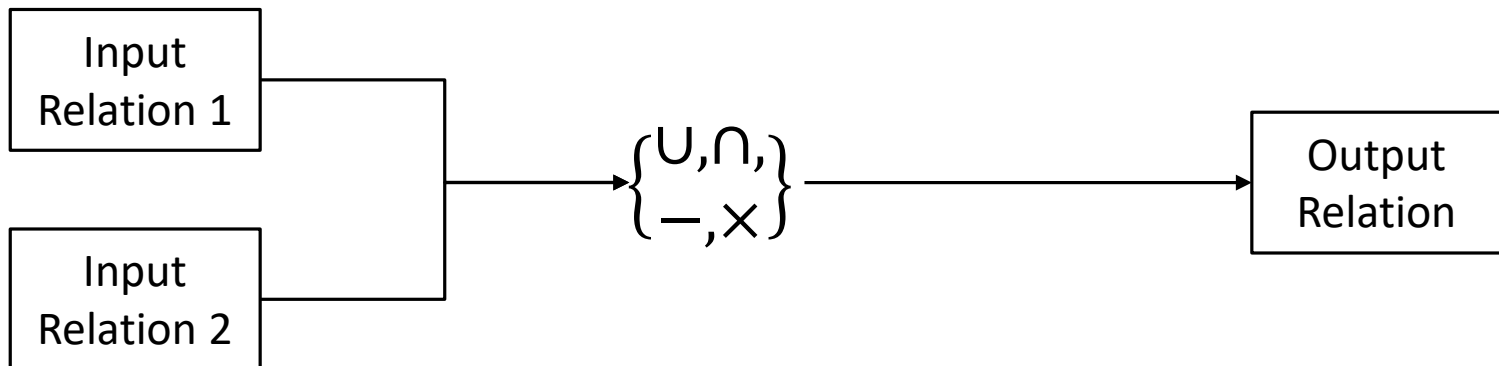
Closure of relation under unary operators (as diagrams)

- Unary operator takes in a relation as input and gives a relation as output



Closure of relation under binary operators (as diagrams)

- Binary operator takes in two relations as inputs and gives a relation as output



- Structured Query Language

- Null values

- Create/drop table

- Table modification

- Constraint checking

- Queries

- Simple queries

Overview

- Structured Query Language

- Null values

- Create/drop table

- Table modification

- Constraint checking

- Queries

- Simple queries

Structured Query Language

Structured Query Language

Introduction

- Designed by D. Chamberlin & R. Boyce (IBM Research) in 1974
- Original name: SEQUEL (*Structured English QUery Language*)
- SQL is not a general-purpose language but a domain-specific language (DSL)
 - Designed for computations on relations
- Unlike relational algebra which is a procedural language, SQL is a **declarative language**
- Focuses on what to compute instead of how to compute

Structured Query Language

Using SQL

- **Directly write SQL statements**
 - *Command line interface*
 - PostgreSQL's **psql**
 - *Graphical user interface*
 - PostgreSQL's **pgAdmin**
- **Included SQL in application programs**
 - *Statement-level interface (SLI)*
 - Embedded SQL
 - Dynamic SQL
 - *Call-level interface (CLI)*
 - JDBC (Java DataBase Connectivity)
 - ODBC (Open DataBase Connectivity)

Structured Query Language

About SQL

- Current ANSI/ISO standard for SQL is called SQL:2016
- Different DBMSs may have minor variations in SQL syntax
- SQL consists of two main parts:
 - **Data Definition Language (DDL)**
 - Used to create/delete/modify schemas
 - **Data Manipulation Language**
 - Used to ask queries/insert/delete/modify data

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

<i>x</i>	<i>y</i>	<i>x AND y</i>	<i>x OR y</i>	<i>NOT x</i>
FALSE FALSE FALSE	FALSE UNKNOWN TRUE			
UNKNOWN UNKNOWN UNKNOWN	FALSE UNKNOWN TRUE			
TRUE TRUE TRUE	FALSE UNKNOWN TRUE			

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Operation

- Result of **comparison operation** involving null value is unknown
- Result of arithmetic operation involving null value is null

Example: assume that the value of x is null

- `x < 100` →
- `x = null` →
- `x <> null` →
- `x + 20` →

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Check

- How to check if a value is equal to null?
- Use the **IS NULL** comparison predicate

<i>x</i>	<i>x IS NULL</i>	<i>x IS NOT NULL</i>
null		
non-null		

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Check

- How to treat `null` values as ordinary values for comparison?
- Use the **IS DISTINCT FROM** comparison predicate

<i>x</i>	<i>y</i>	<i>x IS DISTINCT FROM y</i>
<code>null</code>		
<code>null</code>		
<code>non-null</code>		
<code>non-null</code>		

SQL syntax

Comments

- Comments in SQL are preceded by **two hyphens**
- `-- this is a single-line comment`

C-style comments

- SQL also supports **C-style comments**
- `/* this is also a comment
and can be multi-line */`

Grammar

- Keywords are in **UPPERCASE**, variables are in **lowercase**
- Optional components are in `[square brackets]`
- Choices are separated by `|`
- Non-optional choices are in `{ curly brackets }`
- Possibly infinite repetitions are denoted by `...`

Create/drop table

Create table syntax

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
    | table_constraints }  
    [, ...]  
] );
```

- Reference: <https://www.postgresql.org/docs/11/sql-createtable.html>

Drop table syntax

```
DROP TABLE [ IF EXISTS ] table_name
```

- Reference: <https://www.postgresql.org/docs/11/sql-droptable.html>

Create/drop table

Examples

```
CREATE TABLE Students (  
    -- column1  data_type1  
    -- column2  data_type2  
    -- column3  data_type3  
    -- etc  
);
```

```
DROP TABLE Students;
```

Create/drop table

Specifying constraints

- Domain constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```

Create/drop table

Specifying constraints

- Domain constraints
 - Built-in data types
 - Domain of each includes a special value null
 - Reference: <https://www.postgresql.org/docs/11/datatype.html>

<i>Data Type</i>	<i>Values</i>
boolean	false/true
integer	signed four-byte integer
float8	double-precision floating-point number (8 bytes)
numeric	arbitrary-precision floating-point number
numeric(p,s)	maximum total of p digits with maximum of s in fractional part
char(n)	fixed-length string consisting of n characters
varchar(n)	variable-length string up to n characters
text	variable-length character string
date	calendar date (year, month, day)
timestamp	date and time

Create/drop table

Examples

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100),  
    birthDate    date  
);
```

students

<i>studentID</i>	<i>name</i>	<i>birthDate</i>
3118	Alice	1999-12-25
3118	Trudy	1999-12-25
1423	Bob	2000-05-27
5609	Carol	1999-06-11

Create/drop table

Specifying constraints

- Key constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```

Create/drop table

Specifying constraints

- Key constraints
 - Recap:
 - Superkey *uniquely identifies its tuples*
 - No two distinct tuples of relations have the same values in all attributes of superkey
 - Key *minimal superkey*
 - Candidate key *if there are multiple keys*
 - Primary key *one is selected as primary key*
 - Keyword: PRIMARY KEY

Create/drop table

Examples

```
CREATE TABLE Students ( -- column constraints
    studentID      integer PRIMARY KEY, -- one PK
    name            varchar(100),
    birthDate       date
);
```

Students

<u>studentID</u>	name	birthDate
3118	Alice	1999-12-25
3118	Trudy	1999-12-25
1423	Bob	2000-05-27
5609	Carol	1999-06-11

Students

<u>studentID</u>	name	birthDate
3118	Alice	1999-12-25
1423	Bob	2000-05-27
5609	Carol	1999-06-11

Create/drop table

Examples

```
CREATE TABLE Enrolls ( -- table constraints
    sid    integer,
    cid    integer,
    grade  char(2),
    PRIMARY KEY(sid, cid) -- one or more PK
);
```

Enrolls

<u>sid</u>	<u>cid</u>	grade
3118	112	3.8
1423	101	4.0
1423	311	3.8
3118	112	4.0

Enrolls

<u>sid</u>	<u>cid</u>	grade
1423	101	4.0
1423	311	3.8
3118	112	4.0

Create/drop table

Specifying constraints

- Foreign key constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```

Create/drop table

Specifying constraints

- Foreign key constraints
 - Recap:
 - Foreign key *refers to primary key of second relation*
 - Each foreign key value in **referencing relation** must either
 - Appear as primary key value in **referenced relation**, or
 - Be a **null** value
 - Referencing & referenced relations could be the same relation
 - Also called **referential integrity constraints**
 - Keyword: FOREIGN KEY and REFERENCES

Create/drop table

Examples

```
CREATE TABLE Enrolls (  
    sid      integer REFERENCES Students(studentID),  
    cid      integer, -- assume students and courses  
    grade    char(2), -- table are defined  
    FOREIGN KEY(cid) REFERENCES Courses(courseID)  
);
```

Students

<u>studentID</u>
3118
1423
5609

Enrolls

<i>sid</i>	<i>cid</i>	<i>grade</i>
1423	101	4.0
1423	311	3.8
3118	112	4.0
null	313	4.5
5609	null	5.0
null	null	0.0
1111	312	1.2

Courses

<u>courseID</u>
101
312
112

Create/drop table

Examples

```
CREATE TABLE Enrolls (  
    sid    integer, -- assume Student_Course table  
    cid    integer, -- is defined  
  
    FOREIGN KEY(sid,cid) REFERENCES  
        Student_Course(studentID, courseID)  
);
```

Student_Course

<u>studentID</u>	<u>courseID</u>
3118	112
1423	101
1423	312
5609	112

Enrolls

<i>sid</i>	<i>cid</i>	<i>grade</i>
1423	101	4.0
1423	311	3.8
3118	112	4.0
null	313	4.5
5609	null	5.0
null	null	0.0
1111	312	1.2

Constraint checking

Foreign key constraint violation

- Deletion/update of a referenced tuple could violate a foreign key constraint
- Specify action to deal with violation as part of a foreign key constraint declaration

FOREIGN KEY(...) REFERENCES ...

[ON DELETE **action**] [ON UPDATE **action**]

Students

<u><i>studentID</i></u>
3118
1423
5609

Enrolls

<i>sid</i>	<i>cid</i>	<i>grade</i>
1423	101	4.0
3118	112	4.0
5609	312	1.2

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

[ON DELETE **action**] [ON UPDATE **action**]

Actions

- **NO ACTION** rejects delete/update if it violates constraint (*default option*)
- **RESTRICT** similar to NO ACTION except that constraint checking can't be deferred
- **CASCADE** propagate delete/update to referencing tuple
- **SET DEFAULT** updates foreign keys of referencing tuples to some *default value*
- **SET NULL** updates to null value

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

-- delete 3118

-- change 1423 to 1444

Students

<u>studentID</u>
1444
5609

Enrolls

<i>sid</i>	<i>cid</i>	<i>grade</i>
1444	101	4.0
null	112	4.0
5609	312	1.2

Create/drop table

Specifying constraints

- Other general constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```

Create/drop table

Specifying constraints

- Other general constraints
 - **Not-null** constraints *keyword: NOT NULL*
 - **Unique** constraints *keyword: UNIQUE*
 - Check any two records $x, y \in \text{table}$ such that
$$x.\text{column} \neq y.\text{column}$$
 - **General** constraints *keyword: CHECK*
 - ❖ Constraint is **violated** if it evaluates to FALSE
- Default values
 - Specify default values *keyword: DEFAULT*

Create/drop table

Examples: not-null

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100) NOT NULL,  
    birthDate    date  
);
```

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100),  
    birthDate    date,  
    CHECK (name IS NOT NULL)  
);
```

Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100) UNIQUE,  
    birthDate      date  
);
```

Students

<u>studentID</u>	name	birthDate
3118	Alice	1999-12-25
3119	Alice	1999-12-25
1423	Bob	2000-05-27
5609	Carol	1999-06-11

Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100),  
    birthDate      date,  
    UNIQUE (studentID, name)  
);
```

Students

<u>studentID</u>	name	birthDate
3118	Alice	1999-12-25
3119	Alice	1999-12-25
1423	Bob	2000-05-27
1423	Bob	1999-06-11

Create/drop table

Examples: general

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100),  
    birthDate      date,  
    CHECK (studentID > 2000)  
);
```

Students

<u>studentID</u>	name	birthDate
3118	Alice	1999-12-25
3119	Alice	1999-12-25
1423	Bob	2000-05-27
5609	Carol	1999-06-11

Create/drop table

Examples: default values

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100) DEFAULT 'John Doe',  
    birthDate    date  
);
```

Create/drop table

Assertions

- Complex constraints
 - Multi-table constraints
 - Example:
 - Every student in `students` table *must be enrolled in at least two and at most five courses* in `enrolls` table
- Limitation
 - Create assertion statement is not implemented in many DBMSs
 - We can use **triggers** to enforce complex constraints
 - Lecture 06

Table modification

Insert into syntax

```
INSERT INTO table_name  
[ ( column_name [, ...] ) ]  
VALUES ( { expression | DEFAULT } [, ...] );
```

Delete from syntax

```
DELETE FROM table_name  
[ WHERE condition ];
```

Update syntax

```
UPDATE table_name  
SET column_name = { expression | DEFAULT }  
[ WHERE condition ];
```




Table modification

Example

```
CREATE TABLE Students (  
    studentID    integer PRIMARY KEY,  
    name         varchar(100),  
    dept         varchar(20) DEFAULT 'CS'  
);
```

```
INSERT INTO Students VALUES (12345, 'Alice', 'Eng');
```

```
INSERT INTO Students (studentID) VALUES (23456);
```

```
INSERT INTO Students (studentID) VALUES (12345);
```

Students

<u>studentID</u>	name	dept
12345	Alice	Eng
23456	null	CS

Table modification

Example

```
INSERT INTO Students VALUES (34567, 'Bob', 'Eng');
```

```
INSERT INTO Students (dept, name, studentID)
```

```
VALUES ('Maths', 'Carol', 45678);
```

```
DELETE FROM Students WHERE dept='Eng';
```

```
DELETE FROM Students; -- remove all
```

```
INSERT INTO Students VALUES (12345, 'Alice', 'Eng');
```

```
INSERT INTO Students (studentID) VALUES (23456);
```

Students

<u>studentID</u>	name	dept
12345	Alice	Eng
23456	null	CS
34567	Bob	Eng
45678	Carol	Maths

Table modification

Example

```
UPDATE Students SET name = 'Bob' WHERE dept = 'CS';
```

```
UPDATE Students SET studentID = studentID + 1;
```

```
-- update all
```

students

<u>studentID</u>	name	dept
12346	Alice	Eng
23457	Bob	CS

Schema modification

Alter table

- Add/remove/modify columns
 - `ALTER TABLE Students ALTER COLUMN dept DROP DEFAULT;`
 - `ALTER TABLE Students DROP COLUMN dept;`
 - `ALTER TABLE Students ADD COLUMN faculty varchar(20);`
 - etc
- Add/remove constraints
- etc

- Structured Query Language

- Null values

- Create/drop table

- Table modification

- Constraint checking

- Queries

- Simple queries

Queries

Simple queries

Basic syntax

- Basic form of **SQL query** consists of three clauses

```
SELECT [ DISTINCT ] select_list -- select clause
FROM      from_list    -- from clause
[ WHERE      condition ] -- where clause
```

- **select_list** specifies columns to be included in output
- **from_list** specifies list of relations
- **condition** specifies conditions on relations

- ❖ **Output:** relation generated from `from_list` containing attributes based on `select_list` that satisfies `condition`
 - Output relation could contain duplicate record if `DISTINCT` is not used in the `SELECT` clause

Simple queries


Basic syntax

- Basic form of **SQL query** consists of three clauses

SELECT DISTINCT a_1, a_2, \dots, a_m -- select clause
FROM r_1, r_2, \dots, r_n -- from clause
WHERE c -- where clause

- Relational algebra form**

- $\pi_{a_1, a_2, \dots, a_m}(\sigma_c(r_1 \times r_2 \times \dots \times r_n))$



What about renaming? ρ

Simple queries


Basic syntax

- Basic form of **SQL query** consists of three clauses

SELECT DISTINCT a_1, a_2, \dots, a_m -- select clause
FROM r_1, r_2, \dots, r_n -- from clause
WHERE c -- where clause

- Relational algebra form**

- $\pi_{a_1, a_2, \dots, a_m}(\sigma_c(r_1 \times r_2 \times \dots \times r_n))$



Simple queries

Example

- Find the names of restaurants, the pizzas that they sell, and their prices, where the price is under \$20
- Deconstruct problem
 - Find what?
 - From which relation?
 - What condition?
- Construct query

Output

<u><i>rname</i></u>	<u><i>pizza</i></u>	<i>price</i>
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Pizza King	Diavola	17

Sells

<u><i>rname</i></u>	<u><i>pizza</i></u>	<i>price</i>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Simple queries

Example

- Find the names of restaurants, the pizzas that they sell, and their prices, where the price is under \$20

- Construct query

```
SELECT rname, pizza, price
FROM   Sells
WHERE  price < 20;
```

- Alternative query

```
SELECT *
FROM   Sells
WHERE  price < 20;
```

Output

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Pizza King	Diavola	17

Sells

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Simple queries

Example

- Find all restaurants, the pizzas that they sell, and their prices, where (1) either the price is under \$20 or the pizza is “Marinara”, and (2) the pizza is not “Diavola”
- Construct query

```
SELECT rname, pizza, price -- or simply *
FROM   Sells
WHERE  (price < 20 OR pizza = 'Marinara')
      AND pizza <> 'Diavola';
```

Output

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Mamma's Place	Marinara	22

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Simple queries

Removing duplicates

- q1: SELECT a, c FROM r;
- q2: SELECT DISTINCT a, c FROM r;

r			q1		q2	
a	b	c	a	c	a	c
10	1	2	10	2	10	2
10	7	2	10	2	20	Null
20	3	Null	20	Null	30	2
20	9	Null	20	Null	30	9
30	3	2	30	2		
30	5	9	30	9		

❖ Two tuples (a_1, a_2, \dots, a_n) and (b_1, b_2, \dots, b_n) are considered to be **distinct** if the following evaluates to TRUE:

- ❖ $(a_1 \text{ IS DISTINCT FROM } b_1)$ or
 $(a_2 \text{ IS DISTINCT FROM } b_2)$ or
 ... or
 $(a_n \text{ IS DISTINCT FROM } b_n)$

Simple queries

Renaming column

- q: `SELECT item, price * qty AS cost FROM Orders;`

Orders			q	
<i>item</i>	<i>price</i>	<i>qty</i>	<i>item</i>	<i>Cost</i>
A	2.50	100	A	250.00
B	4.00	100	B	400.00
C	7.50	100	C	750.00

Renaming values

- q: `SELECT 'Price of ' || pizza || ' is ' ||
round(price / 1.3) || ' USD' AS menu
FROM Sells WHERE rname = 'Corleone Corner';`

Sells		
<u><i>rname</i></u>	<u><i>pizza</i></u>	<i>price</i>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Pizza King	Hawaiian	21

q	
<i>menu</i>	
Price of Diavola is 18 USD	
Price of Hawaiian is 19 USD	
Price of Margherita is 15 USD	

Summary

- ❑ SQL is the standard query language for relational DBMS
 - ❑ Table creation `CREATE TABLE table_name`
 - ❑ Table removal `DROP TABLE table_name`
 - ❑ Modification `ALTER TABLE table_name`
 - ❑ Insert `INSERT INTO table_name VALUES (..)`
 - ❑ Delete `DELETE FROM table_name WHERE ..`
 - ❑ Update `UPDATE table_name SET .. WHERE ..`
 - ❑ Queries
`SELECT DISTINCT a_1, a_2, \dots, a_m`
`FROM r_1, r_2, \dots, r_n`
`WHERE c`
 - ❑ $\pi_{a_1, a_2, \dots, a_m}(\sigma_c(r_1 \times r_2 \times \dots \times r_n))$