

CS2102

Database Systems

Slides adapted from Prof. Chan Chee Yong

LECTURE 03

SQL #1

Relational data model

Relation schema

- Each relation has a definition called a **relation schema**
 - Schema specifies **attributes** and **data constraints**
 - Data constraints include **domain constraints**
 - Each row in a relation is called a **tuple/record**

Relations

- A **relation** is defined as a set of tuples

Relational database schema

- A **relational database schema** consists of a set of schemas
- *Relational database schema*
= *relational schemas + data constraints*

Relational database

- A **relational database** is a collection of tables

Integrity constraints (ICs)

Definitions

- Integrity constraint
 - A condition that restricts the data that can be stored in database instance

Types

- Domain constraints
 - Restrict attribute values of relations
- Key constraints
- Foreign key constraints
- Other general constraints

Key constraints

Superkey

- A **superkey** is a subset of attributes in a relation that uniquely identifies its tuples
- No two distinct tuples of relation have the same values in all attributes of superkey

Key

- A **key** is a superkey that satisfies the additional property
 - Not null & no proper subset of a key is a superkey
 - Minimal subset of attributes that uniquely identifies its tuples
- A relation can have multiple keys
 - These are called **candidate keys**
 - One of the candidate keys is then selected as the **primary keys**

Foreign key constraints

Foreign key

- A subset of attributes in a relation is a **foreign key** if it refers to the primary key of a second relation
- **Foreign key constraints**
 - Each foreign key value in **referencing relation** must either
 - Appear as primary key value in **referenced relation**, or
 - Be a **null** value
 - Referencing & referenced relations could be the same relation
 - Also called **referential integrity constraints**

Relational algebra

Selection: σ_c

- $\sigma_c(R)$ selects tuples from relation R that satisfies selection condition c

Projection: π_ℓ

- $\pi_\ell(R)$ projects attributes given by a list ℓ of attributes from relation R

Renaming: $\rho_{S(B_1, B_2, \dots, B_n)}(R)$

- $\rho_{S(B_1, B_2, \dots, B_n)}(R)$ renames $R(A_1, A_2, \dots, A_n)$ to $S(B_1, B_2, \dots, B_n)$
 - When the attributes are not renamed $\rho_S(R)$
 - When the table is not renamed $\rho_{(B_1, B_2, \dots, B_n)}(R)$

Relational algebra

Union: $R \cup S$

- Returns a relation containing all tuples that occur in R , S , or both

Intersection: $R \cap S$

- Returns a relation containing all tuples that occur in both R and S

Set-difference: $R - S$

- Returns a relation containing all tuples that occur in R but not in S

❖ Union (\cup), intersection (\cap), and set-difference ($-$) operators require input relations to be **union compatible**

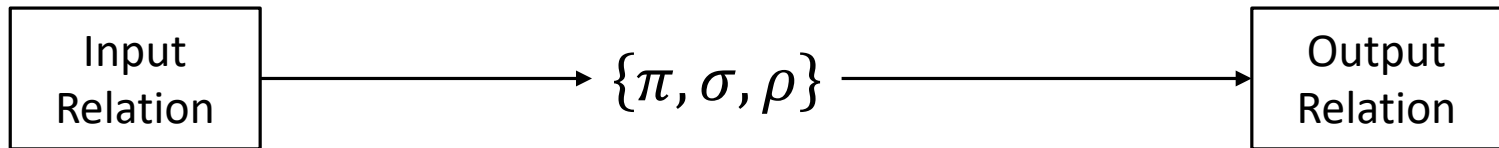
Cross-product: \times

- Consider a relation $R_1(A, B, C)$ and $R_2(X, Y)$
- $R_1 \times R_2$ returns a relation with schema (A, B, C, X, Y) defined as follows:
 - $R_1 \times R_2 = \{(a, b, c, x, y) \mid (a, b, c) \in R_1, (x, y) \in R_2\}$
- Also known as **cartesian product**

Closure properties

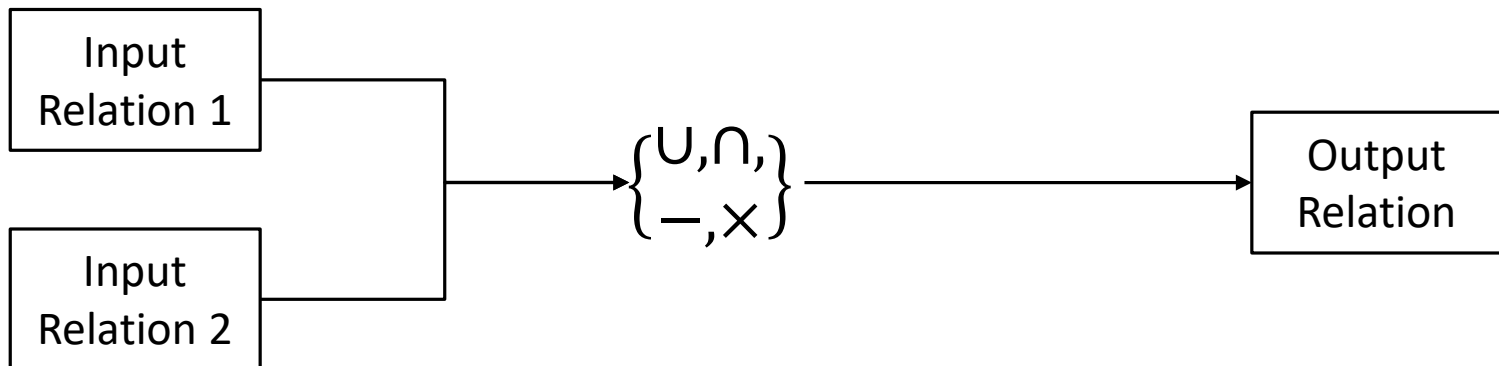
Closure of relation under unary operators (as diagrams)

- Unary operator takes in a relation as input and gives a relation as output



Closure of relation under binary operators (as diagrams)

- Binary operator takes in two relations as inputs and gives a relation as output



- Structured Query Language

- Null values

- Create/drop table

- Table modification

- Constraint checking

- Queries

- Simple queries

Overview

- Structured Query Language

- Null values

- Create/drop table

- Table modification

- Constraint checking

- Queries

- Simple queries

Structured Query Language

Structured Query Language

Introduction

- Designed by D. Chamberlin & R. Boyce (IBM Research) in 1974
- Original name: SEQUEL (*Structured English QUery Language*)
- SQL is not a general-purpose language but a domain-specific language (DSL)
 - Designed for computations on relations
- Unlike relational algebra which is a procedural language, SQL is a **declarative language**
- Focuses on what to compute instead of how to compute

Structured Query Language

Using SQL

- **Directly write SQL statements**
 - *Command line interface*
 - PostgreSQL's **psql**
 - *Graphical user interface*
 - PostgreSQL's **pgAdmin**
- **Included SQL in application programs**
 - *Statement-level interface (SLI)*
 - Embedded SQL
 - Dynamic SQL
 - *Call-level interface (CLI)*
 - JDBC (Java DataBase Connectivity)
 - ODBC (Open DataBase Connectivity)

Structured Query Language

About SQL

- Current ANSI/ISO standard for SQL is called SQL:2016
- Different DBMSs may have minor variations in SQL syntax
- SQL consists of two main parts:
 - **Data Definition Language (DDL)**
 - Used to create/delete/modify schemas
 - **Data Manipulation Language**
 - Used to ask queries/insert/delete/modify data

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

| <i>x</i> | <i>y</i> | <i>x AND y</i> | <i>x OR y</i> | <i>NOT x</i> |
|-------------------------------|--------------------------|----------------|---------------|--------------|
| FALSE FALSE FALSE | FALSE UNKNOWN TRUE | | | |
| UNKNOWN UNKNOWN UNKNOWN | FALSE UNKNOWN TRUE | | | |
| TRUE TRUE TRUE | FALSE UNKNOWN TRUE | | | |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

| <i>x</i> | <i>y</i> | <i>x AND y</i> | <i>x OR y</i> | <i>NOT x</i> |
|-------------------------------|--------------------------|--------------------|-------------------|--------------|
| FALSE FALSE FALSE | FALSE UNKNOWN TRUE | FALSE FALSE | FALSE TRUE | TRUE |
| UNKNOWN UNKNOWN UNKNOWN | FALSE UNKNOWN TRUE | | | |
| TRUE TRUE TRUE | FALSE UNKNOWN TRUE | FALSE TRUE | TRUE TRUE | FALSE |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

| <i>x</i> | <i>y</i> | <i>x AND y</i> | <i>x OR y</i> | <i>NOT x</i> |
|-------------------------------|--------------------------|-----------------------------|---------------|--------------|
| FALSE FALSE FALSE | FALSE UNKNOWN TRUE | FALSE FALSE FALSE | FALSE TRUE | TRUE |
| UNKNOWN UNKNOWN UNKNOWN | FALSE UNKNOWN TRUE | FALSE UNKNOWN UNKNOWN | | |
| TRUE TRUE TRUE | FALSE UNKNOWN TRUE | FALSE UNKNOWN TRUE | TRUE TRUE | FALSE |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

| <i>x</i> | <i>y</i> | <i>x AND y</i> | <i>x OR y</i> | <i>NOT x</i> |
|-------------------------------|--------------------------|-----------------------------|----------------------------|--------------|
| FALSE FALSE FALSE | FALSE UNKNOWN TRUE | FALSE FALSE FALSE | FALSE UNKNOWN TRUE | TRUE |
| UNKNOWN UNKNOWN UNKNOWN | FALSE UNKNOWN TRUE | FALSE UNKNOWN UNKNOWN | UNKNOWN UNKNOWN TRUE | |
| TRUE TRUE TRUE | FALSE UNKNOWN TRUE | FALSE UNKNOWN TRUE | TRUE TRUE TRUE | FALSE |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

| <i>x</i> | <i>y</i> | <i>x AND y</i> | <i>x OR y</i> | <i>NOT x</i> |
|-------------------------------|--------------------------|-----------------------------|----------------------------|--------------|
| FALSE FALSE FALSE | FALSE UNKNOWN TRUE | FALSE FALSE FALSE | FALSE UNKNOWN TRUE | TRUE |
| UNKNOWN UNKNOWN UNKNOWN | FALSE UNKNOWN TRUE | FALSE UNKNOWN UNKNOWN | UNKNOWN UNKNOWN TRUE | UNKNOWN |
| TRUE TRUE TRUE | FALSE UNKNOWN TRUE | FALSE UNKNOWN TRUE | TRUE TRUE TRUE | FALSE |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

| <i>x</i> | <i>y</i> | <i>x AND y</i> | <i>x OR y</i> | <i>NOT x</i> |
|-------------------------------|--------------------------|-----------------------------|----------------------------|--------------|
| FALSE FALSE FALSE | FALSE UNKNOWN TRUE | FALSE FALSE FALSE | FALSE UNKNOWN TRUE | TRUE |
| UNKNOWN UNKNOWN UNKNOWN | FALSE UNKNOWN TRUE | FALSE UNKNOWN UNKNOWN | UNKNOWN UNKNOWN TRUE | UNKNOWN |
| TRUE TRUE TRUE | FALSE UNKNOWN TRUE | FALSE UNKNOWN TRUE | TRUE TRUE TRUE | FALSE |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Operation

- Result of **comparison operation** involving null value is unknown
- Result of arithmetic operation involving null value is null

Example: assume that the value of x is null

- `x < 100` →
- `x = null` →
- `x <> null` →
- `x + 20` →

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Operation

- Result of **comparison operation** involving null value is unknown
- Result of arithmetic operation involving null value is null

Example: assume that the value of x is null

- `x < 100` → unknown
- `x = null` → unknown
- `x <> null` → unknown
- `x + 20` → null

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Operation

- Result of **comparison operation** involving null value is unknown
- Result of arithmetic operation involving null value is null

Example: assume that the value of x is null

- `x < 100` → unknown
- `x = null` → unknown
- `x <> null` → unknown
- `x + 20` → null

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Operation

- Result of **comparison operation** involving `null` value is unknown
- Result of arithmetic operation involving `null` value is null

Example: assume that the value of `x` is `null`

- `x < 100` → unknown
- `x = null` → unknown
- `x <> null` → unknown
- `x + 20` → null

We cannot say if two null values are the same or different from one another because they are unknown

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Check

- How to check if a value is equal to null?
- Use the **IS NULL** comparison predicate

| <i>x</i> | <i>x IS NULL</i> | <i>x IS NOT NULL</i> |
|----------|------------------|----------------------|
| null | | |
| non-null | | |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Check

- How to check if a value is equal to null?
- Use the **IS NULL** comparison predicate

| <i>x</i> | <i>x IS NULL</i> | <i>x IS NOT NULL</i> |
|----------|------------------|----------------------|
| null | TRUE | FALSE |
| non-null | FALSE | TRUE |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Check

- How to treat `null` values as ordinary values for comparison?
- Use the **IS DISTINCT FROM** comparison predicate

| <i>x</i> | <i>y</i> | <i>x IS DISTINCT FROM y</i> |
|-----------------------|----------|-----------------------------|
| <code>null</code> | | |
| <code>null</code> | | |
| <code>non-null</code> | | |
| <code>non-null</code> | | |

Null values

Three-valued logic system

- TRUE
- FALSE
- UNKNOWN

Check

- How to treat `null` values as ordinary values for comparison?
- Use the **IS DISTINCT FROM** comparison predicate

| <i>x</i> | <i>y</i> | <i>x IS DISTINCT FROM y</i> |
|-----------------------|-----------------------|----------------------------------|
| <code>null</code> | <code>null</code> | FALSE |
| <code>null</code> | <code>non-null</code> | TRUE |
| <code>non-null</code> | <code>null</code> | TRUE |
| <code>non-null</code> | <code>non-null</code> | <code>x <> y</code> |

SQL syntax

Comments

- Comments in SQL are preceded by **two hyphens**
- `-- this is a single-line comment`

C-style comments

- SQL also supports **C-style comments**
- `/* this is also a comment
and can be multi-line */`

Grammar

- Keywords are in **UPPERCASE**, variables are in **lowercase**
- Optional components are in `[square brackets]`
- Choices are separated by `|`
- Non-optional choices are in `{ curly brackets }`
- Possibly infinite repetitions are denoted by `...`

Create/drop table

Create table syntax

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
    | table_constraints }  
    [, ...]  
] );
```

- Reference: <https://www.postgresql.org/docs/11/sql-createtable.html>

Drop table syntax

```
DROP TABLE [ IF EXISTS ] table_name
```

- Reference: <https://www.postgresql.org/docs/11/sql-droptable.html>

Create/drop table

Examples

```
CREATE TABLE Students (  
    -- column1  data_type1  
    -- column2  data_type2  
    -- column3  data_type3  
    -- etc  
);
```

```
DROP TABLE Students;
```

Create/drop table

Examples

```
CREATE TABLE Students (  
    -- column1  data_type1  
    -- column2  data_type2  
    -- column3  data_type3  
    -- etc  
);
```

```
DROP TABLE Students;
```

SQL is case-insensitive

Create/drop table

Specifying constraints

- Domain constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```


Create/drop table

Specifying constraints

- Domain constraints
 - Built-in data types
 - Domain of each includes a special value null
 - Reference: <https://www.postgresql.org/docs/11/datatype.html>

| <i>Data Type</i> | <i>Values</i> |
|------------------|--|
| boolean | false/true |
| integer | signed four-byte integer |
| float8 | double-precision floating-point number (8 bytes) |
| numeric | arbitrary-precision floating-point number |
| numeric(p,s) | maximum total of p digits with maximum of s in fractional part |
| char(n) | fixed-length string consisting of n characters |
| varchar(n) | variable-length string up to n characters |
| text | variable-length character string |
| date | calendar date (year, month, day) |
| timestamp | date and time |

Create/drop table

Examples

```
CREATE TABLE students (  
    studentID    integer,  
    name         varchar(100),  
    birthDate    date  
  
);
```

students

| <i>studentID</i> | <i>name</i> | <i>birthDate</i> |
|------------------|-------------|------------------|
| 3118 | Alice | 1999-12-25 |
| 3118 | Trudy | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Specifying constraints

- Key constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```

Create/drop table

Specifying constraints

- Key constraints
 - Recap:
 - Superkey *uniquely identifies its tuples*
 - *No two distinct tuples of relations have the same values in all attributes of superkey*
 - Key *minimal superkey*
 - Candidate key *if there are multiple keys*
 - Primary key *one is selected as primary key*
 - Keyword: PRIMARY KEY

Create/drop table

Examples

```
CREATE TABLE Students ( -- column constraints
    studentID      integer PRIMARY KEY, -- one PK
    name           varchar(100),
    birthDate      date
);
```

Create/drop table

Examples

```
CREATE TABLE Students ( -- column constraints
    studentID      integer PRIMARY KEY, -- one PK
    name           varchar(100),
    birthDate      date
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|-------|------------|
| 3118 | Alice | 1999-12-25 |
| 3118 | Trudy | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Examples

```
CREATE TABLE Students ( -- column constraints
    studentID      integer PRIMARY KEY, -- one PK
    name            varchar(100),
    birthDate       date
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|------------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3118 | Trudy | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Examples

```
CREATE TABLE Students ( -- column constraints
    studentID      integer PRIMARY KEY, -- one PK
    name           varchar(100),
    birthDate      date
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|------------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3118 | Trudy | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Students

| <u>studentID</u> | name | birthDate |
|------------------|-------|------------|
| 3118 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Examples

```
CREATE TABLE Enrolls ( -- table constraints
    sid    integer,
    cid    integer,
    grade  char(2),
    PRIMARY KEY(sid, cid) -- one or more PK
);
```

Create/drop table

Examples

```
CREATE TABLE enrolls ( -- table constraints
    sid    integer,
    cid    integer,
    grade  char(2),
    PRIMARY KEY(sid, cid) -- one or more PK
);
```

Enrolls

| <u><i>sid</i></u> | <u><i>cid</i></u> | <i>grade</i> |
|-------------------|-------------------|--------------|
| 3118 | 112 | 3.8 |
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |

Create/drop table

Examples

```
CREATE TABLE enrolls ( -- table constraints
    sid    integer,
    cid    integer,
    grade  char(2),
    PRIMARY KEY(sid, cid) -- one or more PK
);
```

Enrolls

| <u>sid</u> | <u>cid</u> | grade |
|------------|------------|-------|
| 3118 | 112 | 3.8 |
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |

Create/drop table

Examples

```
CREATE TABLE enrolls ( -- table constraints
    sid    integer,
    cid    integer,
    grade  char(2),
    PRIMARY KEY(sid, cid) -- one or more PK
);
```

Enrolls

| <u>sid</u> | <u>cid</u> | grade |
|-----------------|----------------|----------------|
| 3118 | 112 | 3.8 |
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |

Enrolls

| <u>sid</u> | <u>cid</u> | grade |
|------------|------------|-------|
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |

Create/drop table

Specifying constraints

- Foreign key constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```

Create/drop table

Specifying constraints

- Foreign key constraints
 - Recap:
 - Foreign key *refers to primary key of second relation*
 - Each foreign key value in **referencing relation** must either
 - Appear as primary key value in **referenced relation**, or
 - Be a **null** value
 - Referencing & referenced relations could be the same relation
 - Also called **referential integrity constraints**
 - Keyword: FOREIGN KEY and REFERENCES

Create/drop table

Examples

```
CREATE TABLE Enrolls (  
  sid      integer REFERENCES Students(studentID),  
  cid      integer, -- assume Students and Courses  
  grade    char(2), -- table are defined  
  FOREIGN KEY(cid) REFERENCES Courses(courseID)  
);
```

Create/drop table

Examples

```
CREATE TABLE Enrolls (  
    sid      integer REFERENCES Students(studentID),  
    cid      integer, -- assume Students and Courses  
    grade    char(2), -- table are defined  
    FOREIGN KEY(cid) REFERENCES Courses(courseID)  
);
```

students

| <u>studentID</u> | .. | .. |
|------------------|----|----|
| 3118 | .. | .. |
| 1423 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |
| null | 313 | 4.5 |
| 5609 | null | 5.0 |
| null | null | 0.0 |
| 1111 | 312 | 1.2 |

Courses

| <u>courseID</u> | .. | .. |
|-----------------|----|----|
| 101 | .. | .. |
| 312 | .. | .. |
| 112 | .. | .. |

Create/drop table

Examples

```
CREATE TABLE enrolls (  
  sid      integer REFERENCES Students(studentID),  
  cid      integer, -- assume Students and Courses  
  grade    char(2), -- table are defined  
  FOREIGN KEY(cid) REFERENCES Courses(courseID)  
);
```

students

| <u>studentID</u> | .. | .. |
|------------------|----|----|
| 3118 | .. | .. |
| 1423 | .. | .. |
| 5609 | .. | .. |

enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|-----------------|----------------|----------------|
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |
| null | 313 | 4.5 |
| 5609 | null | 5.0 |
| null | null | 0.0 |
| 1111 | 312 | 1.2 |

courses

| <u>courseID</u> | .. | .. |
|-----------------|----|----|
| 101 | .. | .. |
| 312 | .. | .. |
| 112 | .. | .. |

Create/drop table

Examples

```
CREATE TABLE Enrolls (  
    sid    integer, -- assume Student_Course table  
    cid    integer, -- is defined  
    FOREIGN KEY(sid,cid) REFERENCES  
        Student_Course(studentID, courseID)  
);
```

Create/drop table

Examples

```
CREATE TABLE Enrolls (  
    sid    integer, -- assume Student_Course table  
    cid    integer, -- is defined  
    FOREIGN KEY(sid,cid) REFERENCES  
        Student_Course(studentID, courseID)  
);
```

Student_Course

| <u>studentID</u> | <u>courseID</u> |
|------------------|-----------------|
| 3118 | 112 |
| 1423 | 101 |
| 1423 | 312 |
| 5609 | 112 |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |
| null | 313 | 4.5 |
| 5609 | null | 5.0 |
| null | null | 0.0 |
| 1111 | 312 | 1.2 |

Create/drop table

Examples

```
CREATE TABLE Enrolls (  
    sid    integer, -- assume Student_Course table  
    cid    integer, -- is defined  
    FOREIGN KEY(sid,cid) REFERENCES  
        Student_Course(studentID, courseID)  
);
```

Student_Course

| <u>studentID</u> | <u>courseID</u> |
|------------------|-----------------|
| 3118 | 112 |
| 1423 | 101 |
| 1423 | 312 |
| 5609 | 112 |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|-----------------|----------------|----------------|
| 1423 | 101 | 4.0 |
| 1423 | 311 | 3.8 |
| 3118 | 112 | 4.0 |
| null | 313 | 4.5 |
| 5609 | null | 5.0 |
| null | null | 0.0 |
| 1111 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

- Deletion/update of a referenced tuple could violate a foreign key constraint
- Specify action to deal with violation as part of a foreign key constraint declaration

FOREIGN KEY(...) REFERENCES ...

[ON DELETE **action**] [ON UPDATE **action**]

Students

| <u><i>studentID</i></u> | .. | .. |
|-------------------------|----|----|
| 3118 | .. | .. |
| 1423 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| 3118 | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

[ON DELETE **action**] [ON UPDATE **action**]

Actions

- **NO ACTION** rejects delete/update if it violates constraint (*default option*)
- **RESTRICT** similar to NO ACTION except that constraint checking can't be deferred
- **CASCADE** propagate delete/update to referencing tuple
- **SET DEFAULT** updates foreign keys of referencing tuples to some *default value*
- **SET NULL** updates to null value

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

Students

| <u>studentID</u> | .. | .. |
|------------------|----|----|
| 3118 | .. | .. |
| 1423 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| 3118 | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

```
-- delete 3118
```

Students

| <u>studentID</u> | .. | .. |
|------------------|----|----|
| 3118 | .. | .. |
| 1423 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| 3118 | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

```
-- delete 3118
```

Students

| <u>studentID</u> | .. | .. |
|------------------|---------------|---------------|
| 3118 | .. | .. |
| 1423 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| 3118 | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

```
-- delete 3118
```

Students

| <u>studentID</u> | .. | .. |
|------------------|---------------|---------------|
| 3118 | .. | .. |
| 1423 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| null | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

-- delete 3118

-- change 1423 to 1444

Students

| <u>studentID</u> | .. | .. |
|------------------|----|----|
| 1423 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| null | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

-- delete 3118

-- change 1423 to 1444

Students

| <u>studentID</u> | .. | .. |
|------------------|----|----|
| 1444 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1423 | 101 | 4.0 |
| null | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Constraint checking

Foreign key constraint violation

FOREIGN KEY(...) REFERENCES ...

ON DELETE SET NULL

ON UPDATE CASCADE

-- delete 3118

-- change 1423 to 1444

Students

| <u>studentID</u> | .. | .. |
|------------------|----|----|
| 1444 | .. | .. |
| 5609 | .. | .. |

Enrolls

| <i>sid</i> | <i>cid</i> | <i>grade</i> |
|------------|------------|--------------|
| 1444 | 101 | 4.0 |
| null | 112 | 4.0 |
| 5609 | 312 | 1.2 |

Create/drop table

Specifying constraints

- Other general constraints

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
        | table_constraints }  
    [, ...]  
] )
```

Create/drop table

Specifying constraints

- Other general constraints
 - **Not-null** constraints *keyword: NOT NULL*
 - **Unique** constraints *keyword: UNIQUE*
 - Check any two records $x, y \in \text{table}$ such that
$$x.\text{column} \neq y.\text{column}$$
 - **General** constraints *keyword: CHECK*
 - ❖ Constraint is **violated** if it evaluates to FALSE
- Default values
 - Specify default values *keyword: DEFAULT*

Create/drop table

Examples: not-null

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100) NOT NULL,  
    birthDate    date  
);
```

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100),  
    birthDate    date,  
    CHECK (name IS NOT NULL)  
);
```


Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100) UNIQUE,  
    birthDate      date  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|-------|------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100) UNIQUE,  
    birthDate      date  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|------------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100) UNIQUE,  
    birthDate      date  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|------------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Check

$x, y \in \text{students}$
such that
 $x.\text{name} \neq y.\text{name}$
evaluates to FALSE

Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100),  
    birthDate      date,  
    UNIQUE (studentID, name)  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|-------|------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 1423 | Bob | 1999-06-11 |

Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100),  
    birthDate    date,  
    UNIQUE (studentID, name)  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|----------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 1423 | Bob | 1999-06-11 |

Create/drop table

Examples: unique

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100),  
    birthDate    date,  
    UNIQUE (studentID, name)  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|----------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 1423 | Bob | 1999-06-11 |

Check

$x, y \in \text{students}$
such that
 $(x.\text{studentID} \neq y.\text{studentID})$
 $\vee (x.\text{name} \neq y.\text{name})$
evaluates to FALSE

Create/drop table

Examples: general

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100),  
    birthDate      date,  
    CHECK (studentID > 2000)  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|-------|------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Examples: general

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100),  
    birthDate      date,  
    CHECK (studentID > 2000)  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|----------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Create/drop table

Examples: general

```
CREATE TABLE Students (  
    studentID      integer,  
    name           varchar(100),  
    birthDate      date,  
    CHECK (studentID > 2000)  
);
```

Students

| <u>studentID</u> | name | birthDate |
|------------------|----------------|-----------------------|
| 3118 | Alice | 1999-12-25 |
| 3119 | Alice | 1999-12-25 |
| 1423 | Bob | 2000-05-27 |
| 5609 | Carol | 1999-06-11 |

Check
 $x \in \text{students}$
such that
 $x.\text{studentID} > 2000$
evaluates to FALSE

Create/drop table

Examples: default values

```
CREATE TABLE Students (  
    studentID    integer,  
    name         varchar(100) DEFAULT 'John Doe',  
    birthDate    date  
);
```

Create/drop table

Examples: default values

```
CREATE TABLE Students (  
    studentID      integer,  
    name            varchar(100) DEFAULT 'John Doe',  
    birthDate       date  
);
```

When we insert into students table, if the value is set as DEFAULT, the default value will be used

Create/drop table

Assertions

- Complex constraints
 - Multi-table constraints
 - Example:
 - Every student in `students` table *must be enrolled in at least two and at most five courses* in `enrolls` table
- Limitation
 - Create assertion statement is not implemented in many DBMSs
 - We can use **triggers** to enforce complex constraints
 - Lecture 06