

CS2102

Database Systems

Slides adapted from Prof. Chan Chee Yong

LECTURE 04

ENTITY RELATIONSHIP DATA MODEL

Null values

Three-valued logic system: TRUE, FALSE, UNKNOWN

<i>x</i>	<i>y</i>	<i>x AND y</i>	<i>x OR y</i>	<i>NOT x</i>
FALSE FALSE FALSE	FALSE UNKNOWN TRUE	FALSE FALSE FALSE	FALSE UNKNOWN TRUE	TRUE
UNKNOWN UNKNOWN UNKNOWN	FALSE UNKNOWN TRUE	FALSE UNKNOWN UNKNOWN	UNKNOWN UNKNOWN TRUE	UNKNOWN
TRUE TRUE TRUE	FALSE UNKNOWN TRUE	FALSE UNKNOWN TRUE	TRUE TRUE TRUE	FALSE

<i>x</i>	<i>x IS NULL</i>	<i>x IS NOT NULL</i>
null	TRUE	FALSE
non-null	FALSE	TRUE

<i>x</i>	<i>y</i>	<i>x IS DISTINCT FROM y</i>
null	null	FALSE
null	non-null	TRUE
non-null	null	TRUE
non-null	non-null	<i>x <> y</i>

Data definition language (DDL)

Create table

```
CREATE TABLE [ IF NOT EXISTS ] table_name ( [  
    { column_name data_type  
        [ column_constraints [ ... ] ]  
    | table_constraints }  
    [, ...]  
] );
```

Drop table syntax

```
DROP TABLE [ IF EXISTS ] table_name
```

Data definition language (DDL)

Alter table

- Add/remove/modify columns
 - `ALTER TABLE students ALTER COLUMN dept DROP DEFAULT;`
 - `ALTER TABLE students DROP COLUMN dept;`
 - `ALTER TABLE students ADD COLUMN faculty varchar(20);`
 - etc
- Add/remove constraints
- etc

Data manipulation language (DML)

Insert into syntax

```
INSERT INTO table_name  
[ ( column_name [, ...] ) ]  
VALUES ( { expression | DEFAULT } [, ...] );
```

Delete from syntax

```
DELETE FROM table_name  
[ WHERE condition ];
```

Update syntax

```
UPDATE table_name  
SET column_name = { expression | DEFAULT }  
[ WHERE condition ];
```

Simple queries

Basic syntax

- Basic form of **SQL query** consists of three clauses

```
SELECT [ DISTINCT ] select_list -- select clause
FROM          from_list      -- from clause
[ WHERE       condition ] -- where clause
```

- **select_list** specifies columns to be included in output
 - **from_list** specifies list of relations
 - **condition** specifies conditions on relations
-
- ❖ **Output:** relation generated from `from_list` containing attributes based on `select_list` that satisfies `condition`
 - Output relation could contain duplicate record if `DISTINCT` is not used in the `SELECT` clause

- Entity Relationship Diagram

- ER model

- Relationship constraints

- Participation constraints

- Weak entity sets

- ER to SQL

- ER diagram to SQL

- Additional ER concepts

- ER design and relational mapping

Overview

- Entity Relationship Diagram

- ER model

- Relationship constraints

- Participation constraints

- Weak entity sets

- ER to SQL

- ER diagram to SQL

- Additional ER concepts

- ER design and relational mapping

Entity Relationship Diagram

Conceptual data models

Introduction

- **Entity-relationship (ER) model**
 - Developed by Peter Chen in 1976
 - Designed for conceptual data model specifications
 - Most common data model used for database design
- **Unified modelling language (UML)**
 - Developed by Grady Booch & James Rumbaugh in 1997
 - Goes beyond conceptual data modelling
 - Software design specifications
 - Standardized by Object Management Group (OMG)

Database design process

Steps

1. Requirement analysis

- Find out the data/application/performance requirement of the enterprise

2. Conceptual database design

- Capture data requirements using a conceptual schema

3. Logical database design

- Map conceptual schema to logical schema supported by DBMS

4. Schema refinement

- Improve logical schema design using data constraints

5. Physical database design

- Use performance requirements to design physical schema

6. Application & security design

- Specify access control policies
- 

Database design process

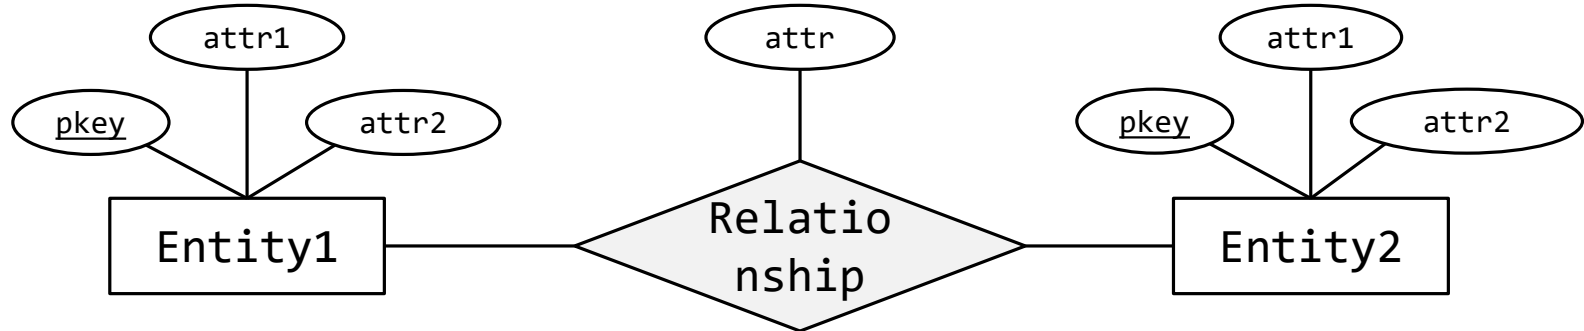
Requirement analysis

I would like my customers to be able to browse my catalog of books and place orders over the Internet. Currently, I take orders over the phone. I have mostly corporate customers who call me and give me the ISBN number of a book and a quantity; they often pay by credit card. If I don't have enough copies in stock, I order additional copies and delay the shipment until the new copies arrive; I want to ship a customer's entire order together. My catalog includes all the books I sell. For each book, the catalog contains its ISBN number, title, author, purchase price, sales price, and the year the book was published. Most of my customers are regulars, and I have records with their names and addresses. New customers have to call me first and establish an account before they can use my website. On my new website, customers should first identify themselves by their unique customer identification number. Then they should be able to browse my catalog and to place orders online.

ER model

Introduction

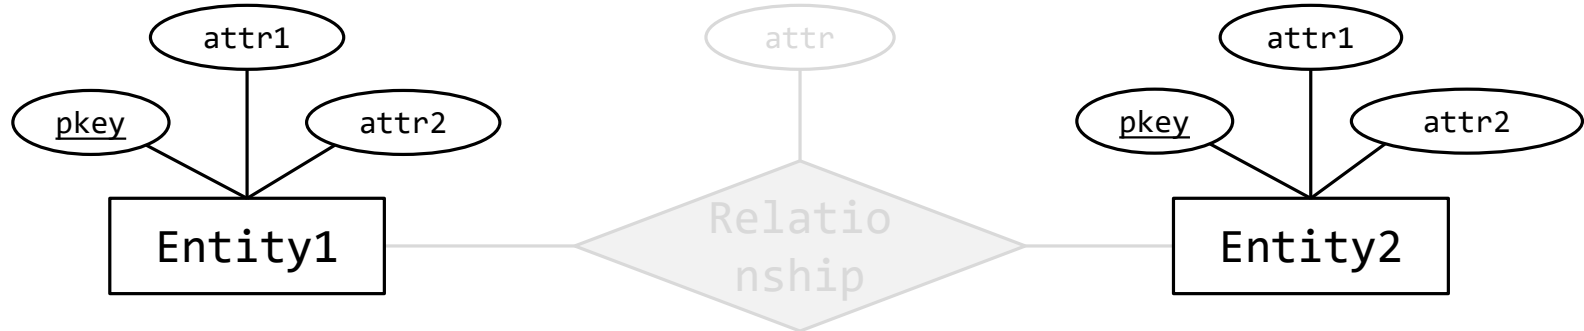
- Data is described in terms of **entities** and their **relationships**
- Information about entities and relationships are described using **attributes**
- Certain data constraints are represented using additional annotations
- ER schemas are presented as **ER diagrams**



ER model

Basic

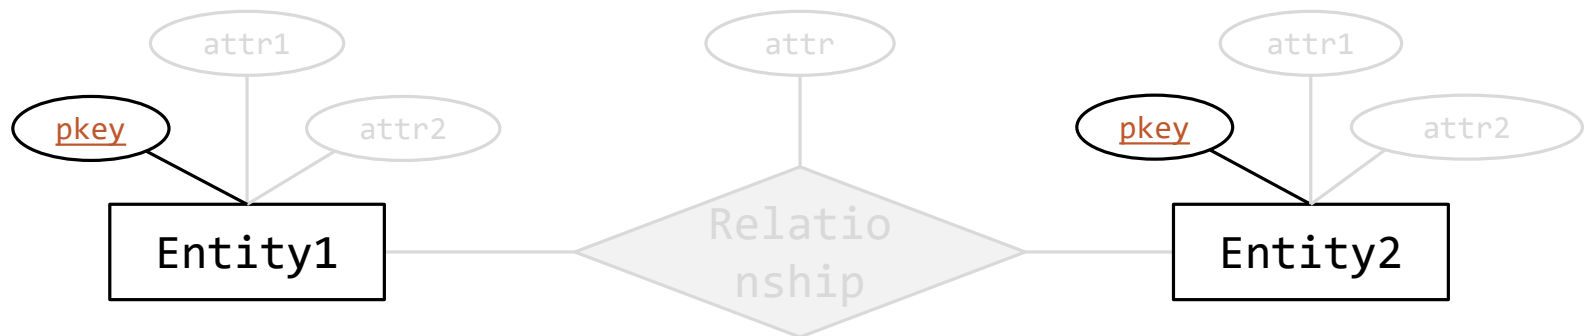
- **Entity** *Real-world object distinguishable from other objects*
- **Attribute** *Specific information describing an entity*
each attribute has an atomic domain (e.g, integer, string)
 - Represented by *ovals*
- **Entity set** *A collection of similar entities*
 - Represented by *rectangles*



ER model

Keys


- Each entity set has a **key** (*i.e., minimal set of attributes whose values uniquely identify an entity*)
- An entity set could have multiple keys called **candidate keys**
- One of the candidate keys is chosen as the **primary key**
- The attributes that formed a primary key are underlined
- Example: **sid** in Students and **cid** in Courses



ER model

Relationships

- **Relationship** *An association among two or more entities attributes are used to describe information about relationships*
- **Relationship set** *A collection of similar relationships*
 - Represented by *diamonds*



ER model


Entities, relationships, and attributes

I would like my customers to be able to browse my catalog of books and place orders over the Internet. Currently, I take orders over the phone. I have mostly corporate customers who call me and give me the ISBN number of a book and a quantity; they often pay by credit card. If I don't have enough copies in stock, I order additional copies and delay the shipment until the new copies arrive; I want to ship a customer's entire order together. My catalog includes all the books I sell. For each book, the catalog contains its ISBN number, title, author, purchase price, sales price, and the year the book was published. Most of my customers are regulars, and I have records with their names and addresses. New customers have to call me first and establish an account before they can use my website. On my new website, customers should first identify themselves by their unique customer identification number. Then they should be able to browse my catalog and to place orders online.

Relationship constraints

Many-to-many relationship sets

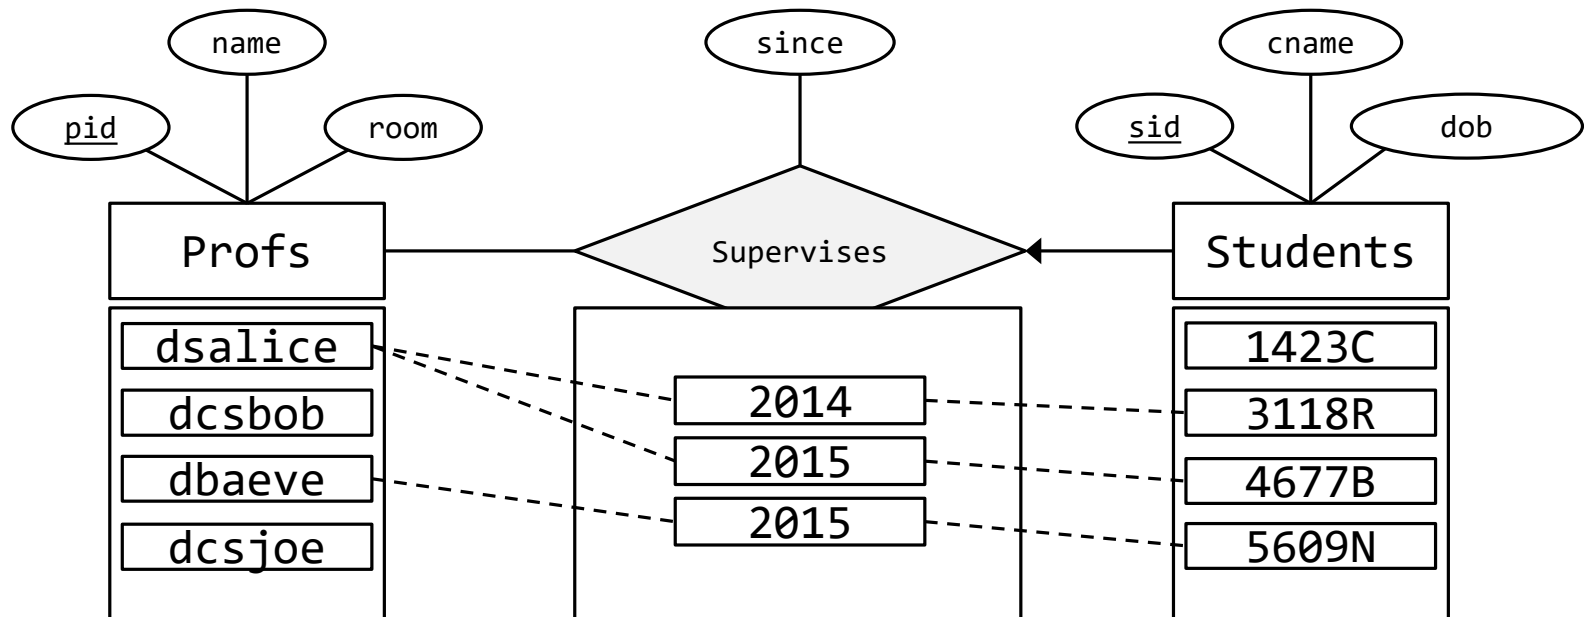
- Example
 - Each student can enroll in 0 or more courses
 - Each course can be enrolled by 0 or more students



Relationship constraints

Key constraints: one-to-many relationships

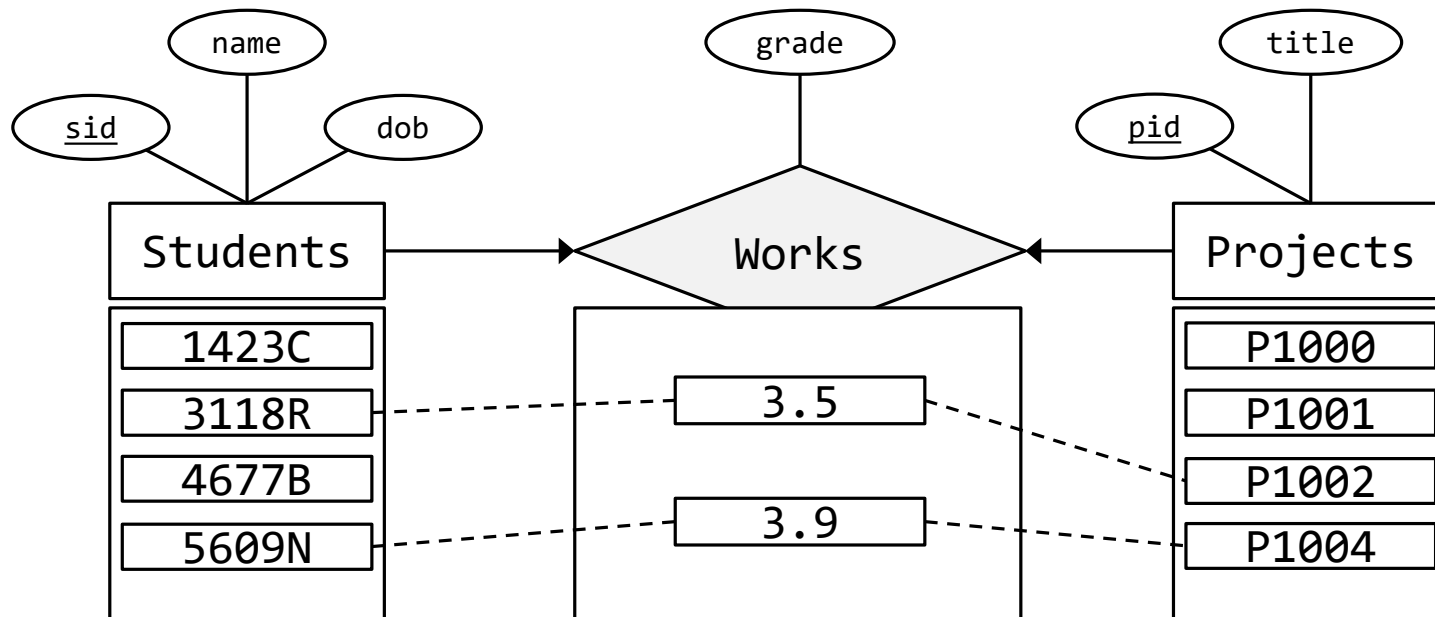
- Let R be a relationship set that involves entity set E
- Key constraint** on E w.r.t. R
 - Each instance E can participate in at most one instance of R
- Example:
 - Each professor can supervise 0 or more students
 - Each student can be supervised by at most one professor



Relationship constraints

Key constraints: one-to-one relationships



- Example:
 - Each student can work on at most one project
 - Each project can be worked on by at most one student



Relationship constraints

Key constraints: N-ary relationships


- Example:
 - Each professor can work in at most one department located at some building




Participation constraints

Introduction

- **Participation constraints:** is the participation of an entity set in a relationship set mandatory?
- **Partial participation constraint**
 - Example: each student can enroll in 0 or more courses



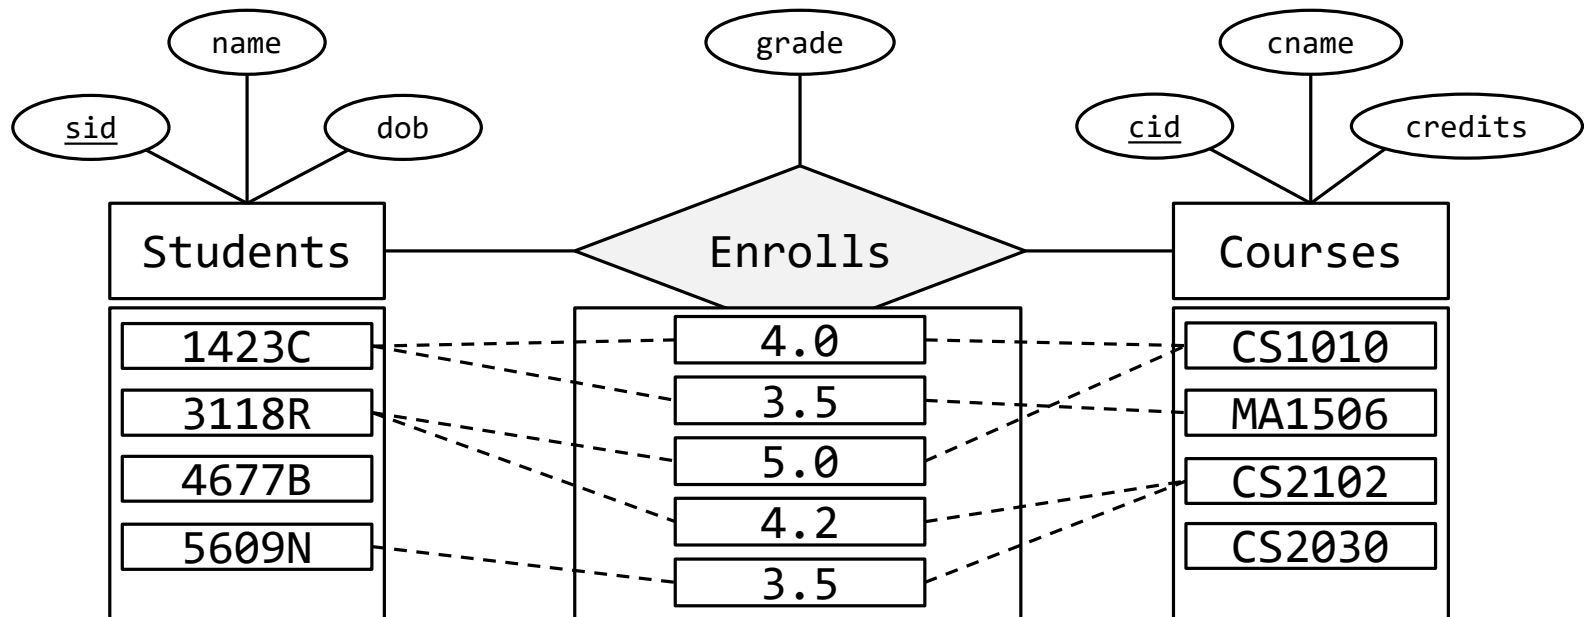
- **Total participation constraint**
 - Example: each student must enroll in at least one course



Participation constraints

Partial participation constraints

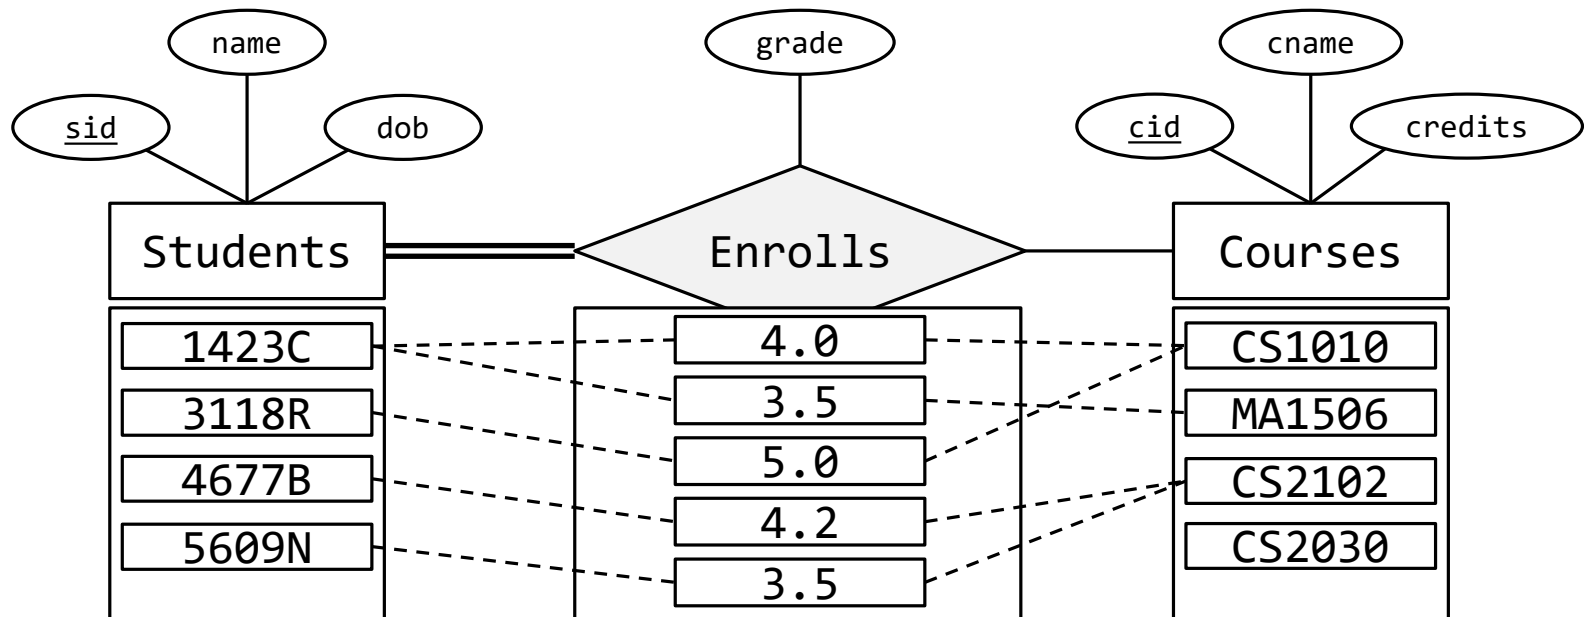
- Example
 - Each student can enroll in 0 or more courses
 - Each course can be enrolled by 0 or more students



Participation constraints

Total participation constraints

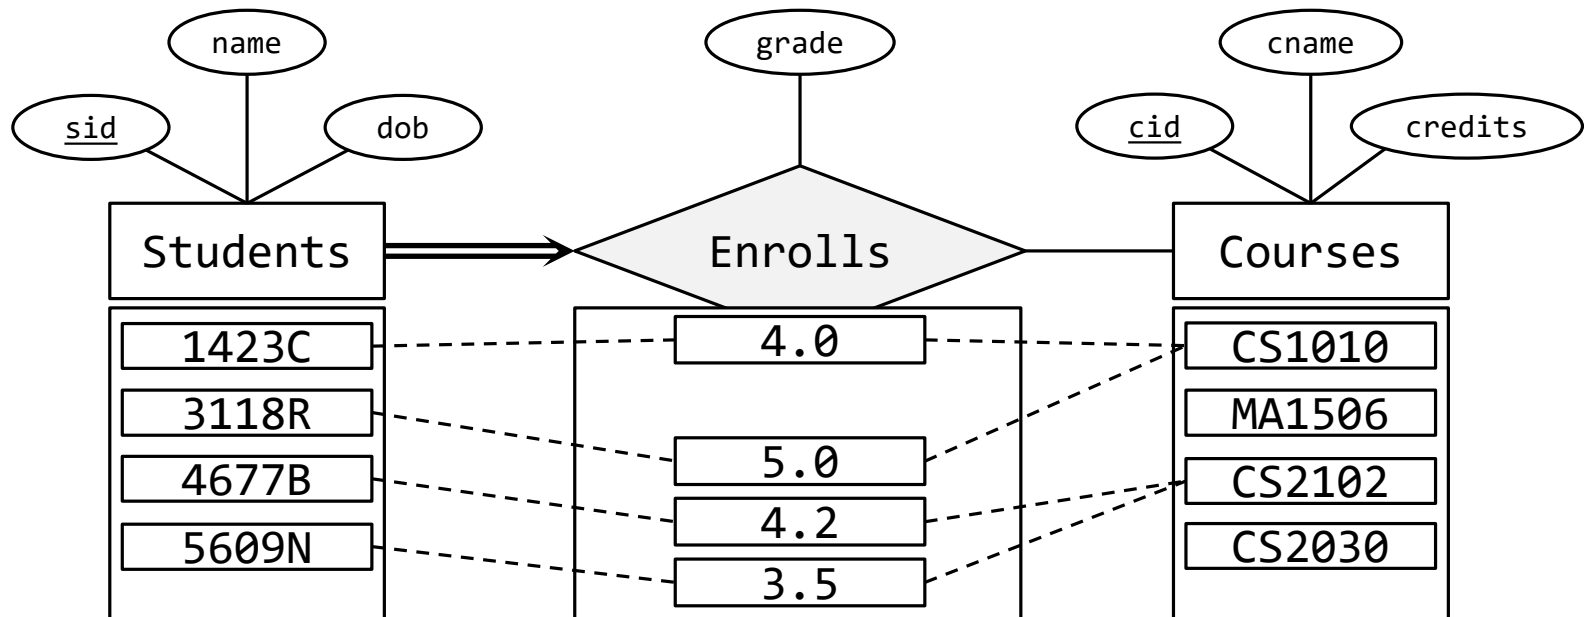
- Example
 - Each student must enroll in at least one courses
 - Each course can be enrolled by 0 or more students



Participation constraints

Key & Total participation constraints


- Example
 - Each student must enroll in at most one & at least one courses
 - Each course can be enrolled by 0 or more students



Roles in relationship

Roles


- **Roles** are used when one entity set appears two or more times in a relationship set



Weak entity sets

Definition


- **Weak entity set** is an entity set that does not have its own key
- A weak entity can only be uniquely identified by considering the primary key of another entity (called **identifying owner**)
- There must be a many-to-one relationship (called **identifying relationship**) from the weak entity set to an owner entity set
- Weak entity set must have total participation in identifying relationship



Weak entity sets

Properties


- **Partial key** of a weak entity set is a set of attributes of weak entity set that uniquely identifies a weak entity *for a given owner entity*
- A weak entity's existence is dependent on the existence of its owner entity
- Weak entity sets are represented by double-lined rectangles
- Identifying relationship sets are represented by double-lined diamonds



Summary

Relationship constraints

- **Many-to-many** Each instance of E participates in 0 or more instance of R
- **Key** Each instance of E participates in at most 1 instance of R
- **Total** Each instance of E participates in at least 1 instance of R
- **Key & total** Each instance of E participates in exactly one instance of R
- **Weak entity** E is a weak entity set with identifying owner E' and identifying relationship set R



- Entity Relationship Diagram

 - ER model

 - Relationship constraints

 - Participation constraints

 - Weak entity sets

- ER to SQL

 - ER diagram to SQL

 - Additional ER concepts


 - ER design and relational mapping

ER to SQL


ER diagram to SQL

Entity sets

```
CREATE TABLE Students (  
    sid      integer,  
    name     varchar(30),  
    dob      date,  
    PRIMARY KEY(sid)  
);
```




```
CREATE TABLE Offices (  
    building    char(10),  
    room        integer,  
    area        varchar(20),  
    PRIMARY KEY(building,room)  
);
```



ER diagram to SQL


Entity sets with candidate keys



```
CREATE TABLE Offices (  
    officeID      char(10) PRIMARY KEY,  
    building      char(10) NOT NULL,  
    room          integer  NOT NULL,  
    area          varchar(20),  
    UNIQUE(building,room)  
);
```

ER diagram to SQL


Relationship sets without constraints



```
CREATE TABLE Enrolls (  
    sid      integer REFERENCES Students,  
    cid      char(5) REFERENCES Parts,  
    qty      numeric,  
    PRIMARY KEY(sid,cid)  
);
```


ER diagram to SQL


Relationship sets without constraints



```
CREATE TABLE Supplies (  
    sid      char(10) REFERENCES Suppliers,  
    pid      char(10) REFERENCES Parts,  
    pjid     char(10) REFERENCES Projects,  
    price    numeric,  
    qty      integer,  
    PRIMARY KEY(sid,pid,pjid)  
);
```

ER diagram to SQL

Relationship sets with key constraints



First approach

Represent Supervises with a separate table

- Profs (pid, name, room)
- Students (sid, name, dob)
- Supervises (sid, pid, since)


Second approach

Combine Supervises & Students into a single table

- Profs (pid, name, room)
- SupervisedStudents (sid, name, dob, pid, since)

ER diagram to SQL

Relationship sets with key constraints



First approach


```
CREATE TABLE Supervises (  
    sid    integer,  
    pid    char(7),  
    since  date,  
    PRIMARY KEY(sid),  
    FOREIGN KEY(sid) REFERENCES  
        Students,  
    FOREIGN KEY(pid) REFERENCES  
        Profs  
);
```

Second approach

```
CREATE TABLE SupervisedStudents (  
    sid    integer,  
    name   varchar(30),  
    dob    date,  
    pid    char(7),  
    since  date,  
    PRIMARY KEY(sid),  
    FOREIGN KEY(pid) REFERENCES  
        Profs  
);
```

ER diagram to SQL

Relationship sets with key & total constraints



First approach

```
CREATE TABLE Supervises (  
  sid integer,  
  pid char(7),  
  since date,  
  PRIMARY KEY(sid),  
  FOREIGN KEY(sid) REFERENCES  
    Students,  
  FOREIGN KEY(pid) REFERENCES  
    Profs  
);
```


Total participation constraint of
Students w.r.t Supervises is not
captured!

Second approach

```
CREATE TABLE SupervisedStudents (  
  sid integer,  
  name varchar(30),  
  dob date,  
  pid char(7) NOT NULL,  
  since date,  
  PRIMARY KEY(sid),  
  FOREIGN KEY(pid) REFERENCES  
    Profs  
);
```

ER diagram to SQL


Roles in relationships



```
CREATE TABLE Mentors (  
    seniorSID integer,  
    juniorSID integer,  
    PRIMARY KEY(seniorSID, juniorSID),  
    FOREIGN KEY(seniorSID) REFERENCES Students(sid),  
    FOREIGN KEY(juniorSID) REFERENCES Students(sid)  
);
```

ER diagram to SQL

Weak entity sets



Books table

```
CREATE TABLE Books (  
  isbn    char(30),  
  title   char(50),  
  author  char(60),  
  PRIMARY KEY(isbn)  
);
```

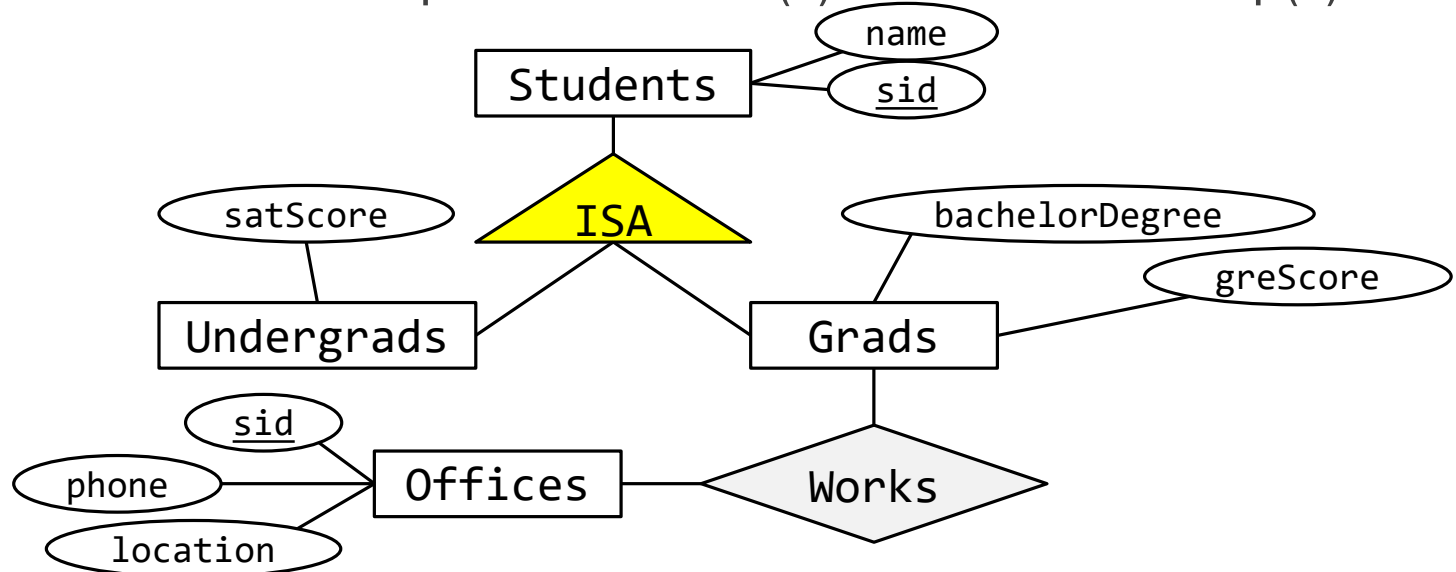
BookChapters table

```
CREATE TABLE BookChapters (  
  num      char(30),  
  title    char(50),  
  isbn     char(60),  
  PRIMARY KEY(num,isbn),  
  FOREIGN KEY(isbn)  
    REFERENCES Books  
    ON DELETE cascade  
);
```

Additional ER concepts

ISA hierarchies

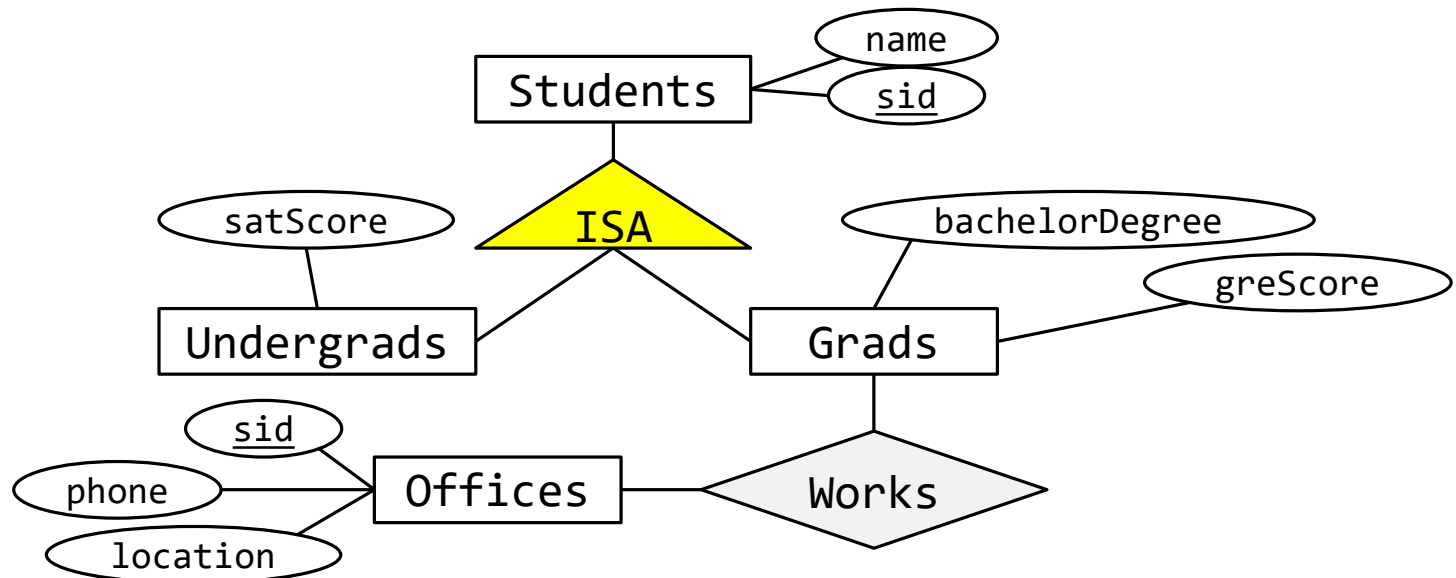
- Based on “is-a” relationship in OOP
 - Subclass-superclass relationship
 - Describing an entity sets into subclasses
- Every entity in a subclass entity set is an entity in its superclass entity set
- Each subclass has specific attribute(s) and/or relationship(s)



Additional ER concepts

ISA hierarchies

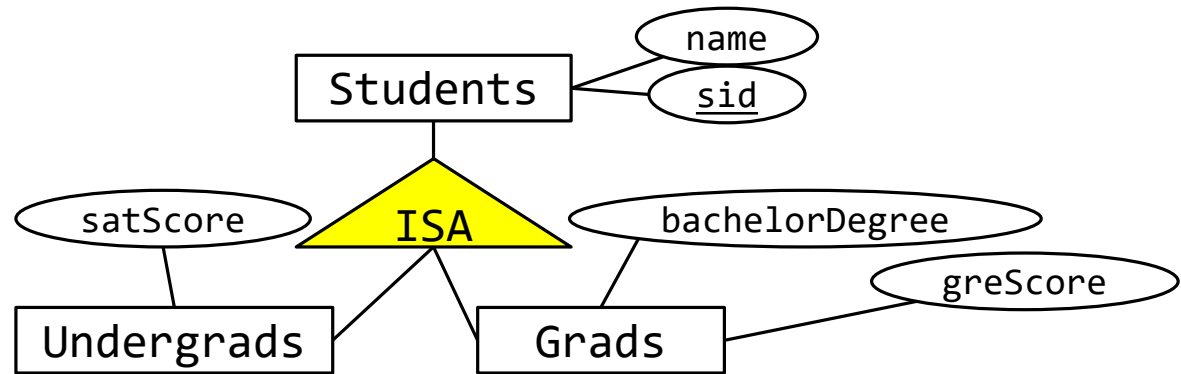
- Constraints:
 - **Overlap constraints:** can entity belong to multiple subclasses?
 - Satisfied if entity in superclass could belong to multiple subclasses
 - **Covering constraints:** does an entity in a superclass have to belong to some subclass?
 - Satisfied if every entity in a superclass has to belong to some subclass



Additional ER concepts

ISA hierarchies

- **Approach #1:**
one relation
per subclass
or superclass



```
CREATE TABLE Students (  
    sid      integer PRIMARY KEY,  
    name     char(30) );
```


```
CREATE TABLE Undergrads (  
    sid      integer PRIMARY KEY REFERENCES Students  
            ON DELETE cascade,  
    satScore numeric );
```

```
CREATE TABLE Grads (  
    sid      integer PRIMARY KEY REFERENCES Students  
            ON DELETE cascade,  
    greScore numeric );
```

Additional ER concepts

ISA hierarchies

- **Approach #2:**
one relation
per subclass



Applicable if covering
constraint is satisfied

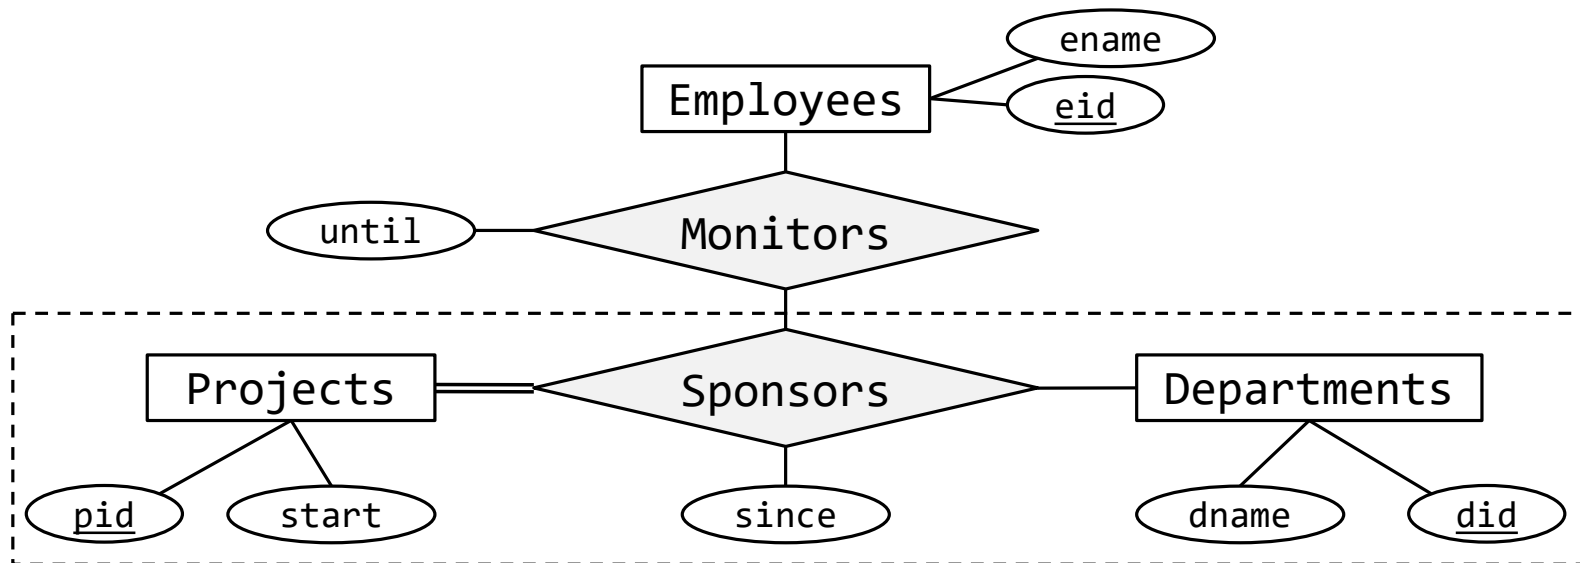
```
CREATE TABLE Undergrads (  
    sid          integer PRIMARY KEY,  
  
    satScore     numeric );
```

```
CREATE TABLE Grads (  
    sid          integer PRIMARY KEY,  
  
    greScore     numeric );
```

Additional ER concepts

Aggregation

- How to model a relationship between entities & relationships?
- Example:
 - Every project is sponsored by at least one department
 - Each sponsorship has a “since” attribute & might be monitored by 0 or more employees
 - Each monitoring has an “until” attribute



Additional ER concepts

Aggregation

- Relational mapping

```
CREATE TABLE Projects (  
  pid      char(20),  
  start    date,  
  PRIMARY KEY(pid)  
);
```

```
CREATE TABLE Departments (  
  did      char(20),  
  dname    char(30),  
  PRIMARY KEY(did)  
);
```

```
CREATE TABLE Sponsors (  
  pid      char(20)  
    REFERENCES Projects,  
  did      char(30),  
    REFERENCES Departments,  
  since    date,  
  PRIMARY KEY(pid,did)  
);
```

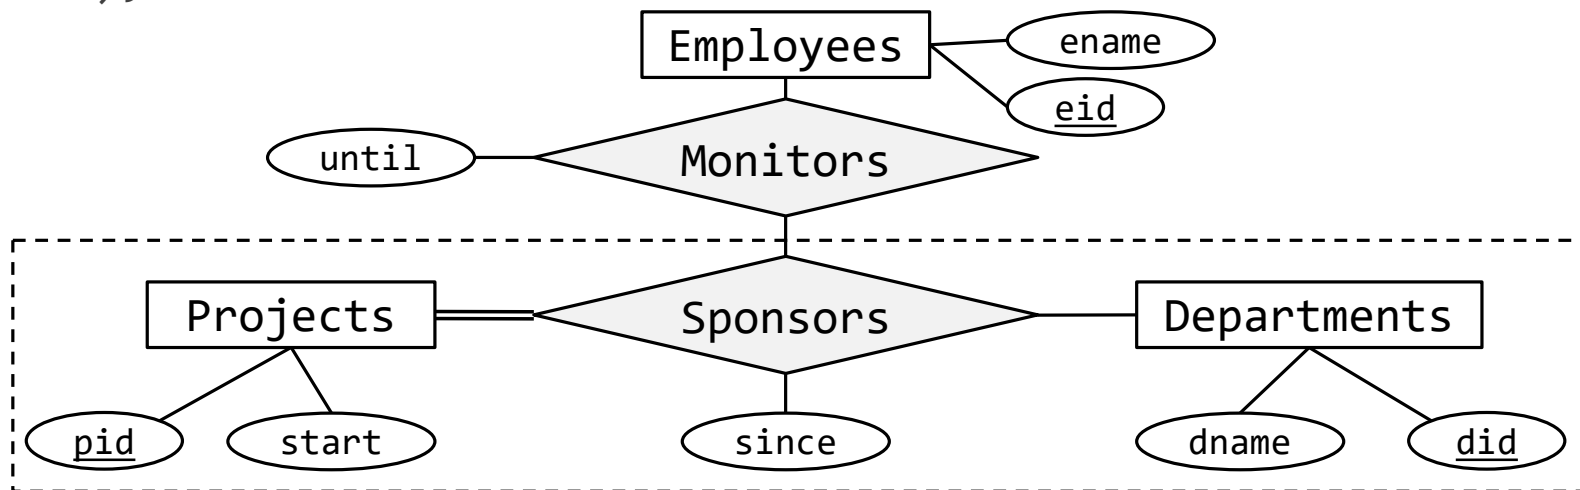
```
CREATE TABLE Employees (  
  eid      char(20),  
  ename    char(30),  
  PRIMARY KEY(eid)  
);
```

Additional ER concepts

Aggregation

- Relational mapping

```
CREATE TABLE Monitors (  
    eid      char(20) REFERENCES Employees,  
    pid      char(30),  
    did      char(30),  
    until    date,  
    PRIMARY KEY(eid,pid,did),  
    FOREIGN KEY(pid,did) REFERENCES Sponsors(pid,did)  
);
```




ER design and relational mapping

Guidelines for ER design

- ER design should capture as many of the application's constraints as possible
- ER design must not impose any constraint that is not required in the application

Guidelines for relational mapping

- Relational schema should enforce as many of the application's constraints as possible using column/table constraints
 - Relational schema must not impose any constraint that is not required in the application
- 

Summary

- ❑ ER model has expressive constructs for conceptual data design
 - ❑ Concepts: *entities; relationships; attributes; weak entities; ISA hierarchies; aggregation*
 - ❑ Constraints: *key constraints; participation constraints*
- ❑ ER design is subjective
- ❑ Rules for mapping entity-relationship model to relational model
 - ❑ Entity & relationship sets
 - ❑ Key constraints
 - ❑ Participation constraints
 - ❑ Relationship roles
 - ❑ Weak entity sets
 - ❑ ISA hierarchies
 - ❑ Aggregation