

# CS2102

# Database Systems

*Slides adapted from Prof. Chan Chee Yong*

---

LECTURE 06


SQL #3

# Relationship constraints


## Types

- Many-to-many
- Key
- Total
- Key & total
- Weak entity


Each instance of E participates in 0 or more instance of R




Each instance of E participates in at most 1 instance of R




Each instance of E participates in at least 1 instance of R



Each instance of E participates in exactly one instance of R




E is a weak entity set with identifying owner E' and identifying relationship set R



# ER diagram to SQL


## Entity sets

```
CREATE TABLE E (  
    pk1    type,  
    pk2    type,  
    attr   type,  
    PRIMARY KEY (pk1, pk2)  
);
```



## Many-to-many


```
CREATE TABLE R (  
    pk1    type REFERENCES E1,  
    pk2    type REFERENCES E2,  
    attr   type,  
    PRIMARY KEY (pk1, pk2)  
);
```



# ER diagram to SQL


## Key constraints approach #1

```
CREATE TABLE R (  
  pk1    type REFERENCES E1,  
  pk2    type REFERENCES E2,  
  attr   type,  
  PRIMARY KEY (pk1)  
);
```



## Key constraints approach #2


```
CREATE TABLE R (  
  pk1    type PRIMARY KEY,  
  pk2    type REFERENCES E2,  
  attr   type,  
  attr1  type  
);
```



# ER diagram to SQL


## Key & total constraints

```
CREATE TABLE R (
  pk1    type PRIMARY KEY,
  pk2    type NOT NULL,
  attr   type,
  attr1  type,
  FOREIGN KEY (pk2) REFERENCES E2
);
```



## Weak entity sets

```
CREATE TABLE R (
  pk1    type,
  pk2    type REFERENCES E2
         ON DELETE cascade,
  attr   type,
  attr1  type,
  PRIMARY KEY (pk1, pk2)
);
```




# Additional ER concepts

## ISA hierarchies approach #1

```
CREATE TABLE E (  
  pk      type PRIMARY KEY  
);
```

```
CREATE TABLE E1 (  
  pk      type  
          PRIMARY KEY  
          REFERENCES E ON DELETE cascade  
);
```

```
CREATE TABLE E2 (  
  pk      type  
          PRIMARY KEY  
          REFERENCES E ON DELETE cascade  
);
```



# Additional ER concepts


## ISA hierarchies approach #2

```
CREATE TABLE E1 (  
  pk      type  
          PRIMARY KEY
```

```
);
```

```
CREATE TABLE E2 (  
  pk      type  
          PRIMARY KEY
```

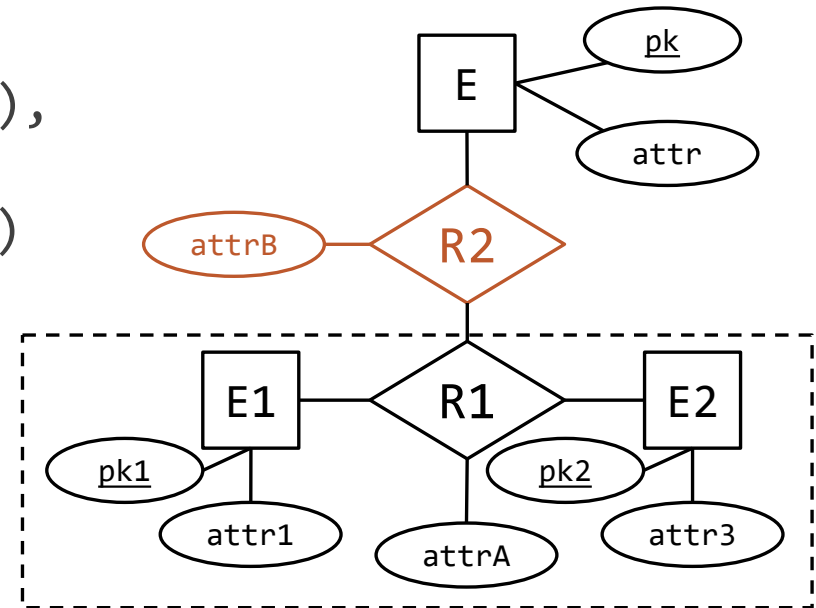
```
);
```



# Additional ER concepts

## Aggregation

```
CREATE TABLE R2 (  
  pk      type REFERENCES E,  
  pk1     type,  
  pk2     type,  
  attrB   type,  
  PRIMARY KEY (pk, pk1, pk2),  
  FOREIGN KEY (pk1, pk2)  
    REFERENCES R1 (pk1, pk2)  
);
```





- **Aggregate functions**

- Introduction

- Order by clause

- Group by clause

- Having clause

- **Subqueries**

- select clause

- from clause

- where clause

- Queries with universal quantification

---

## Overview

- **Aggregate functions**

- Introduction

- Order by clause

- Group by clause

- Having clause

- **Subqueries**

- select clause

- from clause

- where clause

- Queries with universal quantification

---

# Aggregate functions

---

# Aggregate functions

## Introduction

- **Aggregate function** computes a single value from a set of tuples
- Types of aggregate functions:
  - MIN, MAX, AVG, COUNT, SUM
- Aggregate function can be used in different parts of SQL queries
  - select clause
  - having clause (*to be discussed later*)
  - order by clause (*to be discussed later*)

# Aggregate functions

Query	Meaning
SELECT MIN(A) FROM R	Minimum values in A
SELECT MAX(A) FROM R	Maximum values in A
SELECT AVG(A) FROM R	Average values in A
SELECT SUM(A) FROM R	Sum of values in A
SELECT COUNT(A) FROM R	Count number of non-null values in A
SELECT COUNT(*) FROM R	Count number of rows in R
SELECT AVG(DISTINCT A) FROM R	Average of distinct values in A
SELECT SUM(DISTINCT A) FROM R	Sum of distinct values in A
SELECT COUNT(DISTINCT A) FROM R	Count number of distinct non-null values in A

Query	Empty relation?	All null values with cardinality n?
SELECT MIN(A) FROM R	null	null
SELECT MAX(A) FROM R	null	null
SELECT AVG(A) FROM R	null	null
SELECT SUM(A) FROM R	null	null
SELECT COUNT(A) FROM R	0	0
SELECT COUNT(*) FROM R	0	n

---

# Aggregate functions

## Example

- Find the number, minimum price, maximum price, and average price of pizzas sold by Corleone Corner

# Aggregate functions

## Example

- Find the number, minimum price, maximum price, and average price of pizzas sold by Corleone Corner

```
SELECT COUNT(price), MIN(price),  
       MAX(price)   , AVG(price)  
FROM   Sells WHERE rname = 'Corleone Corner';
```

# Aggregate functions

## Example

- Find the number, minimum price, maximum price, and average price of pizzas sold by Corleone Corner

```
SELECT COUNT(price), MIN(price),  
       MAX(price) , AVG(price)  
FROM   Sells WHERE rname = 'Corleone Corner';
```

Sells

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
:	:	:

Output

count	min	max	avg
3	19	25	22.6667

# Order by clause

## Introduction

- **Order by clause** sort the result of SQL query
- Syntax
  - **ORDER BY** **column1** [ ASC | DESC ]  
[ , **column2** [ ASC | DESC ] [ ... ] ]
- Example:
  - Find all customer name and the pizza the customer likes. Show the output in ascending order of the pizza, followed by in descending order of the customer name

```
SELECT    cname, pizza
FROM      Likes
;
```

Likes

<u>cname</u>	<u>pizza</u>
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Siciliana
Ralph	Diavola

Output

<u>cname</u>	<u>pizza</u>
Ralph	Diavola
Moe	Funghi
Maggie	Funghi
Lisa	Funghi
Homer	Hawaiian
Homer	Margherita
Moe	Siciliana



# Order by clause

## Introduction

- **Order by clause** sort the result of SQL query
- Syntax
  - **ORDER BY** **column1** [ ASC | DESC ]  
[ , **column2** [ ASC | DESC ] [ ... ] ]
- Example:
  - Find all customer name and the pizza the customer likes. Show the output in **ascending order of the pizza, followed by in descending order of the customer name**

```
SELECT    cname, pizza
FROM      Likes
ORDER BY pizza ASC,
         cname DESC;
```

Likes

<u>cname</u>	<u>pizza</u>
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Siciliana
Ralph	Diavola

Output

<u>cname</u>	<u>pizza</u>
Ralph	Diavola
Moe	Funghi
Maggie	Funghi
Lisa	Funghi
Homer	Hawaiian
Homer	Margherita
Moe	Siciliana

# Order by clause

## Limit and offset

- **Limit** and **offset** allows retrieval of only a portion of rows
- Syntax
  - `LIMIT { number | ALL }`
  - `OFFSET number`
- Example:
  - Find the top 3 most expensive pizzas. Show the pizza name, the name of the restaurant that sells it, and its selling price for each output tuple; and show the output in descending order of price

```
SELECT    pizza, rname, price
FROM      Sells
ORDER BY  price DESC
          ;
```

**Output**

<i><b>pizza</b></i>	<i><b>rname</b></i>	<i><b>price</b></i>
Hawaiian	Corleone Corner	25
Diavola	Corleone Corner	24
Funghi	Lorenzo Tavern	23

# Order by clause

## Limit and offset

- **Limit** and **offset** allows retrieval of only a portion of rows
- Syntax
  - `LIMIT { number | ALL }`
  - `OFFSET number`
- Example:
  - Find the **top 3** most expensive pizzas. Show the pizza name, the name of the restaurant that sells it, and its selling price for each output tuple; and show the output in descending order of price

```
SELECT    pizza, rname, price
FROM      Sells
ORDER BY  price DESC
LIMIT    3;
```

**Output**

<i><b>pizza</b></i>	<i><b>rname</b></i>	<i><b>price</b></i>
Hawaiian	Corleone Corner	25
Diavola	Corleone Corner	24
Funghi	Lorenzo Tavern	23

# Order by clause

## Limit and offset

- **Limit** and **offset** allows retrieval of only a portion of rows
- Syntax
  - `LIMIT { number | ALL }`
  - `OFFSET number`
- Example:
  - For each pizza that is sold by some restaurant, find the pizza name, the restaurant name, and its selling price; show the output in descending order of price and exclude the top 3 pizzas

```
SELECT    pizza, rname, price
FROM      Sells
ORDER BY  price DESC
          ;
```

### Output

<i>pizza</i>	<i>rname</i>	<i>price</i>
Marinara	Mamma's Place	23
Hawaiian	Pizza King	21
Margherita	Corleone Corner	19
Diavola	Pizza King	17
Sciliana	Gambino Oven	16

# Order by clause

## Limit and offset

- **Limit** and **offset** allows retrieval of only a portion of rows
- Syntax
  - LIMIT { **number** | ALL }
  - OFFSET **number**
- Example:
  - For each pizza that is sold by some restaurant, find the pizza name, the restaurant name, and its selling price; show the output in descending order of price and **exclude the top 3 pizzas**

```
SELECT    pizza, rname, price
FROM      Sells
ORDER BY  price DESC
OFFSET   3;
```

### Output

<i><b>pizza</b></i>	<i><b>rname</b></i>	<i><b>price</b></i>
Marinara	Mamma's Place	23
Hawaiian	Pizza King	21
Margherita	Corleone Corner	19
Diavola	Pizza King	17
Sciliana	Gambino Oven	16

# Order by clause

## Limit and offset

- **Limit** and **offset** allows retrieval of only a portion of rows
- Syntax
  - `LIMIT { number | ALL }`
  - `OFFSET number`
- Example:
  - Find the 4<sup>th</sup> and 5<sup>th</sup> most expensive pizzas. Show the pizza name, the name of the restaurant the sells it, and its selling price; show the output in descending order of price

```
SELECT    pizza, rname, price
FROM      Sells
ORDER BY  price DESC
```

;

### Output

<i>pizza</i>	<i>rname</i>	<i>price</i>
Marinara	Mamma's Place	23
Hawaiian	Pizza King	21

# Order by clause

## Limit and offset

- **Limit** and **offset** allows retrieval of only a portion of rows
- Syntax
  - `LIMIT { number | ALL }`
  - `OFFSET number`
- Example:
  - Find the 4<sup>th</sup> and 5<sup>th</sup> most expensive pizzas. Show the pizza name, the name of the restaurant the sells it, and its selling price; show the output in descending order of price

```
SELECT    pizza, rname, price
FROM      Sells
ORDER BY  price DESC
LIMIT    2
OFFSET    3;
```

### Output

<i>pizza</i>	<i>rname</i>	<i>price</i>
Marinara	Mamma's Place	23
Hawaiian	Pizza King	21

# Group by clause

## Introduction

- **Group by clause** divides the rows into groups such that aggregate functions can be applied to each group
- Syntax
  - GROUP BY **column1** [ , **column2** [ ... ] ]
- Example:
  - For each restaurant that sells some pizza, find the minimum and maximum prices of its pizzas

**Sells**

<u><i>rname</i></u>	<u><i>pizza</i></u>	<i>price</i>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21



# Group by clause

## Introduction

- Example:
  - For each restaurant that sells some pizza, find the minimum and maximum prices of its pizzas
- Conceptual steps:
  1. Partition the tuples in `Sells` into groups based on `rname`

**Sells**

<u><i>rname</i></u>	<u><i>pizza</i></u>	<i>price</i>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

# Group by clause

## Introduction

- Example:
  - For each restaurant that sells some pizza, find the minimum and maximum prices of its pizzas
- Conceptual steps:
  1. Partition the tuples in `Sells` into groups based on `rname`
  2. Compute `MIN(price)` and `MAX(price)` for each group

**Sells**

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

# Group by clause

## Introduction

- Example:
  - For each restaurant that sells some pizza, find the minimum and maximum prices of its pizzas
- Conceptual steps:
  1. Partition the tuples in `Sells` into groups based on `rname`
  2. Compute `MIN(price)` and `MAX(price)` for each group
  3. Output one tuple for each group

**Sells**

<u><i>rname</i></u>	<u><i>pizza</i></u>	<i>price</i>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

**Output**

<i>rname</i>	<i>min</i>	<i>max</i>
Corleone Corner	19	25
Gambino Oven	16	16
Lorenzo Tavern	23	23
Mamma's Place	22	22
Pizza King	17	21

# Group by clause

## Introduction

- Example:
  - For each restaurant that sells some pizza, find the minimum and maximum prices of its pizzas
- Conceptual steps:
  1. Partition the tuples in `Sells` into groups based on `rname`
  2. Compute `MIN(price)` and `MAX(price)` for each group
  3. Output one tuple for each group

```
SELECT rname, MIN(price), MAX(price)
FROM   Sells
WHERE  rname = 'Corleone Corner'
UNION
SELECT rname, MIN(price), MAX(price)
FROM   Sells
WHERE  rname = 'Gambino Oven'
UNION  ...
;
```

Output

<i>rname</i>	<i>min</i>	<i>max</i>
Corleone Corner	19	25
Gambino Oven	16	16
Lorenzo Tavern	23	23
Mamma's Place	22	22
Pizza King	17	21

# Group by clause

## Introduction

- Example:
  - For each restaurant that sells some pizza, find the minimum and maximum prices of its pizzas
- Conceptual steps:
  1. Partition the tuples in `Sells` into groups based on `rname`
  2. Compute `MIN(price)` and `MAX(price)` for each group
  3. Output one tuple for each group

```
SELECT rname, MIN(price), MAX(price)
FROM   Sells
GROUP BY rname;
```

Output

<i>rname</i>	<i>min</i>	<i>max</i>
Corleone Corner	19	25
Gambino Oven	16	16
Lorenzo Tavern	23	23
Mamma's Place	22	22
Pizza King	17	21

# Group by clause

## Introduction

- Example:
  - For each restaurant that sells some pizza, find its average pizza price. Show the restaurant in descending order of their average pizza price.

Sells

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Output

<u>rname</u>	<u>avgPrice</u>
Lorenzo Tavern	23.0000
Corleone Corner	22.6667
Mamma's Place	22.0000
Pizza King	19.0000
Gambino Oven	16.0000

```
SELECT  rname, AVG(price) AS avgPrice
FROM    Sells
GROUP BY rname
ORDER BY avgPrice DESC;
```

# Group by clause

## Introduction

- Example:
  - Find the number of students for each (dept, year) combination. Show the output in ascending order of (dept, year)

Students

<i>sid</i>	<i>name</i>	<i>year</i>	<i>dept</i>
12345	Alice	1	CS
67890	Bob	1	Eng
11123	Carol	1	CS
20135	Eve	2	Eng
87012	Fred	2	CS

Output

<i>dept</i>	<i>year</i>	<i>count</i>
CS	1	2
CS	2	1
Eng	1	1
Eng	2	1

```
SELECT  dept, year, COUNT(*)
FROM    Students
GROUP BY dept, year
ORDER BY dept, year;
```

# Group by clause

## Grouping

- In a query with “GROUP BY  $a_1, a_2, \dots, a_n$ ”, two tuples  $t$  &  $t'$  belong to the **same group** if the following expression evaluates to true
  - $(t.a_1 \text{ IS NOT DISTINCT FROM } t'.a_1) \wedge \dots \wedge (t.a_n \text{ IS NOT DISTINCT FROM } t'.a_n)$
  - In other words, all values are NOT DISTINCT
    - Two null values are considered non-distinct
- Example:
  - How many groups in  $R$  if  $R$  is grouped by  $\{A, C\}$

R		
A	B	C
null	1	19
null	2	19
6	1	null
6	20	null
20	2	10
1	1	2
1	18	2



# Group by clause

## Grouping

- In a query with “GROUP BY  $a_1, a_2, \dots, a_n$ ”, two tuples  $t$  &  $t'$  belong to the **same group** if the following expression evaluates to true
  - $(t.a_1 \text{ IS NOT DISTINCT FROM } t'.a_1) \wedge \dots \wedge (t.a_n \text{ IS NOT DISTINCT FROM } t'.a_n)$
  - In other words, all values are NOT DISTINCT
    - Two null values are considered non-distinct
- Each output tuple corresponds to one group

**R**

<b>A</b>	<b>B</b>	<b>C</b>
null	1	19
null	2	19
6	1	null
6	20	null
20	2	10
1	1	2
1	18	2

---

# Group by clause

## Grouping

- Condition
  - For each column  $A$  in relation  $R$  that appears in SELECT, one of the following condition must hold
    1. Column  $A$  appears in the GROUP BY clause
    2. Column  $A$  appears in aggregated expression in SELECT (e.g., MIN( $A$ ))
    3. The primary (or candidate) key of  $R$  appears in the GROUP BY clause
  - If an aggregate function appears in SELECT, and there is no GROUP BY clause
    - Then the SELECT must not contain any column that is not in an aggregated expression

# Group by clause

## Grouping

- Condition
  - For each column  $A$  in relation  $R$  that appears in SELECT, one of the following condition must hold
    1. Column  $A$  appears in the GROUP BY clause
    2. Column  $A$  appears in aggregated expression in SELECT (e.g., MIN( $A$ ))
    3. The primary (or candidate) key of  $R$  appears in the GROUP BY clause

Students

<i>sid</i>	<i>name</i>	<i>year</i>	<i>dept</i>
12345	Alice	1	CS
11123	Carol	1	CS
87012	Fred	2	CS
67890	Bob	1	Eng
20135	Eve	2	Eng

```
SELECT    dept, year, COUNT(*)  
FROM      Students  
GROUP BY dept;
```

---

# Group by clause

## Grouping

- Condition
  - If an aggregate function appears in SELECT, and there is no GROUP BY clause
  - Then the SELECT must not contain any column that is not in an aggregated expression

```
SELECT rname, MIN(price), MAX(price)
FROM   Sells;
```

```
SELECT MIN(price), MAX(price)
FROM   Sells;
```

# Group by clause

## Example

- Question:
  - For each restaurant that sells some pizza, find its name, area, and the average price of its pizzas
- Conceptual steps
  1. Get restaurants that sells pizza
  2. Partition by rname
  3. Output one tuple for each group

Sells NATURAL JOIN Restaurants

<u>rname</u>	<u>area</u>	<u>price</u>
Corleone Corner	North	24
Corleone Corner	North	25
Corleone Corner	North	19
Gambino Oven	Central	16
Lorenzo Tavern	Central	23
Mamma's Place	South	22
Pizza King	East	17
Pizza King	East	21

output

<u>rname</u>	<u>area</u>	<u>avg</u>
Corleone Corner	North	19.00
Gambino Oven	Central	16.00
Lorenzo Tavern	Central	22.67
Mamma's Place	South	23.00
Pizza King	East	23.00

# Group by clause

## Example

- Question:
  - For each restaurant that sells some pizza, find its name, area, and the average price of its pizzas
- Conceptual steps
  1. Get restaurants that sells pizza
  2. Partition by rname
  3. Output one tuple for each group

Sells NATURAL JOIN Restaurants

<u>rname</u>	<u>area</u>	<u>price</u>
Corleone Corner	North	24
Corleone Corner	North	25
Corleone Corner	North	19
Gambino Oven	Central	16
Lorenzo Tavern	Central	23
Mamma's Place	South	22
Pizza King	East	17
Pizza King	East	21

```
SELECT    R.rname, R.area
          AVG(S.price)
FROM      Sells S NATURAL JOIN
          Restaurants R
GROUP BY  R.rname;
```

output

<u>rname</u>	<u>area</u>	<u>avg</u>
Corleone Corner	North	19.00
Gambino Oven	Central	16.00
Lorenzo Tavern	Central	22.67
Mamma's Place	South	23.00
Pizza King	East	23.00

# Having clause

## Introduction

- Example:
  - Find restaurants that sell pizzas with an average selling price of at least \$22
- Conceptual steps
  1. Get restaurants that sells pizza
  2. Partition by rname
  3. Output one tuple for each group

```
SELECT    rname
FROM      Sells S
GROUP BY  rname
WHERE     AVG(price) >= 22;
```

Sells

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

output

<u>rname</u>
Corleone Corner
Lorenzo Tavern
Mamma's Place

# Having clause

## Introduction

- Example:
  - Find restaurants that sell pizzas with an average selling price of at least \$22
- Conceptual steps
  1. Get restaurants that sells pizza
  2. Partition by rname
  3. Output one tuple for each group

```
SELECT    rname
FROM      Sells S
GROUP BY  rname
HAVING    AVG(price) >= 22;
```

Sells

<u>rname</u>	<u>pizza</u>	<u>price</u>
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

output

<u>rname</u>
Corleone Corner
Lorenzo Tavern
Mamma's Place



# Having clause

## Properties

- Condition
  - For each column  $A$  in relation  $R$  that appears in HAVING, one of the following condition must hold
    1. Column  $A$  appears in the GROUP BY clause
    2. Column  $A$  appears in aggregated expression in HAVING (e.g.,  $\text{MIN}(A)$ )
    3. The primary (or candidate) key of  $R$  appears in the GROUP BY clause
  
- ❖ Similar to GROUP BY but SELECT is replaced with HAVING

# Summary

## □ Conceptual evaluation of queries

### □ Query:

SELECT	<b>DISTINCT select-list</b>
FROM	<b>from-list</b>
WHERE	<b>where-condition</b>
GROUP BY	<b>groupby-list</b>
HAVING	<b>having-condition</b>
ORDER BY	<b>orderby-list</b>
LIMIT	<b>limit-spec</b>
OFFSET	<b>offset-spec</b>

1. Compute the cross-product of the tables in **from-list**
2. Select the tuples in the cross-product that evaluate to TRUE for the **where-condition**
3. Partition the selected tuples into groups using the **groupby-list**
4. Select the groups that evaluate to TRUE for the **having-condition** condition
5. For each selected group, generate an output tuple by selecting/computing the attributes/expressions that appear in the **select-list**
6. Remove any duplicate output tuples because of **DISTINCT**
7. Sort the output tuples based on the **orderby-list**
8. Remove the appropriate output tuples based on the **limit-spec** & **offset-spec**

- Aggregate functions

- Introduction

- Order by clause

- Group by clause

- Having clause

- Subqueries

- select clause

- from clause

- where clause

- Queries with universal quantification

---

## Subqueries

# Subqueries

## Introduction

- We let an SQL query be of the form

```
SELECT [ DISTINCT ] select_list  -- select clause
FROM      from_list             -- from clause
[ WHERE      expression ]       -- where clause
```

- A **subquery** is an SQL query that may be added into
  - **select clause**
    - Scalar subqueries
  - **from clause**
    - Common table expression
  - **where clause**
    - EXISTS; IN; ANY/SOME; ALL

# Subqueries

## Scoping rules

- Queries with subquery expressions are also called **nested queries**
- A subquery expression is referred to as an **inner query** that is nested within an **outer query**
- **Scoping rules for table alias (a.k.a. tuple variable):**
  - A tuple variable declared in a subquery/query  $Q$  can be used only in  $Q$  and any subquery nested within  $Q$
  - If a tuple variable is declared both locally in a subquery  $Q$  as well as in an outer query, the local declaration applies in  $Q$

# Subqueries

## Scoping rules

- **Scoping rules for table alias (a.k.a. tuple variable):**
  - A tuple variable declared in a subquery/query  $Q$  can be used only in  $Q$  and any subquery nested within  $Q$
  - If a tuple variable is declared both locally in a subquery  $Q$  as well as in an outer query, the local declaration applies in  $Q$
- **Example:**
  - Given  $R(a, b)$ 
    - `SELECT S.a FROM ( SELECT * FROM R ) AS S;`
    - `SELECT * FROM ( SELECT * FROM R ) AS S, R;`
    - `SELECT * FROM ( SELECT * FROM R ) AS S  
WHERE S.a = 1;`
    - `SELECT *  
FROM ( SELECT * FROM R ) AS S,  
      ( SELECT * FROM S ) AS T;`

# select clause

## Scalar subqueries

- A **scalar subquery** is a subquery that returns at most one tuple with one column
  - If the result of the subquery is empty, its return value is `null`
  - A scalar subquery can be used as a scalar expression
  - Example:

- ```
SELECT rname
  FROM Restaurants
 WHERE rname = 'Corleone Corner';
```

Output

| <u>rname</u>    |
|-----------------|
| Corleone Corner |

- ```
SELECT *
  FROM Restaurants
 WHERE rname = 'Corleone Corner';
```

Output

<u>rname</u>	<i>area</i>
Corleone Corner	North

# select clause

## Scalar subqueries

- For each restaurant that sells Funghi, find its name, area, and selling price
  - name and price can be obtained from Sells
  - area can be obtained from Restaurants
  - Attempt #1

- Natural join

```
SELECT R.rname, R.area, S.price
FROM   Sells S NATURAL JOIN Restaurants R
WHERE  S.pizza = 'Funghi';
```

- Attempt #2

- Scalar subquery

```
SELECT rname,
       (SELECT R.area FROM Restaurants R
        WHERE R.rname = S.rname), price
FROM   Sells S   WHERE pizza = 'Funghi';
```



# from clause

## Common table expressions (CTEs)

- A **table expression** computes a table
  - Computing a table is a subquery!
- A **common table expression** is a temporary named result set that can be queried

- Syntax:

**WITH**

```
cte1 AS (subquery1) [,
cte2 AS (subquery2) [ ... ] ]
```

**query**

- Property:
  - Each  $cte_i$  is the name of a temporary relation defined by  $subquery_i$
  - **query** is the SQL statement that references **cte**
  - CTEs can be used for writing recursive queries (*not covered*)

---

# from clause

## Example

- Given
  - Courses (cid, cname, credits)
  - Enrolls (sid, cid, grade)
- Find the courses where the total number of enrolled students is higher than that for the course named “Database Systems”. Output the cname and the total number of enrolled students for each selected course.
  - Assume that cname is a candidate key for Courses
  - Idea:
    - Get the count for “Database Systems”
    - Compare the count
    - Get the pair (courses, number enrolled)

# from clause

## Example

- Idea:
  - Get the count for “Database Systems”
  - Compare the count
  - Get the pair (courses, number enrolled)

```
SELECT      C.cname, ( SELECT COUNT(*)
                        FROM      Enrolls E
                        WHERE     E.cid = C.cid ) AS num
FROM        Courses C NATURAL JOIN Enrolls E
GROUP BY    C.cid
HAVING      COUNT(*) >
            ( SELECT COUNT(*)
              FROM    Courses C NATURAL JOIN Enrolls E
              WHERE   C.cname = 'Database Systems' );
```

# from clause

## Example

- Somewhat duplicated?
- If we rearrange

```
SELECT    cname, num
FROM      ( SELECT C.cid, C.cname, COUNT(*) AS num
            FROM    Courses C NATURAL JOIN Enrolls E
            GROUP BY C.cid ) AS X
WHERE     num >
          ( SELECT COUNT(*)
            FROM    Courses C NATURAL JOIN Enrolls E
            WHERE   C.cname = 'Database Systems' );
```

# from clause

## Example

- Somewhat duplicated?
- If we rearrange
- But we cannot refer to X!
  - Without CTE

```
WITH X AS
( SELECT  C.cid, C.cname, COUNT(*) AS num
  FROM    Courses C NATURAL JOIN Enrolls E
  GROUP BY C.cid )
SELECT  cname, num
FROM    X
WHERE   num >
      ( SELECT num FROM    X
        WHERE  cname = 'Database Systems' );
```

# where clause

## Types of subqueries

- **EXISTS** subqueries
  - Returns true if the result subquery is non-empty
  - Otherwise, false
- **IN** subqueries
  - Subquery must return exactly one column
  - Returns false if the result of subquery is empty
  - Otherwise, for subquery result  $= \{v_1, \dots v_n\}$  and expression result  $v$ ;  
return the result of the boolean expression
$$((v = v_1) \vee (v = v_2) \vee \dots \vee (v = v_n))$$

# where clause

## Types of subqueries

- **ANY/SOME** subqueries

- Subquery must return exactly one column
- Returns false if the result of subquery is empty
- Otherwise, for subquery result  $= \{v_1, \dots v_n\}$ , expression result  $v$ , and operator  $\oplus$ ; return the result of the boolean expression
$$((v \oplus v_1) \vee (v \oplus v_2) \vee \dots \vee (v \oplus v_n))$$

- **ALL** subqueries

- Subquery must return exactly one column
- Returns false if the result of subquery is empty
- Otherwise, for subquery result  $= \{v_1, \dots v_n\}$ , expression result  $v$ , and operator  $\oplus$ ; return the result of the boolean expression
$$((v \oplus v_1) \wedge (v \oplus v_2) \wedge \dots \wedge (v \oplus v_n))$$

# where clause

## EXISTS subqueries

- Example:
  - Find distinct customers who like some pizza sold by “Corleone Corner”
- Idea:
  - Find pizza sold by “Corleone Corner”
  - Check if customer likes the pizza
  - Remove duplicates
- Query:

```
SELECT DISTINCT L.cname
FROM    Likes L INNER JOIN Sells S
        ON S.pizza = L.pizza
WHERE   S.rname = 'Corleone Corner';
```



# where clause

## EXISTS subqueries

- Example:
  - Find distinct customers who like some pizza sold by “Corleone Corner”
- Idea:
  - Find pizza sold by “Corleone Corner”
  - Check if customer likes the pizza
  - Remove duplicates
- Query:

```
SELECT DISTINCT cname
FROM   Likes L
WHERE  EXISTS (
        SELECT 1 FROM Sells S
        WHERE   S.rname = 'Corleone Corner'
        AND     S.pizza = L.pizza );
```

# where clause

## NOT EXISTS subqueries

- Example:
  - Find distinct customers who does not like some pizza sold by “Corleone Corner”
- Idea:
  - Find customers who likes pizza sold by “Corleone Corner”
  - Set difference

- Query:

```
SELECT cname FROM Customers
EXCEPT
SELECT L.cname
FROM Likes L INNER JOIN Sells S
      ON S.pizza = L.pizza
WHERE S.rname = 'Corleone Corner';
```

# where clause

## NOT EXISTS subqueries

- Example:
  - Find distinct customers who does not like some pizza sold by “Corleone Corner”
- Idea:
  - Invert condition
  - NOT EXISTS

- Query:

```
SELECT DISTINCT cname
FROM   Likes L
WHERE  NOT EXISTS (
        SELECT 1 FROM Sells S
        WHERE   S.rname = 'Corleone Corner'
        AND     S.pizza = L.pizza );
```

# where clause

## IN subqueries

- Example:

- Find pizzas that contain ham or seafood

- UNION

```
SELECT pizza FROM Contains WHERE ingredient = 'ham'  
UNION
```

```
SELECT pizza FROM Contains WHERE ingredient =  
'seafood';
```

- OR

```
SELECT DISTINCT pizza FROM Contains  
WHERE ingredient = 'ham' OR ingredient = 'seafood';
```

- IN

```
SELECT DISTINCT pizza FROM Contains  
WHERE ingredient IN ('ham', 'seafood');
```

# where clause

## IN subqueries

- Example:
  - Find distinct customers who like some pizza sold by “Corleone Corner”
- Idea:
  - Find pizza sold by “Corleone Corner”
  - Check if customer likes the pizza
  - Remove duplicates
- Query:

```
SELECT DISTINCT cname
FROM   Likes L
WHERE  pizza IN (
        SELECT pizza FROM Sells S
        WHERE   rname = 'Corleone Corner'
      );
```

# where clause

## ANY/SOME subqueries

- Example:
  - Find distinct restaurants that sell some pizza P1 that is more expensive than some pizza P2 sold by “Corleone Corner”. P1 and P2 are not necessarily the same pizza. Exclude “Corleone Corner” from the query result.
- Idea:
  - Find pizza sold by “Corleone Corner”
  - Check if price more expensive
  - Remove duplicates and “Corleone Corner”

- Query:

```
SELECT DISTINCT rname FROM Sells S
WHERE  rname <> 'Corleone Corner' AND EXISTS (
    SELECT 1 FROM Sells S2 WHERE S.price > S2.price
    AND rname = 'Corleone Corner' );
```

# where clause

## ANY/SOME subqueries

- Example:
  - Find distinct restaurants that sell some pizza P1 that is more expensive than some pizza P2 sold by “Corleone Corner”. P1 and P2 are not necessarily the same pizza. Exclude “Corleone Corner” from the query result.
- Idea:
  - Find pizza sold by “Corleone Corner”
  - Check if price more expensive
  - Remove duplicates and “Corleone Corner”
- Query:

```
SELECT DISTINCT rname FROM Sells
WHERE rname <> 'Corleone Corner' AND price > ANY (
    SELECT price FROM Sells
    WHERE rname = 'Corleone Corner' );
```

# where clause

## ANY/SOME subqueries

- Example:
  - For each restaurant, find the name, and price of its most expensive pizzas. Exclude restaurants that do not sell any pizza.
- Idea:
  - Find pizza that has other more expensive pizza
  - Remove these pizzas

- Query:

```
SELECT rname, pizza, price FROM Sells
EXCEPT
SELECT rname, pizza, price FROM Sells S1
WHERE price < ANY (
    SELECT S2.price FROM Sells S2
    WHERE S2.rname = S1.rname
);
```



# where clause

## ALL subqueries

- Example:
  - For each restaurant, find the name, and price of its most expensive pizzas. Exclude restaurants that do not sell any pizza.
- Idea:
  - Find pizza that is more expensive or as expensive than all other pizzas

- Query:

```
SELECT rname, pizza, price FROM Sells S1
WHERE price >= ALL (
    SELECT S2.price FROM Sells S2
    WHERE S2.rname = S1.rname
);
```

---

# Queries with universal quantification

## Example

- Given:
  - Courses (courseID, name, dept)
  - Students (studentID, name, birthdate)
  - Enrolls (sid, cid, grade)
- Question:
  - Find the names of all students who have enrolled in **all** the courses offered by CS department

# Queries with universal quantification

## Example

- Question:
  - Find the names of all students who have enrolled in **all** the courses offered by CS department
- Analysis:
  - Let  $R$  denote the set of all students who have enrolled in **all** the courses offered by CS department
  - Let  $\bar{R} = \text{Students} - R$   
= set of all students who have NOT enrolled ... CS department
  - A student  $s \in \bar{R}$  iff there exists some CS course  $c$  such that  $s$  is not enrolled in  $c$
  - Given a studentID  $x$  let  $F(x)$  = set of courseIDs of CS courses that are not enrolled by student with studentID  $x$
  - $\bar{R} = \{s \in \text{Students} \mid F(s.\text{studentID}) \neq \emptyset\}$

# Queries with universal quantification

## Example

- Question:
  - Find the names of all students who have enrolled in **all** the courses offered by CS department
- Analysis:
  - $\bar{R} = \{s \in \text{Students} \mid F(s.\text{studentID}) \neq \emptyset\}$
  - $\bar{R}$  can be computed by the following pseudo SQL query:  

```
SELECT s.studentID FROM Students
WHERE EXISTS (F.s.studentID));
```
  - $R$  can be computed by the following pseudo SQL query:  

```
SELECT s.studentID FROM Students
WHERE NOT EXISTS (F.s.studentID));
```

# Queries with universal quantification

## Example

- Question:
  - Find the names of all students who have enrolled in **all** the courses offered by CS department

- Solution:  $F(x)$

```
SELECT courseID FROM Courses C
WHERE dept = 'CS'
  AND NOT EXISTS (
    SELECT 1 FROM Enrolls E
    WHERE E.cid = C.courseID
      AND E.sid = x );
```

# Queries with universal quantification

## Example

- Question:
  - Find the names of all students who have enrolled in **all** the courses offered by CS department
- Solution: *R*

```
SELECT name FROM Students S
WHERE NOT EXISTS (
    SELECT courseID FROM Courses C
    WHERE dept = 'CS'
    AND NOT EXISTS (
        SELECT 1 FROM Enrolls E
        WHERE E.cid = C.courseID
        AND E.sid = x )
);
```

---

# Summary

- ❑ A **subquery** is an SQL query that may be added into
  - ❑ **select clause**
    - ❑ Scalar subqueries
    - ❑ One column and one tuple
  - ❑ **from clause**
    - ❑ Common table expression
  - ❑ **where clause**
    - ❑ EXISTS; IN; ANY/SOME; ALL