

CS2102

Database Systems

Slides adapted from Prof. Chan Chee Yong

LECTURE 07A

STORED PROCEDURE, FUNCTIONS, AND TRIGGERS

SQL

Conceptual evaluation of queries

Query:

SELECT	DISTINCT select-list
FROM	from-list
WHERE	where-condition
GROUP BY	groupby-list
HAVING	having-condition
ORDER BY	orderby-list
LIMIT	limit-spec
OFFSET	offset-spec

1. Compute the cross-product of the tables in **from-list**
2. Select the tuples in the cross-product that evaluate to TRUE for the **where-condition**
3. Partition the selected tuples into groups using the **groupby-list**
4. Select the groups that evaluate to TRUE for the **having-condition** condition
5. For each selected group, generate an output tuple by selecting/computing the attributes/expressions that appear in the **select-list**
6. Remove any duplicate output tuples because of **DISTINCT**
7. Sort the output tuples based on the **orderby-list**
8. Remove the appropriate output tuples based on the **limit-spec** & **offset-spec**

Summary

- ❑ A **subquery** is an SQL query that may be added into
 - ❑ **select clause**
 - ❑ Scalar subqueries
 - ❑ One column and one tuple
 - ❑ **from clause**
 - ❑ Common table expression
 - ❑ **where clause**
 - ❑ EXISTS; IN; ANY/SOME; ALL

- Transaction
 - Stored procedures and functions
 - Introduction
 - Creation
 - Usage
 - plpgsql
 - Triggers
 - Creation
 - Learn by example
-

Overview

- Transaction
 - Stored procedures and functions
 - Introduction
 - Creation
 - Usage
 - plpgsql
 - Triggers
 - Creation
 - Learn by example
-

Transaction

Transaction

Introduction

- A **transaction** consists of one or more update/retrieval operations (i.e., SQL statements)
 - An abstraction of a **logical unit of work**

Syntax

- **BEGIN code { COMMIT | ROLLBACK }**
- The **BEGIN** command **starts a new transaction**
- Each transaction must end with either
 - **COMMIT** **success** → *update database*
 - **ROLLBACK** **failure** → *restore database to state before BEGIN*

Transaction

Example

- Transfer \$1000 from aid = 1 to aid = 2
 - What happens when error occurred?

BEGIN

```
UPDATE accounts  
SET balance = balance  
+ 1000  
WHERE accountID = 2;
```

```
UPDATE accounts  
SET balance = balance  
- 1000  
WHERE accountID = 1;
```

COMMIT

VS

```
UPDATE accounts  
SET balance = balance  
+ 1000  
WHERE accountID = 2;
```

```
UPDATE accounts  
SET balance = balance  
- 1000  
WHERE accountID = 1;
```

Transaction

ACID properties

- **Atomicity** Either all the effects of the transactions are reflected in the database or none are
all or none
- **Consistency** The execution of a transaction in isolation preserves the consistency of the database
user-defined property should be preserved
- **Isolation** The execution of a transaction is isolated from the effects of other concurrent transaction execution
can run concurrently
- **Durability** The effect of a committed transaction persists in the database even in the presence of system failures
commit is permanent

Transaction

Constraint check

- By default constraints are checked **immediately** at the end of SQL statement execution
 - A violation will cause the statement to be **rolledback**
- Can be **deferred** to the end of transaction execution
 - A violation will cause the transaction to be **aborted**
- The type of constraint checking should be specified as part of constraint declaration

```
CREATE TABLE students (  
    name      varchar(100),  
    CHECK (name IS NOT NULL)  
        DEFERRABLE INITIALLY DEFERRED  
);
```

Transaction

Circular reference?

- Given T1 (a, b) and T2 (a, b) such that
 - T1.a is a foreign key that refers to T2.a
 - T2.a is a foreign key that refers to T1.a
- How?
 - ALTER TABLE T1 ADD FOREIGN KEY (a)
REFERENCES T2 (a) DEFERRABLE INITIALLY DEFERRED;
 - ALTER TABLE T2 ADD FOREIGN KEY (a)
REFERENCES T1 (a) DEFERRABLE INITIALLY DEFERRED;
 - Transaction!
BEGIN;
INSERT INTO T1 VALUES (1,2);
INSERT INTO T2 VALUES (1,2);
COMMIT;

- Transaction
 - Stored procedures and functions
 - Introduction
 - Creation
 - Usage
 - plpgsql
 - Triggers
 - Creation
 - Learn by example
-

Stored procedures and function

History

History of stored procedures

- **SQL 1999** standard proposes a language for stored procedures and triggers
- PostgreSQL (*and other DBMS*) can **store**, **share**, and **execute** code on the server
- **PL/pgSQL** is PostgreSQL's language specifically designed to seamlessly embed SQL code and interact with the database SQL code
- PL/pgSQL partially complies with the SQL standard
- PL/pgSQL is similar to Oracle PL/SQL, **SQLServer Transact-SQL**, and **DB2 SQL Procedural Language** as well as other languages implementing variants of the **Persistent Storage Modules** portion of the SQL standard

History

Performance

- Stored procedures are **compiled**. The code is cached and shared by all users. The code is executed on the server's side (*typically a powerful machine*). They generally incur **fewer data transfer** across the network.
- What it really says: The optimization of the SQL code they contain is **not adaptive**. The client's computing power is **underutilized**.

Productivity

- Stored procedures are **shared and reused** under access control. The application logic they encode is implemented and maintained **in a single place**.
- What it really says: The languages and concepts are **complicated**. The code is **not portable** from one DBMS to the next.

Security:

- Data manipulation can be **restricted** to calling stored procedures under strict access control.
- What it really says: It is **easy to make tragic mistake** in the coding

Introduction


What?

- Stored procedure is a procedure is a **procedure** / **function** / **subroutine** that is available to application that access a DBMS and is stored in the database
- Includes typical **control structures** such as **IF** and **WHILE**

Why?

- Uses includes (*but not limited to*) **data validation**, **access control mechanism**, **consolidation of logic**, **consistency check**, etc

How?

- Create function syntax
 - Function definition
- 

Creation

Create function syntax

- **CREATE** [OR REPLACE] **FUNCTION**
 func_name ([arg1 type1 [, arg2 type2 [...]])
 RETURNS ret_type AS func_def
 LANGUAGE lang;

Create procedure syntax

- **CREATE** [OR REPLACE] **PROCEDURE**
 proc_name ([arg1 type1 [, arg2 type2 [...]])
 AS proc_def
 LANGUAGE lang;

Removal

- DROP { FUNCTION | PROCEDURE } [IF EXISTS] name;

❖ Reference: <https://www.postgresql.org/docs/11/sql-createfunction.html>

Creation

Value conversion

```

○ CREATE OR REPLACE FUNCTION
  miles_to_km (mile numeric)
  RETURNS numeric AS
  ' BEGIN RETURN mile * 1.609; END; '
  LANGUAGE plpgsql;

```

Insertion

```

○ CREATE OR REPLACE PROCEDURE
  add_restaurant (name varchar(50),
                  area varchar(10)) AS
  ' BEGIN
    INSERT INTO Restaurants VALUES (name, area);
    END; '
LANGUAGE plpgsql;

```


Creation

Dollar-quoted string

- Problem

```
CREATE OR REPLACE FUNCTION hello_world ()  
RETURNS CHAR(11) AS  
‘ BEGIN RETURN ‘Hello World’; END; ’  
LANGUAGE plpgsql;
```

- Solution

```
CREATE OR REPLACE FUNCTION hello_world ()  
RETURNS CHAR(11) AS  
    BEGIN RETURN ‘Hello World’; END;  
LANGUAGE plpgsql;
```

Usage

Stored functions

- Can be used in an **expression**
 - Given Shops (name, distance)

```
SELECT S.name, miles_to_km(S.distance)
FROM   Shops S
WHERE  miles_to_km(S.distance) < 10;
```

Stored procedure

- Can be called

```
CALL add_restaurant('Domino''s', 'West');
```

plpgsql

What can we write?

- Declaration (before BEGIN) `DECLARE var type;`
- Assignment `var := expr;`
- Selection
 - `IF cond THEN stmt;`
 `[ELIF cond THEN stmt [...]]`
 `[ELSE stmt];`
- Iteration
 - `WHILE cond LOOP stmt; END LOOP;`
 - `FOR var IN (expr ...) LOOP`
 `{ stmt; | EXIT; | EXIT WHEN cond; } [...]`
 `END LOOP;`
 - `LOOP { stmt; | EXIT; | EXIT WHEN cond; } [...]`
 `END LOOP;`

plpgsql

Operations

- Arithmetic and bitwise

- **Simple** +, -, *, / (integer division truncates), %, ^ (exponent)
- **Bitwise** &, |, # (xor), ~ (not), <<, >>
- **Others** | / (square root), || / (cube root), @ (absolute value),
! (factorial postfix), !! (factorial prefix)

- References: <https://www.postgresql.org/docs/11/functions-math.html>

- Comparison

- **Simple** <, >, <=, >=, =, <>

- References: <https://www.postgresql.org/docs/11/functions-comparison.html>

plpgsql

Example

- Given two points (x1, y1) and (x2, y2), write a function to return the distance between the two points

```
CREATE OR REPLACE FUNCTION
dist (x1 numeric, y1 numeric,
      x2 numeric, y2 numeric) RETURNS numeric AS
$$ DECLARE distx numeric; disty numeric;
BEGIN
    distx := x1 - x2; disty := y1 - y2;
    RETURN |/ ((distx * distx) + (disty * disty));
END; $$ LANGUAGE plpgsql;
```

- Transaction
 - Stored procedures and functions
 - Introduction
 - Creation
 - Usage
 - plpgsql
 - Triggers
 - Creation
 - Learn by example
-

Triggers

History

History of triggers

- A trigger is a **procedure** or **function** that is executed **when a database event occurs** on a table
 - Example: INSERT, DELETE, UPDATE, CREATE TABLE, etc
- Triggers are used to **maintain integrity**, **propagate updates**, and **repair** the database
- Their syntax and semantics vary from one DBMS to another

History

Performance

- Triggers are **compiled**. The code is cached and shared by all users. The code is executed on the server's side (*typically a powerful machine*). They generally incur **no data transfer** across the network.
- What it really says: The optimization of the SQL code they contain is **not adaptive**. The client's computing power is **underutilized**.

Productivity

- Triggers are applied to **all interactions**. The application logic they encode is implemented and maintained **in a single place**.
- What it really says: The languages and concepts are **complicated**. The code is **not portable**. Interactions among triggers (chain reactions) and between triggers, constraints, and transactions are **difficult to control**.

Security:

- Data manipulation can be **restricted** to calling stored procedures under strict access control.
- What it really says: It is **easy to make tragic mistake** in the coding

Introduction


What?

- Trigger is a procedure is a **special stored procedure** that is executed when a specific event occurs in the database
- Triggers may occur **before** or **after** the event

Why?

- Uses includes (*but not limited to*) **data validation, access control mechanism, consolidation of logic, consistency check, etc**

How?

- Trigger syntax
 - Learn by example
- 

Creation

Create trigger syntax

- **CREATE TRIGGER** **trigger_name**
 { BEFORE | AFTER | INSTEAD OF }
 { event [OR event [...]] } ON **table**
 [FOR [EACH] { ROW | STATEMENT }]
 [WHEN **cond**] EXECUTE PROCEDURE **func_name()**;

Special variables

- **NEW** (for row only), **OLD** (for row only)
- **TG_WHEN** (returns 'BEFORE', 'AFTER')
- **TG_OP** (returns 'INSERT', 'UPDATE', 'DELETE', 'TRUNCATE')

Removal

- **DROP TRIGGER** [IF EXISTS] **trigger_name**
 ON **table** [CASCADE];

Learn by example

Preliminary

- Given Games (name, price)
- CREATE TABLE Games (
 name varvhar(50) PRIMARY KEY,
 price numeric NOT NULL CHECK (price > 0)
);

Games

<i>name</i>	<i>price</i>
A	100
B	200
C	300

Learn by example

For each row

- CREATE OR REPLACE FUNCTION t_func4()
RETURNS TRIGGER AS \$\$ BEGIN
RAISE NOTICE 'Trigger 4'; RETURN NULL;
END; \$\$ LANGUAGE plpgsql;
- CREATE TRIGGER trig4
BEFORE INSERT OR UPDATE ON Games
FOR EACH ROW WHEN (NEW.price > 100)
EXECUTE PROCEDURE t_func4();
- Behavior
 - Message 'Trigger 4' is displayed on psql whenever a row is inserted/updated
 - Row is NOT inserted when price > 100

Learn by example

For each row

- CREATE OR REPLACE FUNCTION t_func3()
RETURNS TRIGGER AS \$\$ BEGIN
RAISE NOTICE 'Trigger 3'; RETURN NEW;
END; \$\$ LANGUAGE plpgsql;
- CREATE TRIGGER trig3
BEFORE INSERT OR UPDATE ON Games
FOR EACH ROW WHEN (NEW.price > 100)
EXECUTE PROCEDURE t_func3();
- Behavior
 - Message 'Trigger 3' is displayed on psql whenever a row is inserted/updated
 - Row is STILL inserted when price > 100

Learn by example

For each row

- CREATE OR REPLACE FUNCTION t_func2()
RETURNS TRIGGER AS \$\$ BEGIN
RETURN (NEW.name, OLD.price * 2);
END; \$\$ LANGUAGE plpgsql;
- CREATE TRIGGER trig2
BEFORE UPDATE ON Games
FOR EACH ROW WHEN (NEW.price <= OLD.price)
EXECUTE PROCEDURE t_func2();
- Behavior
 - When the new price is lower than old price
 - Row is inserted such that the old price is doubled

Learn by example

For each statement

- CREATE OR REPLACE FUNCTION t_func1()
RETURNS TRIGGER AS \$\$ BEGIN
RAISE NOTICE 'Trigger 1'; RETURN NULL;
END; \$\$ LANGUAGE plpgsql;
- CREATE TRIGGER trig1
BEFORE INSERT OR UPDATE ON Games
FOR EACH STATEMENT
EXECUTE PROCEDURE t_func1();
- Behavior
 - Message 'Trigger 1' is displayed on psql once every statement
 - Values is STILL inserted
 - Reference: <https://www.postgresql.org/docs/current/plpgsql-trigger.html>

Multiple triggers

Trigger interactions

- Triggers **interact** with other triggers
 - Interaction happens in **alphabetical order**
 - Example
 - Create `<t_func4(), trig4>` on Games
 - Update 'A' to \$75
 - Create `<t_func2(), trig2>` on Games
 - Update 'A' to \$50
- ❖ What happened here?

Summary

- ❑ A **transaction** starts with **BEGIN** ends with either **COMMIT** or **ROLLBACK**
 - ❑ We assume ACID property is maintained
- ❑ **Stored function**
 - ❑ **CREATE** [**OR REPLACE**] **FUNCTION** **func_name** ...
 - ❑ **SELECT** **func_name** (...);
- ❑ **Stored procedure**
 - ❑ **CREATE** [**OR REPLACE**] **PROCEDURE** **proc_name** ...
 - ❑ **CALL** **proc_name** (...);
- ❑ **Triggers**
 - ❑ **CREATE TRIGGER** **trigger_name**
 { **BEFORE** | **AFTER** | **INSTEAD OF** }
 { **event** [**OR event** [...]] } **ON table**
 [**FOR** [**EACH**] { **ROW** | **STATEMENT** }]
 [**WHEN cond**] **EXECUTE PROCEDURE** **func_name**();