

# CS2102

# Database Systems

*Slides adapted from Prof. Chan Chee Yong*

---

LECTURE 11

MODERN DBMS

- Distributed DBMS
  - OLAP
    - Multidimensional aggregation
  - NoSQL
    - Graph database
  - Future
    - NewSQL
    - Database-as-a-Service
- 

## Overview

- Distributed DBMS
  - OLAP  
Multidimensional aggregation
  - NoSQL  
Graph database
  - Future  
NewSQL  
Database-as-a-Service
- 

# Distributed DBMS

# Distributed DBMS

## History

- Targeted to support the organizational structure of distributed enterprises
  - 1976: SDD-1 (*Computer Corporation of America*)
  - 1977: Distributed INGRES (*U.C. Berkeley*)
  - 1981: R\* (*IBM Research*)



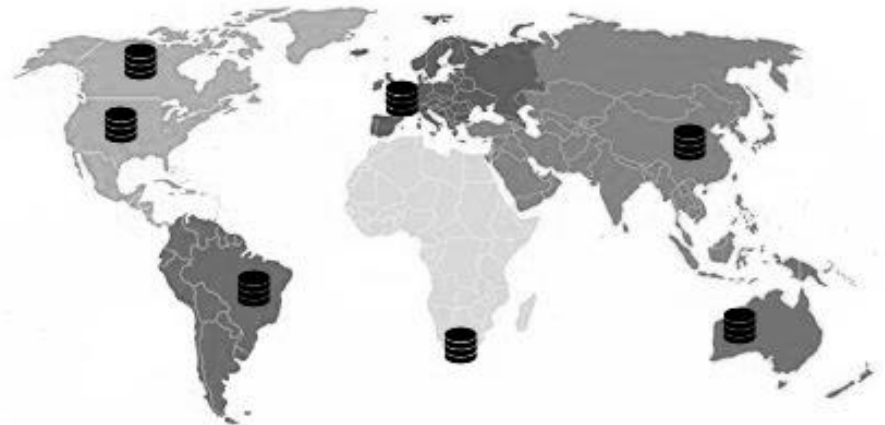
# Distributed DBMS

## Features

- Data is physically stored across multiple sites
  - But form a single logical database
  - Query and update feels like it is one database

## Why

- Enterprise is often distributed
- Database recovery




# Distributed DBMS

## Design

- Replication
  - Is table  $T$  available in more than one location?
  - Non-replicated
  - Partially replicated
  - Fully replicated

Loc	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
1	X			X	
2			X		
3		X			X




# Distributed DBMS

## Design

- Replication
  - Is table  $T$  available in more than one location?
  - Non-replicated
  - Partially replicated
  - Fully replicated

Loc	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
1	X			X	X
2	X		X		
3		X	X		X




# Distributed DBMS

## Design

- Replication
  - Is table  $T$  available in more than one location?
  - Non-replicated
  - Partially replicated
  - Fully replicated

Loc	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
1	X	X	X	X	X
2	X	X	X	X	X
3	X	X	X	X	X





# Distributed DBMS

## Design

- Fragmentation
  - For a single table  $T$ , is all its data in a one location?
  - Non-fragmented
  - Fragmented

Loc	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
1	$T_1$			$T_4$	
2		$T_2$	$T_3$	$T_4$	
3		$T_2$			$T_5$



# Distributed DBMS

## Design

- Fragmentation
  - For a single table  $T$ , is all its data in a one location?
  - Non-fragmented
  - Fragmented

Loc	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
1	$T_1^1$			$T_4$	$T_5^2$
2	$T_1^2$	$T_2^1$	$T_3$		$T_5^3$
3		$T_2^2$			$T_5^1$




# Distributed DBMS

## Design

- Replication & fragmentation can be mixed

Loc	$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
1	$T_1^1 T_1^2$			$T_4$	$T_5^1 T_5^2$
2	$T_1^2$	$T_2^1$	$T_3$		$T_5^1 T_5^3$
3	$T_1^1$	$T_2^2$			$T_5^1$



# Distributed DBMS

## Consistency? Availability? Partition tolerance?

- **Consistency:** *every read receives the most recent write or an error*
- **Availability:** *every request receives a (non-error) response, without guarantee that it contains the most recent write*
- **Partition tolerance:** *the system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes*




# Parallel DBMS

## Motivation

- Advantages of distributed DBMS without the disadvantages of distributed DBMS
- Use of shared file system across multiple machine

## Main advantage

- High performance
- High availability



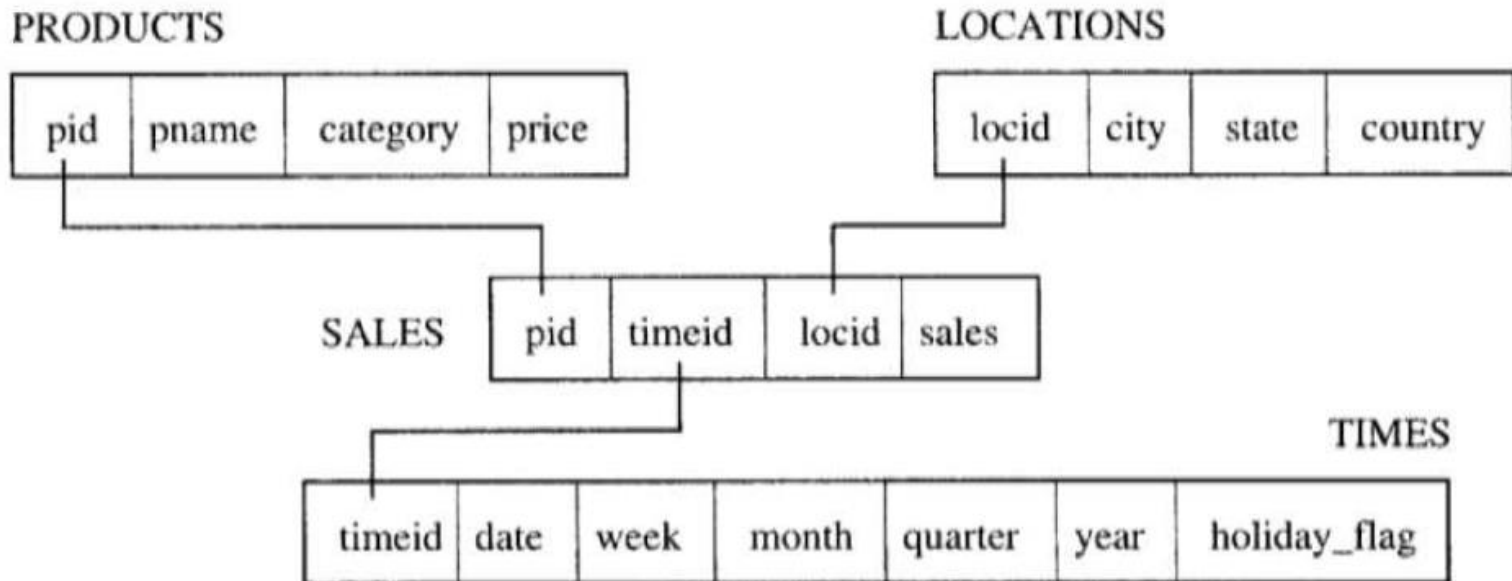
- Distributed DBMS
  - OLAP
    - Multidimensional aggregation
  - NoSQL
    - Graph database
  - Future
    - NewSQL
    - Database-as-a-Service
- 

OLAP

# OLAP

## Star schema

- Data is modeled using a **fact table** and **dimensions tables**



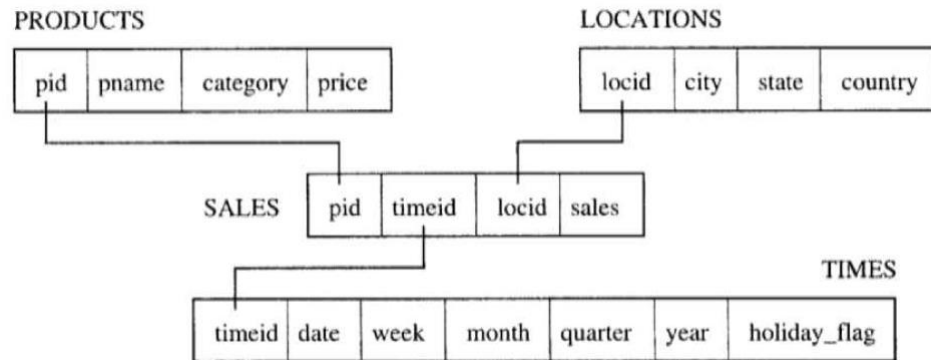
(Ramakrishnan & Gehrke, 2003)

# OLAP

## Multidimensional aggregation

- Total sales
  - for each product?
    - and each time?
      - and each location?

```
SELECT P.category,  
       T.year,  
       L.city,  
       SUM(S.sales)  
FROM Sales S  
   JOIN Products P on ...  
   JOIN Times T on ...  
   JOIN Locations L on ...  
GROUP BY P.category,  
         T.year,  
         L.city
```



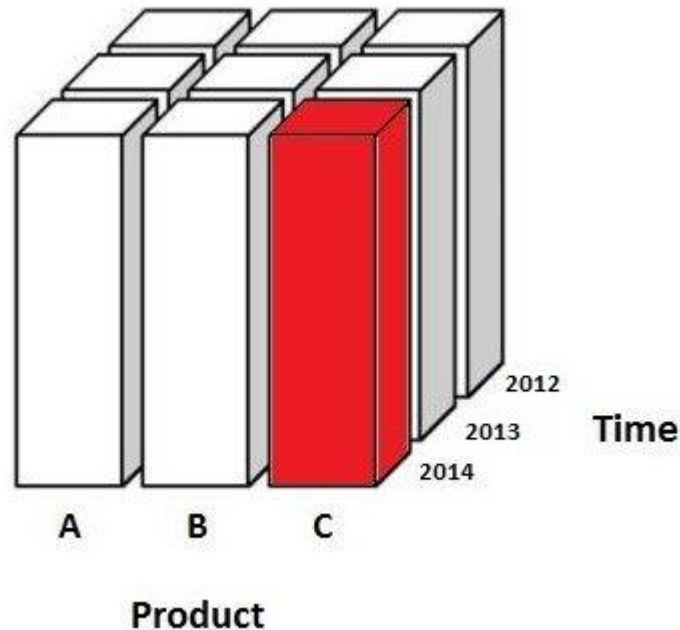
(Ramakrishnan & Gehrke, 2003)



# OLAP

## Multidimensional aggregation


- Total sales
  - for each product?
    - and each time?
      - and each location?



# OLAP

## Multidimensional aggregation


- Total sales
  - for each product?
    - and each time?
      - and each location?



# OLAP

## Multidimensional aggregation


- Total sales
  - for each product?
    - and each time?
      - and each location?
      - and?



# OLAP

## Multidimensional aggregation

- Total sales
  - for each product?
    - and each time?
      - and each location?



```
SELECT P.category,  
       T.year,  
       L.city,  
       SUM(S.sales)  
FROM Sales S  
      JOIN Products P on ...  
      JOIN Times T on ...  
      JOIN Locations L on ...  
GROUP BY CUBE( P.category,  
               T.year,  
               L.city )
```

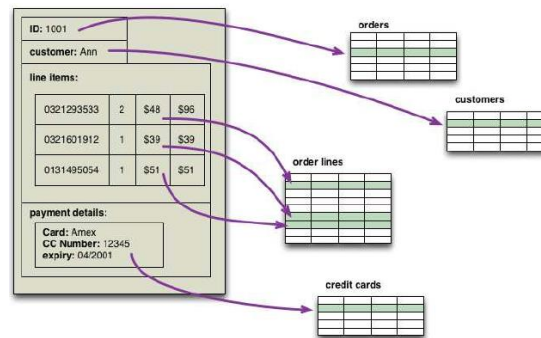
- Distributed DBMS
  - OLAP
    - Multidimensional aggregation
  - NoSQL
    - Graph database
  - Future
    - NewSQL
    - Database-as-a-Service
- 

# NoSQL

# NoSQL

## Not only SQL

- Early NoSQL systems
  - Google's Bigtable; Amazon's Dynamo; Yahoo!'s PNUTS
- Data models
  - Key-value
  - Document
  - Graphs
  - etc...



(Martin Fowler, 2012)

- Features?
  - Schema-less data
  - Simple access API (put, get, and delete)
    - Instead of query language
  - Limited/no ACID transactional support
  - Weak consistency for replicated data

---

# Graph database system

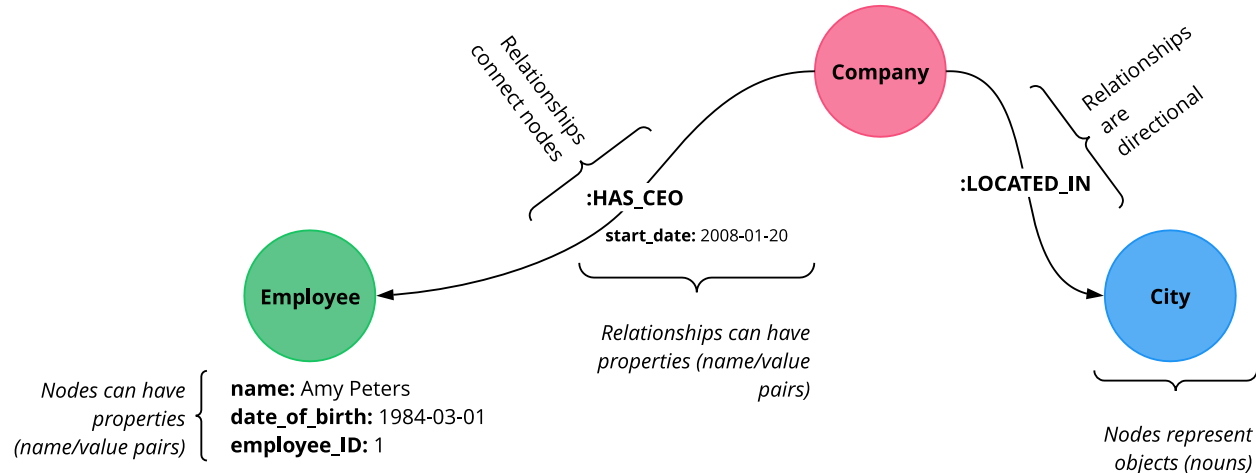
## Introduction

- Based on different graph data models:
  - **Property graphs**
    - *Systems:* JanusGraph, Neo4J, etc
  - **RDF graphs**
    - RDF = resource description framework
    - Data stores known as triplestores/semantic graph databases
      - Store data as (subject, predicate, object) triples
    - Query language: SPARQL
    - Supports RDF schema (RDFS) & web ontology language (OWL) inference
    - *Systems:* AllegroGraph, GraphDB, etc
  - **Hypergraphs**
    - *Systems:* HyperGraphDB, Microsoft Graph Engine, etc

# Property graph data model

## Model

- Nodes represent entities
  - Each node has at least one label and possibly properties
- Directed edges represent relationships between entities
  - Each relationship has a type and possibly properties
- Each property is a key-value pair

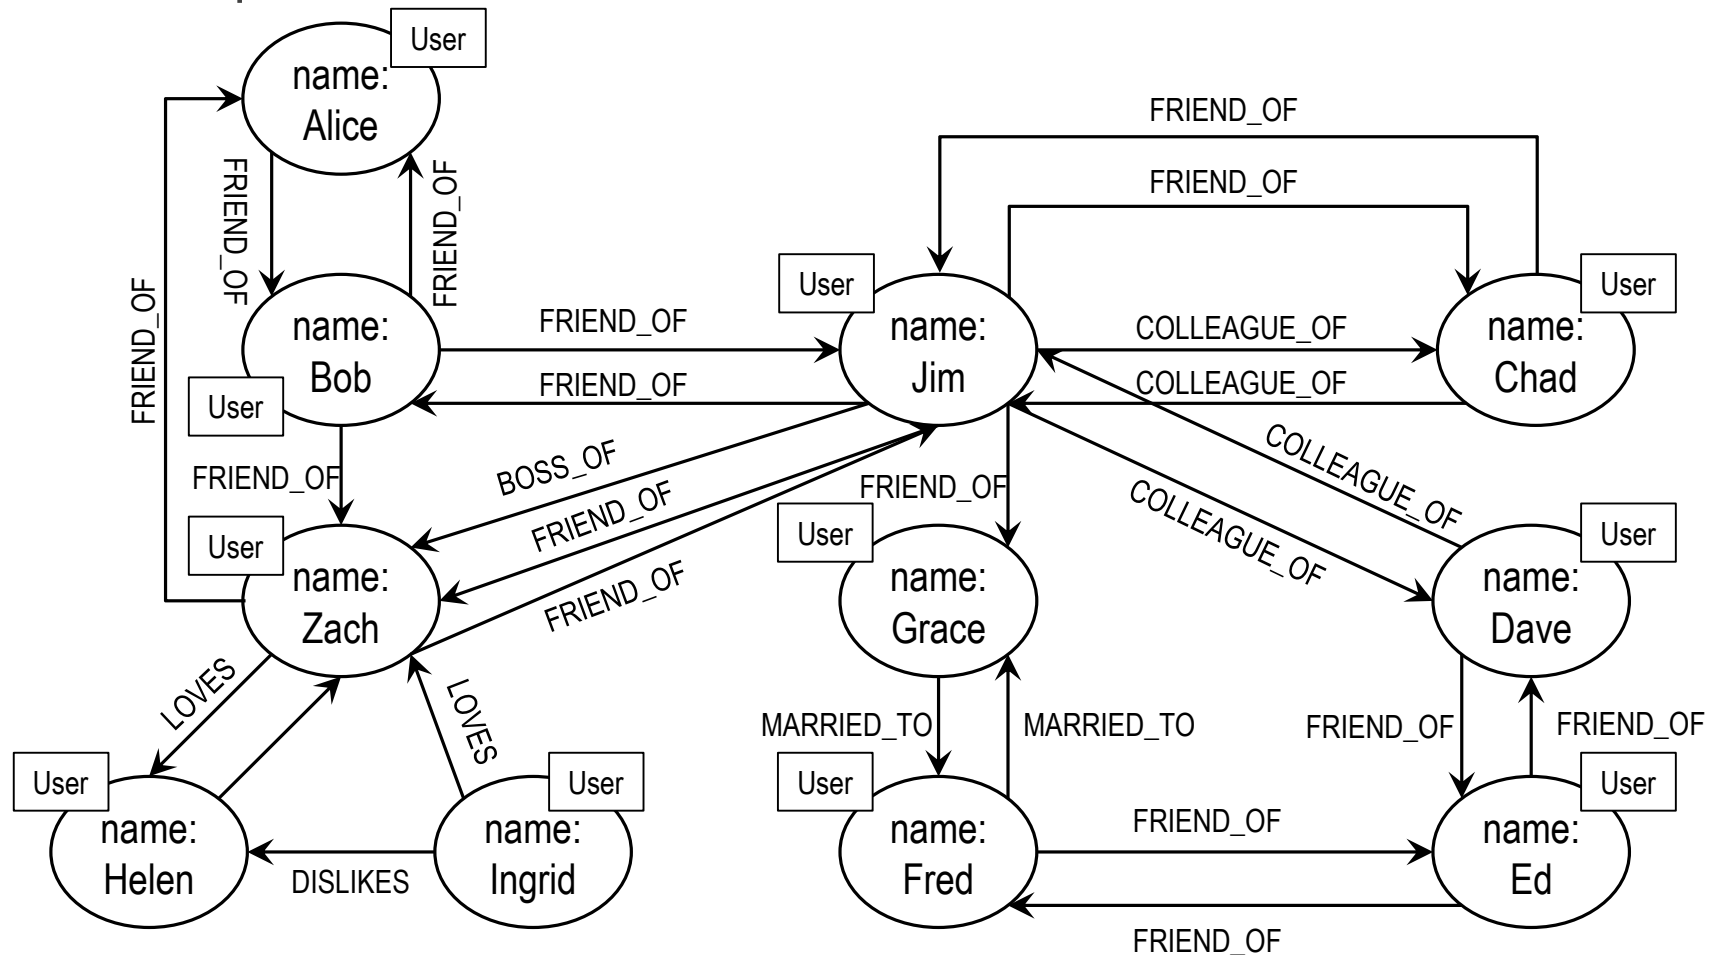




# Property graph data model

## Example

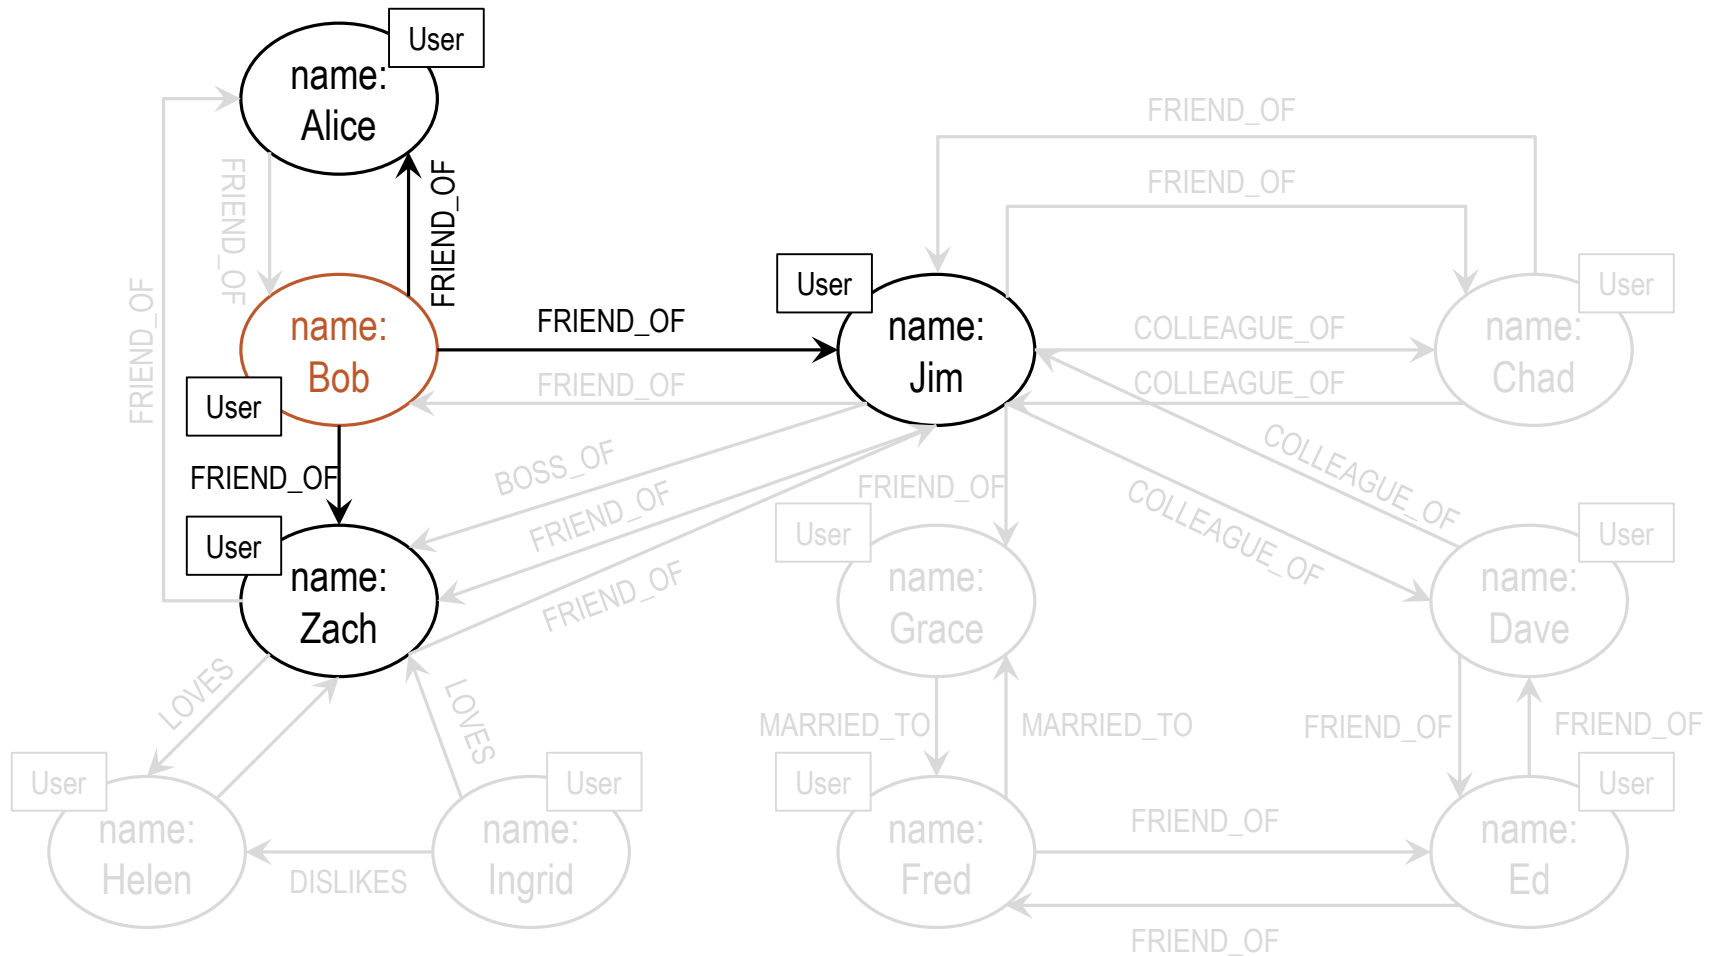
- Adapted from Robinson, Webber, & Eifrem, 2015



# Property graph data model

## Query

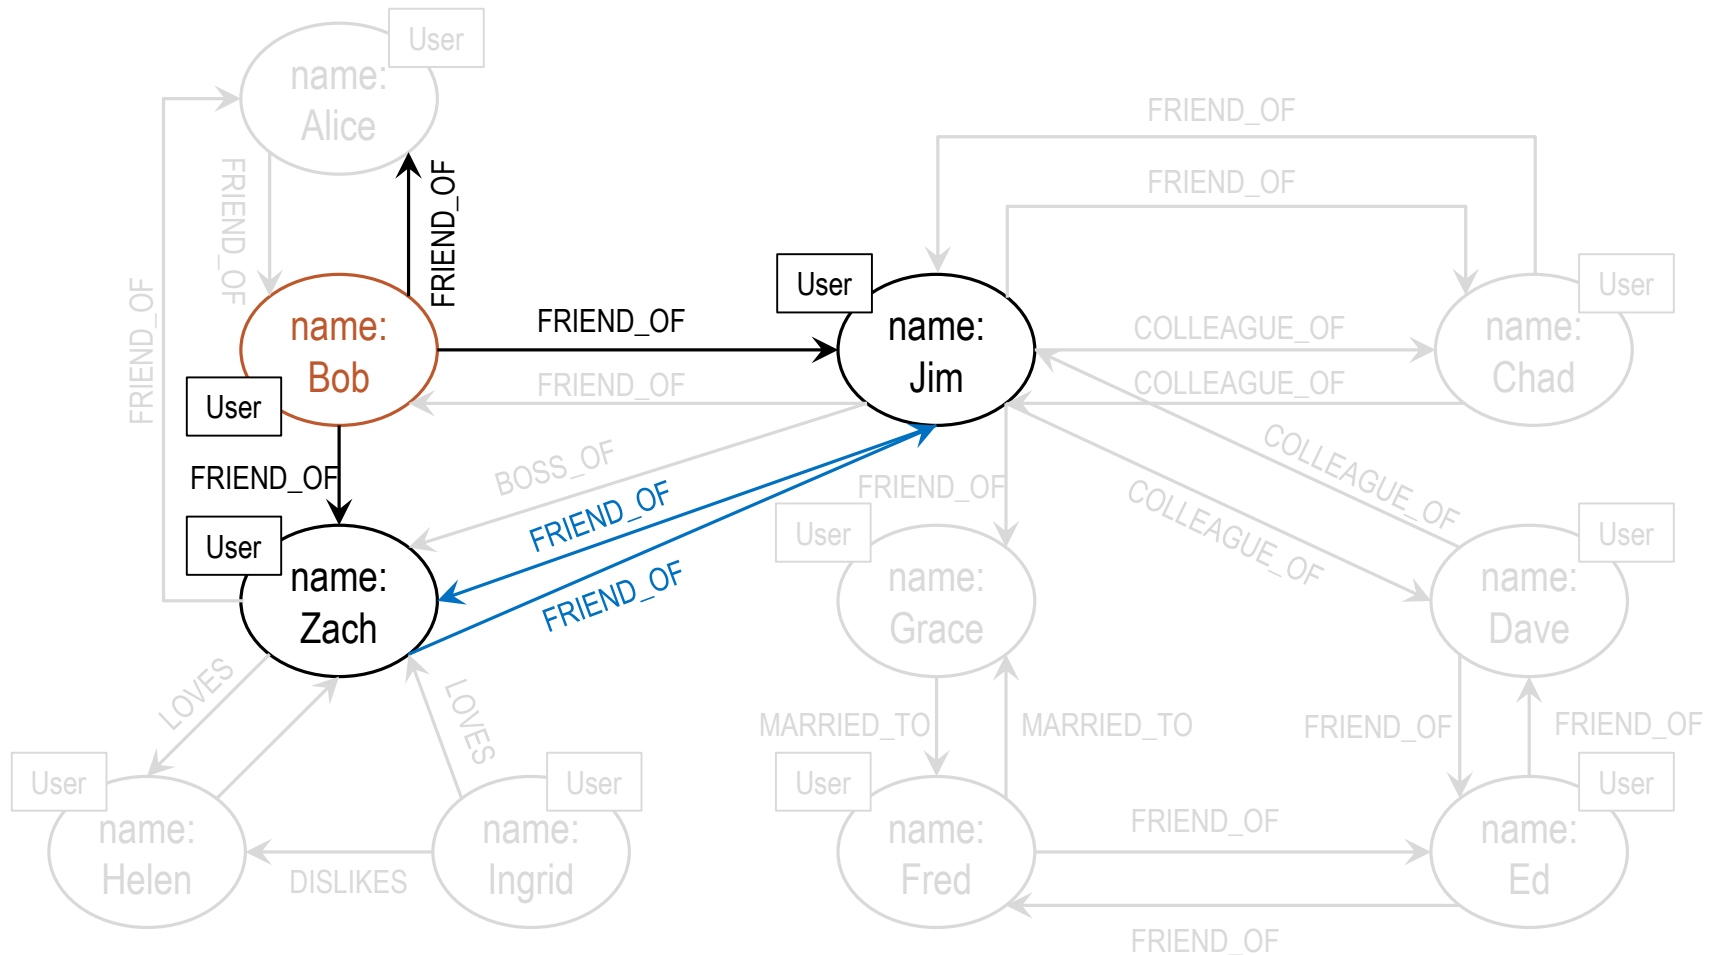
- Find all friends of Bob



# Property graph data model

## Query

- Find all users who are friends of Bob that shares similar friend as Bob



# Property graph data model

## Neo4j cypher query language

- Declarative language based on property graph model
- Example:
  - Find all users who are friends of Bob that shares similar friend as Bob

- Query:

```
MATCH  (e)<-[:FRIEND_OF]-(bob)-[:FRIEND_OF]->(f)-[:FRIEND_OF]->(e)
WHERE  bob.name = "Bob"
RETURN f.name AS name,
        count(e) AS score,
        collect(e.name) AS friends
ORDER BY score DESC
```

name	score	friends
"Zach"	2	["Alice", "Jim"]
"Jim"	1	["Zach"]

# Property graph data model

## Neo4j cypher query language

- Declarative language based on property graph model
- Example:
  - For each user, find the number of his/her direct/indirect friends
- Query:

```
MATCH (u:User)
OPTIONAL MATCH (u)-[:FRIEND_OF*]->(u2:User)
RETURN u.name AS name,
       count(DISTINCT u2) AS num
ORDER BY name
```

name	num
"Alice"	5
"Bob"	5
"Chad"	5
"Dave"	2
"Ed"	2
"Fred"	2
"Grace"	0
"Helen"	0
"Ingrid"	0
"Jim"	5
"Zach"	5

- Distributed DBMS
  - OLAP  
Multidimensional aggregation
  - NoSQL  
Graph database
  - Future  
NewSQL  
Database-as-a-Service
- 

Future

---

# NewSQL

## NewSQL database systems

- Targeted at OLTP workloads
- Features
  - Relational data model
  - SQL query language
  - ACID transactions
  - Runs on distributed cluster of shared-nothing nodes
- Some examples:
  - Clustrix
  - CockroachDB
  - Google Spanner
  - MemSQL
  - Microsoft Azure Cosmos DB
  - VoltDB

---

# Database-as-a-Service (DBaaS)

## RDBMS

- Amazon RDS (Amazon, Aurora, MySQL, MariaDB, SQL Server, Oracle, PostgreSQL)
  - <https://aws.amazon.com/rds/>
- Google Cloud SQL (MySQL, PostgreSQL)
  - <https://cloud.google.com/sql/>

## NoSQL

- Amazon DynamoDB
  - <https://aws.amazon.com/dynamodb/>
- Microsoft Azure Cosmos DB
  - <https://azure.microsoft.com/en-us/services/cosmos-db/>

## NewSQL

- Google Cloud Spanner
  - <https://cloud.google.com/spanner/>