

CS2102 Database Systems

Semester 1 2019/2020

Tutorial 05 (*Selected Answers*)

Quiz

Questions 1-6 uses the following database schema as Tutorial 03 Question 07. They are reproduced here for convenience.


- **Pizzas** (pizza)
- **Customers** (cname, area)
- **Restaurants** (rname, area)
- **Contains** (pizza, ingredient)
- **Sells** (rname, pizza, price)
- **Likes** (cname, pizza)

Pizzas indicates all the pizzas of interest. **Customers** indicates the name and location of each customer. **Restaurants** indicates the name and location of each restaurant. **Contains** indicates the ingredients used in each pizza. **Sells** indicates the pizzas sold by restaurants and their prices. **Likes** indicates the pizzas that customers like.

The following are all the foreign key constraints on the database schema:

- **Contains.pizza** is a foreign key that refers to **Pizzas.pizza**
- **Sells.rname** is a foreign key that refers to **Restaurants.name**
- **Sells.pizza** is a foreign key that refers to **Pizzas.pizza**
- **Likes.cname** is a foreign key that refers to **Customers.pizza**
- **Likes.pizza** is a foreign key that refers to **Pizzas.pizza**

Furthermore, we provide the simplified ER diagram corresponding to the schema above minus **Contains**.



1. For each pizza, find the restaurant that sells it at the most expensive price. Your output should be (rname, pizza, price). Exclude pizza that is not sold by any restaurants. Can you do it without using any aggregate function?

Solution:

```
SELECT S1.rname, S1.pizza, S1.price FROM Sells S1
WHERE NOT EXISTS (
    SELECT 1 FROM Sells S2
    WHERE S2.pizza = S1.pizza AND S2.price > S1.price
);
```

Alternative solution:

```
SELECT rname, pizza, price FROM Sells
WHERE price >= ALL ( SELECT price FROM Sells );
```

2. Find all restaurant pairs (rname1, rname2) such that the price of the most expensive pizza sold by rname1 is higher than that for rname2.

ER and SQL

Solution:

```
SELECT R1.rname, R2.rname
FROM   Restaurants R1, Restaurants R2
WHERE  ( SELECT MAX(price) FROM Sells WHERE rname = R1.rname )
      > ( SELECT MAX(price) FROM Sells WHERE rname = R2.rname );
```

Alternative solution:

```
WITH MaxPrice AS (
    SELECT rname, (SELECT MAX(price) FROM Sells WHERE rname = R.rname) AS maxP
    FROM   Restaurant R
)
SELECT R1.rname, R2.rname FROM MaxPrice R1 JOIN MaxPrice R2 ON R1.maxP > R2.maxP;
```

Alternative solution with restaurant not selling any pizza included with maxP = 0:

```
WITH MaxPrice AS (
    SELECT rname,
           COALESCE((SELECT MAX(price) FROM Sells WHERE rname = R.rname),0) AS maxP
    FROM   Restaurant R
)
SELECT R1.rname, R2.rname FROM MaxPrice R1 JOIN MaxPrice R2 ON R1.maxP > R2.maxP;
```

3. For each restaurant that sells some pizza, find the restaurant name and the average price of its pizza if its average price is higher than \$22. Can you do it without using HAVING clause in your answer?

Solution:

```
WITH RestaurantAvgPrice AS
( SELECT rname, AVG(price) AS avgPrice FROM Sells GROUP BY rname )
SELECT * FROM RestaurantAvgPrice WHERE avgPrice > 22;
```

4. For each restaurant R that sells some pizza, let totalPrice(R) denote the total price of all the pizzas sold by R. Find all pairs (rname, totalPrice(rname)) where totalPrice(rname) is higher than the average totalPrice() over all restaurants.

Solution:

```
SELECT rname, SUM(price) AS totalPrice FROM Sells S GROUP BY rname
HAVING SUM(price) > ( SELECT SUM(price) / COUNT(DISTINCT rname) FROM Sells );
```

Alternative solution:

```
WITH TotalPrice AS (
    SELECT rname, SUM(price) AS totalP FROM Sells GROUP BY rname
)
SELECT rname, totalP FROM TotalPrice
WHERE totalP > (Select AVG(totalP) FROM TotalPrice);
```

ER and SQL

Invalid solution:

```
SELECT rname, SUM(price) AS totalPrice FROM Sells S GROUP BY rname
HAVING totalPrice > ( SELECT SUM(price) / COUNT(DISTINCT rname) FROM Sells );
```

NOTE: totalPrice is undefined in the HAVING clause as the SELECT clause is evaluated after the HAVING clause.

Invalid solution:

```
SELECT rname, SUM(price) AS totalPrice FROM Sells S GROUP BY rname
HAVING SUM(price) > SUM(price) / COUNT(*);
```

NOTE: all three aggregate functions are computed w.r.t. a group

5. Find the customer pair (cname1, cname2) such that $\text{cname1} < \text{cname2}$ and they like exactly the same pizzas. Exclude customer pairs that do not like any pizza. Can you do it without using EXCEPT operator?

Solution:

```
SELECT C1.cname, C2.cname
FROM Customers C1 JOIN Customers C2 ON C1.cname < C2.cname
WHERE EXISTS ( SELECT 1 FROM Likes WHERE cname = C1.cname )
  AND NOT EXISTS (
    SELECT 1 FROM Likes L1
    WHERE cname = C1.cname
      AND NOT EXISTS ( SELECT 1 FROM Likes L2
        WHERE cname = C2.cname AND pizza = L1.pizza )
  )
  AND NOT EXISTS (
    SELECT 1 FROM Likes L2
    WHERE cname = C2.cname
      AND NOT EXISTS ( SELECT 1 FROM Likes L1
        WHERE cname = C1.cname AND pizza = L2.pizza )
  );
```

NOTE: C1 and C2 likes exactly the same pizza if and only if (a) for all pizza that C1 likes, C2 likes the pizza, and (b) for all pizza that C2 likes, C1 also likes that pizza. Using universal quantification trick of $\neg\neg(\forall F) \equiv \neg(\exists\neg F)$.

Alternative solution:

```
WITH NumLike AS
( SELECT cname, COUNT(*) AS num FROM Likes L GROUP BY cname ),
NumBothLike AS
( SELECT L1.cname AS cname1, L2.cname AS cname2, COUNT(*) AS num
  FROM Likes L1 JOIN Likes L2
    ON ( L1.pizza = L2.pizza ) AND ( L1.cname < L2.cname )
  GROUP BY L1.cname, L2.cname )
```

ER and SQL

```
SELECT cname1, cname2 FROM NumBothLike NBL
WHERE num = ( SELECT num FROM NumLike WHERE cname = B.cname1 )
AND num = ( SELECT num FROM NumLike WHERE cname = B.cname2 );
```

NOTE: Let S_1 be the set of pizza that C1 likes and S_2 be the set of pizza that C2 likes. Then $S_1 = S_2$ if and only if $|S_1 \cap S_2| = |S_1| = |S_2|$ where $|S|$ denotes the cardinality of the set S .

6. Write an SQL statement to increase the selling prices of pizzas as follows:
- Increase by \$3 if the restaurant is located in 'Central'
 - Increase by \$2 if the restaurant is located in 'East'
 - Otherwise, increase by \$1

Solution:

```
UPDATE Sells S
SET Price = CASE ( SELECT area FROM Restaurants WHERE rname = S.rname )
                WHEN 'Central' THEN price + 3
                WHEN 'East'     THEN price + 2
                ELSE price + 1
END;
```

Tutorial Questions

[Discussion: 7(a), 7(b), 7(c), 7(d), 7(e)]

7. Consider the following relational schema.

```
CREATE TABLE Students (
    sid      integer PRIMARY KEY,
    name     varchar(50) NOT NULL
);
CREATE TABLE Presenters (
    week     integer CHECK (week > 0),
    qnum     integer NOT NULL CHECK (qnum > 0),
    sid      integer REFERENCES Students (sid),
    PRIMARY KEY (week, sid)
);
```

The `Students` relation maintains information about students, and the `Presenters` relation maintains information about students who have presented solutions for tutorial questions. Specifically, a tuple (w, q, s) in `Presenters` relation means that the student with `sid s` presented tutorial questions `q` in week `w`. For each week, a student can present at most one question.

Assume that if the maximum week value in `Presenters` relation is `w`, then there is at least one record in `Presenters` for each week value from $\{1, 2, \dots, w\}$, and the next tutorial class will be in Week `w+1`.

Write an SQL query for each of the following questions. Remove duplicate records from all results. Try not to use `DISTINCT` keywords unless necessary.

- Find the `sid` of all students who have presented the most often.
- Find all `sid` pairs (s_1, s_2) such that $s_1 < s_2$ and both students have presented in the same week for at least 5 different weeks.
- Find all students who did not present for any three consecutive weeks. For example, if Alice presented only twice so far in week 2 and 6, then Alice should be in the query result as she did not present in weeks 3, 4, and 5.

- d) This question considers how to choose presenters for the next tutorial. Given a student with $sid\ s$, let $numQ(s)$ denote the total number of questions that s has presented so far. Let $lastWk(s)$ denote the most recent week number that s has presented. If s has not presented at all, then $numQ(s) = 0$ and $lastWk(s) = 0$.

Given two students with $sid\ s1$ and $s2$, $s1$ has a higher priority than $s2$ if one of the following conditions hold:

- i. $numQ(s1) < numQ(s2)$
- ii. $(numQ(s1) = numQ(s2))$ and $(lastWk(s1) < lastWk(s2))$
- iii. $(numQ(s1) = numQ(s2))$ and $(lastWk(s1) = lastWk(s2))$ and $(s1 < s2)$

Find all sets of two students S to be presenters for the next tutorial such that none of the students in $(Students - S)$ has higher priority than any of the students in S . The schema of your query result should be a pair of $sid\ (s1, s2)$, where $s1 < s2$ and each tuple in the query result represents a possible value for S .

- e) Let W denote the maximum week value in `Presenters`. We say that the `Presenters` table is consistent if it satisfies both the following conditions:

- i. There must exist at least one tuple t in `Presenters` with $t.week = w$ for each value of w in $\{1, 2, \dots, W\}$
- ii. For each value i in $\{1, 2, \dots, W\}$, if the maximum $qnum$ value for week i in `Presenters` is Q , then there must exist at least one tuple t in `Presenters` with $t.week = i$ and $t.qnum = j$ for each value of j in $\{1, 2, \dots, Q\}$.

Write SQL query to output 0 if `Presenters` is consistent; and 1 otherwise.

Solution:

- a) Solution #1:

```
SELECT sid FROM Presenters
GROUP BY sid
HAVING COUNT(*) >= ALL
      ( SELECT COUNT(*) FROM Presenters GROUP BY sid );
```

Solution #2:

```
WITH numPres AS
      ( SELECT sid, COUNT(*) AS num FROM Presenters GROUP BY sid )
SELECT sid FROM numPres WHERE num = ( SELECT MAX(num) FROM numPres );
```

- b) SQL Query:

```
SELECT P1.sid, P2.sid
FROM   Presenters P1 JOIN Presenters P2
      ON P1.sid < P2.sid AND P1.week = P2.week
GROUP BY P1.sid, P2.sid
HAVING COUNT(*) >= 5;
```

- c) Solution #1: If we have the students who did not present in any particular week, this problem is quite trivial. So what if we compute those first!

```
WITH NotPresent AS (
  SELECT sid, week
  FROM   ( SELECT sid FROM Students ) AS st_tbl,
        ( SELECT DISTINCT week FROM Presenters ) AS wk_tbl
  EXCEPT
  SELECT sid, week
  FROM   Presenters
)
SELECT DISTINCT P1.sid FROM NotPresent P1, NotPresent P2, NotPresent P3
```

ER and SQL

```
WHERE P1.sid = P2.sid AND P2.sid = P3.sid
AND P2.week = P1.week + 1 AND P3.week = P1.week + 2;
```

Solution #2: Let w denote the maximum week number from Presenters. A student with identifier s is in the output if $w \geq 3$ and there exists three consecutive weeks, say starting from week w to week $w+2$ such that s did not present in these three weeks. More precisely, s is in the output if $w \geq 3$ and there exists some w in $\{1, \dots, w-2\}$ where $\{w, w+1, w+2\} \cap \pi_{\text{week}}(\sigma_{\text{sid}=s}(\text{Presenters})) = \emptyset$.

```
WITH Weeks AS (
    SELECT DISTINCT Week FROM Presenters
    WHERE week + 2 <= ( SELECT MAX(week) FROM Presenters )
)
SELECT S.sid FROM Students S
WHERE ( SELECT MAX(week) FROM Presenters ) >= 3
AND EXISTS (
    SELECT 1 FROM Weeks W
    WHERE NOT EXISTS (
        SELECT week FROM Weeks W2
        WHERE W2.week >= W.week
        AND W2.week <= W.week + 2
        INTERSECT
        SELECT week FROM Presenters P
        WHERE P.sid = S.sid
    )
);
```

Solution #3: Case analysis; case (4) are saying that for any 2 presentations (p_1, p_2) , the distance is ≥ 3 weeks and there are no other presentations (p_3) in between for ALL (p_1, p_2) . The rest are simply edge cases.

```
-- Case 1: not presented and  $w \geq 3$ 
SELECT sid FROM Students
WHERE sid NOT IN (SELECT sid FROM Presenters)
AND (SELECT MAX(week) FROM Presenters) >= 3
UNION
SELECT sid FROM Presenters P1
WHERE
-- Case 2: first presentation is week 4 or after
( SELECT MIN(week) FROM Presenters WHERE sid = P1.sid ) >= 4
OR
-- Case 3: last presentation is week  $w-3$  or before
( SELECT MAX(week) FROM Presenters ) -
( SELECT MAX(week) FROM Presenters WHERE sid = P1.sid ) >= 3
OR
-- Case 4: forall(P1,P2), not-exists(P3)
-- or, forall(P1), exists(P2), not-exists(P3)
EXISTS (
    SELECT 1 FROM Presenters P2
    WHERE P1.sid = P2.sid AND P2.week
    AND NOT EXISTS (
        SELECT 1 FROM Presenters P3
        WHERE P3.sid = P1.sid AND P1.week < P3.week AND P3.week < P2.week
    )
);
```

ER and SQL

Wrong Answer #1:

```
SELECT sid FROM Students
WHERE ( SELECT MAX(week) FROM Presenters ) >= 3
EXCEPT
SELECT sid FROM Students S
WHERE EXISTS (
    SELECT 1 FROM Presenters P1, Presenters P2, Presenters P3
    WHERE P1.sid = S.sid AND P2.sid = S.sid AND P3.sid = S.sid
    AND P2.week = P1.week + 1 AND P3.week = P1.week + 2
);
```

The answer above is *incorrect* because a student who has presented for 3 consecutive weeks (e.g., in week 1, 2 and 3) but did not present in the next 3 consecutive weeks (e.g., in week 4, 5, and 6) would be incorrectly excluded from the output.

d) Solution #1:

```
WITH StudentInfo AS (
    SELECT sid, COUNT(*) AS numQ, MAX(week) AS lastWk
    FROM Presenters
    GROUP BY sid
    UNION
    SELECT sid, 0, 0
    FROM Students
    WHERE sid NOT IN ( SELECT sid FROM Presenters )
),
StudentSet AS (
    SELECT sid FROM StudentInfo
    ORDER BY numQ, lastWk, sid
    LIMIT 2
)
SELECT MIN(sid) AS sid1, MAX(sid) AS sid2 FROM StudentSet;
```

Solution #2:

```
WITH StudentInfo AS (
    SELECT P.sid, COUNT(*) AS numQ, MAX(week) AS lastWk
    FROM Presenters P
    GROUP BY sid
    UNION
    SELECT sid, 0, 0
    FROM Students
    WHERE sid NOT IN ( SELECT sid FROM Presenters )
)
SELECT S1.sid, S2.sid
FROM StudentInfo S1 JOIN StudentInfo S2 ON S1.sid < S2.sid
WHERE NOT EXISTS (
    SELECT 1 FROM StudentInfo S3
    WHERE S3.sid <> S1.sid AND S3.sid <> S2.sid
    AND ( ( S3.numQ < S1.numQ )
    OR ( S3.numQ < S2.numQ )
    OR ((S3.numQ = S1.numQ) AND (S3.lastWk < S1.lastWk))
    OR ((S3.numQ = S2.numQ) AND (S3.lastWk < S2.lastWk))
    OR ((S3.numQ = S1.numQ) AND (S3.lastWk = S1.lastWk)
    AND (S3.sid < S1.sid))
    OR ((S3.numQ = S2.numQ) AND (S3.lastWk = S2.lastWk)
    AND (S3.sid < S2.sid))
    )
);
```

ER and SQL

);

e) SQL Query:

```
WITH Status AS (  
    SELECT 1 AS num FROM Presenters  
    WHERE ( SELECT MAX(week) FROM Presenters )  
           <>  
           ( SELECT COUNT(DISTINCT week) FROM Presenters )  
    OR EXISTS (  
        SELECT 1 FROM Presenters  
        GROUP BY week  
        HAVING MAX(qnum) <> COUNT(DISTINCT qnum)  
    )  
)  
SELECT COALESCE(MIN(num), 0) FROM Status;
```